

# Teach your (micro)services speak Protocol Buffers with gRPC.

Mihai Iachimovschi

 @mishunika

 [mihai.iachimovschi@gmail.com](mailto:mihai.iachimovschi@gmail.com)

# What's inside?

# What's inside?

- Message serialization and deserialization

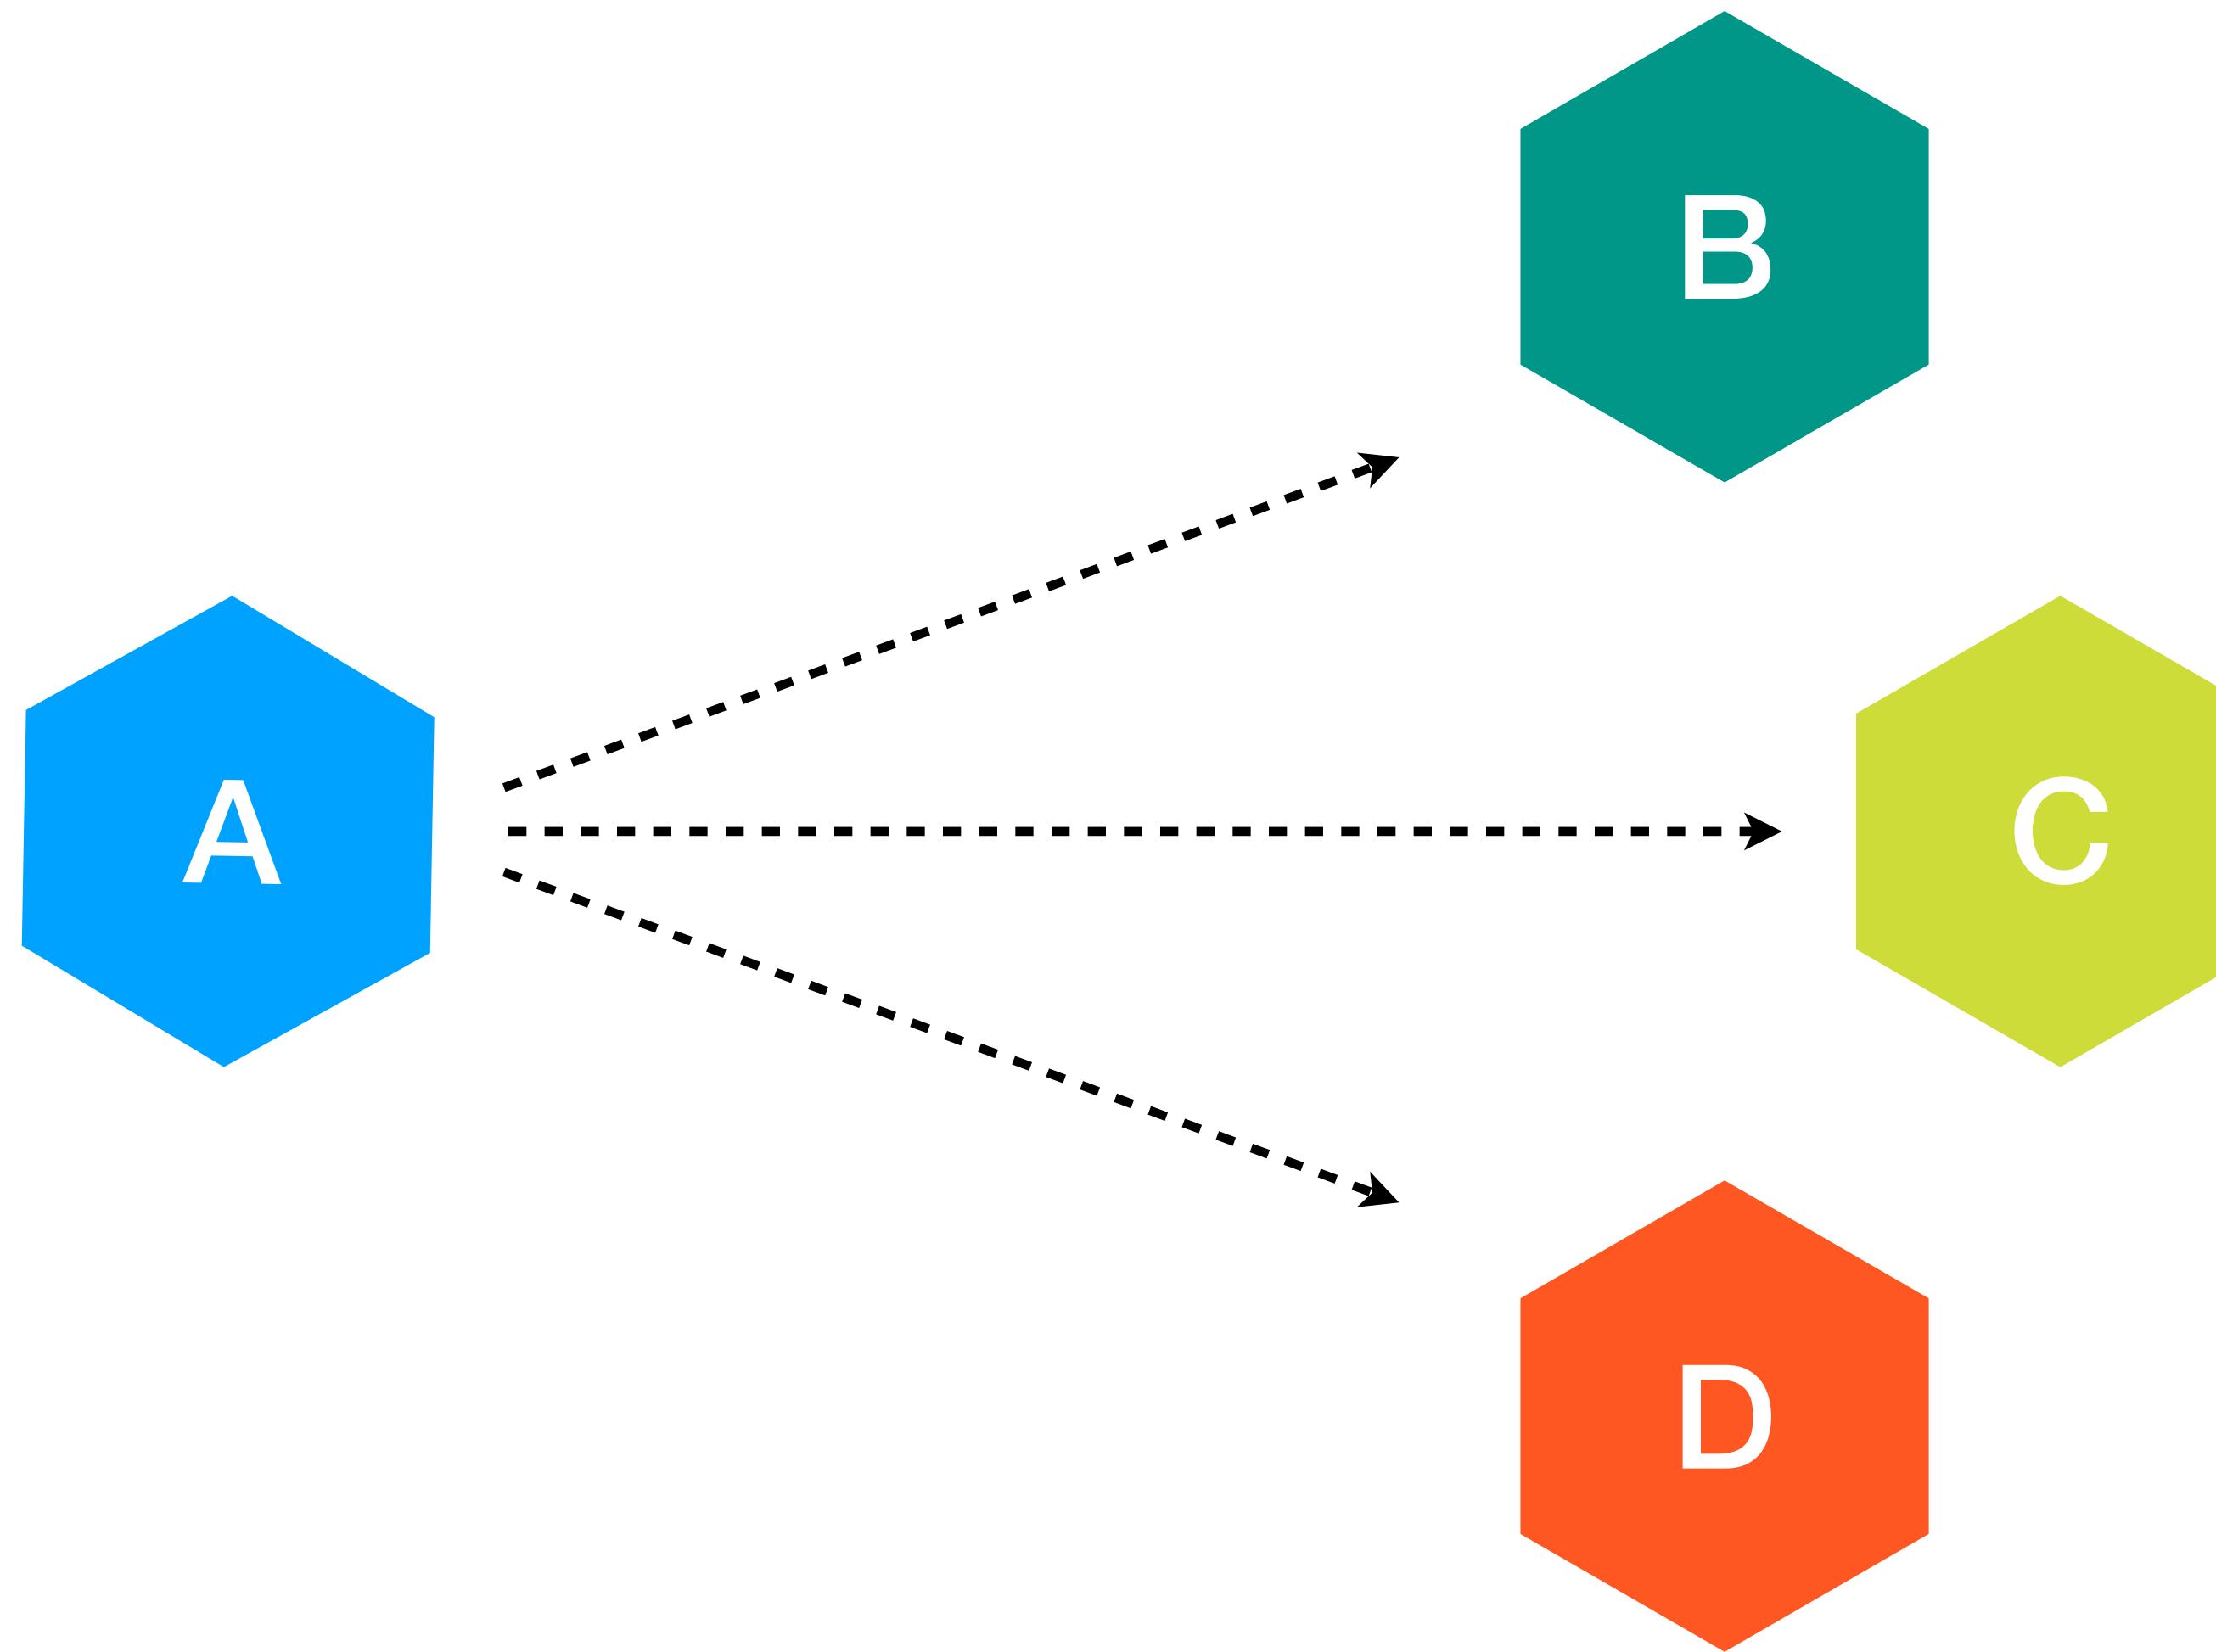
# What's inside?

- Message serialization and deserialization
- Message transport

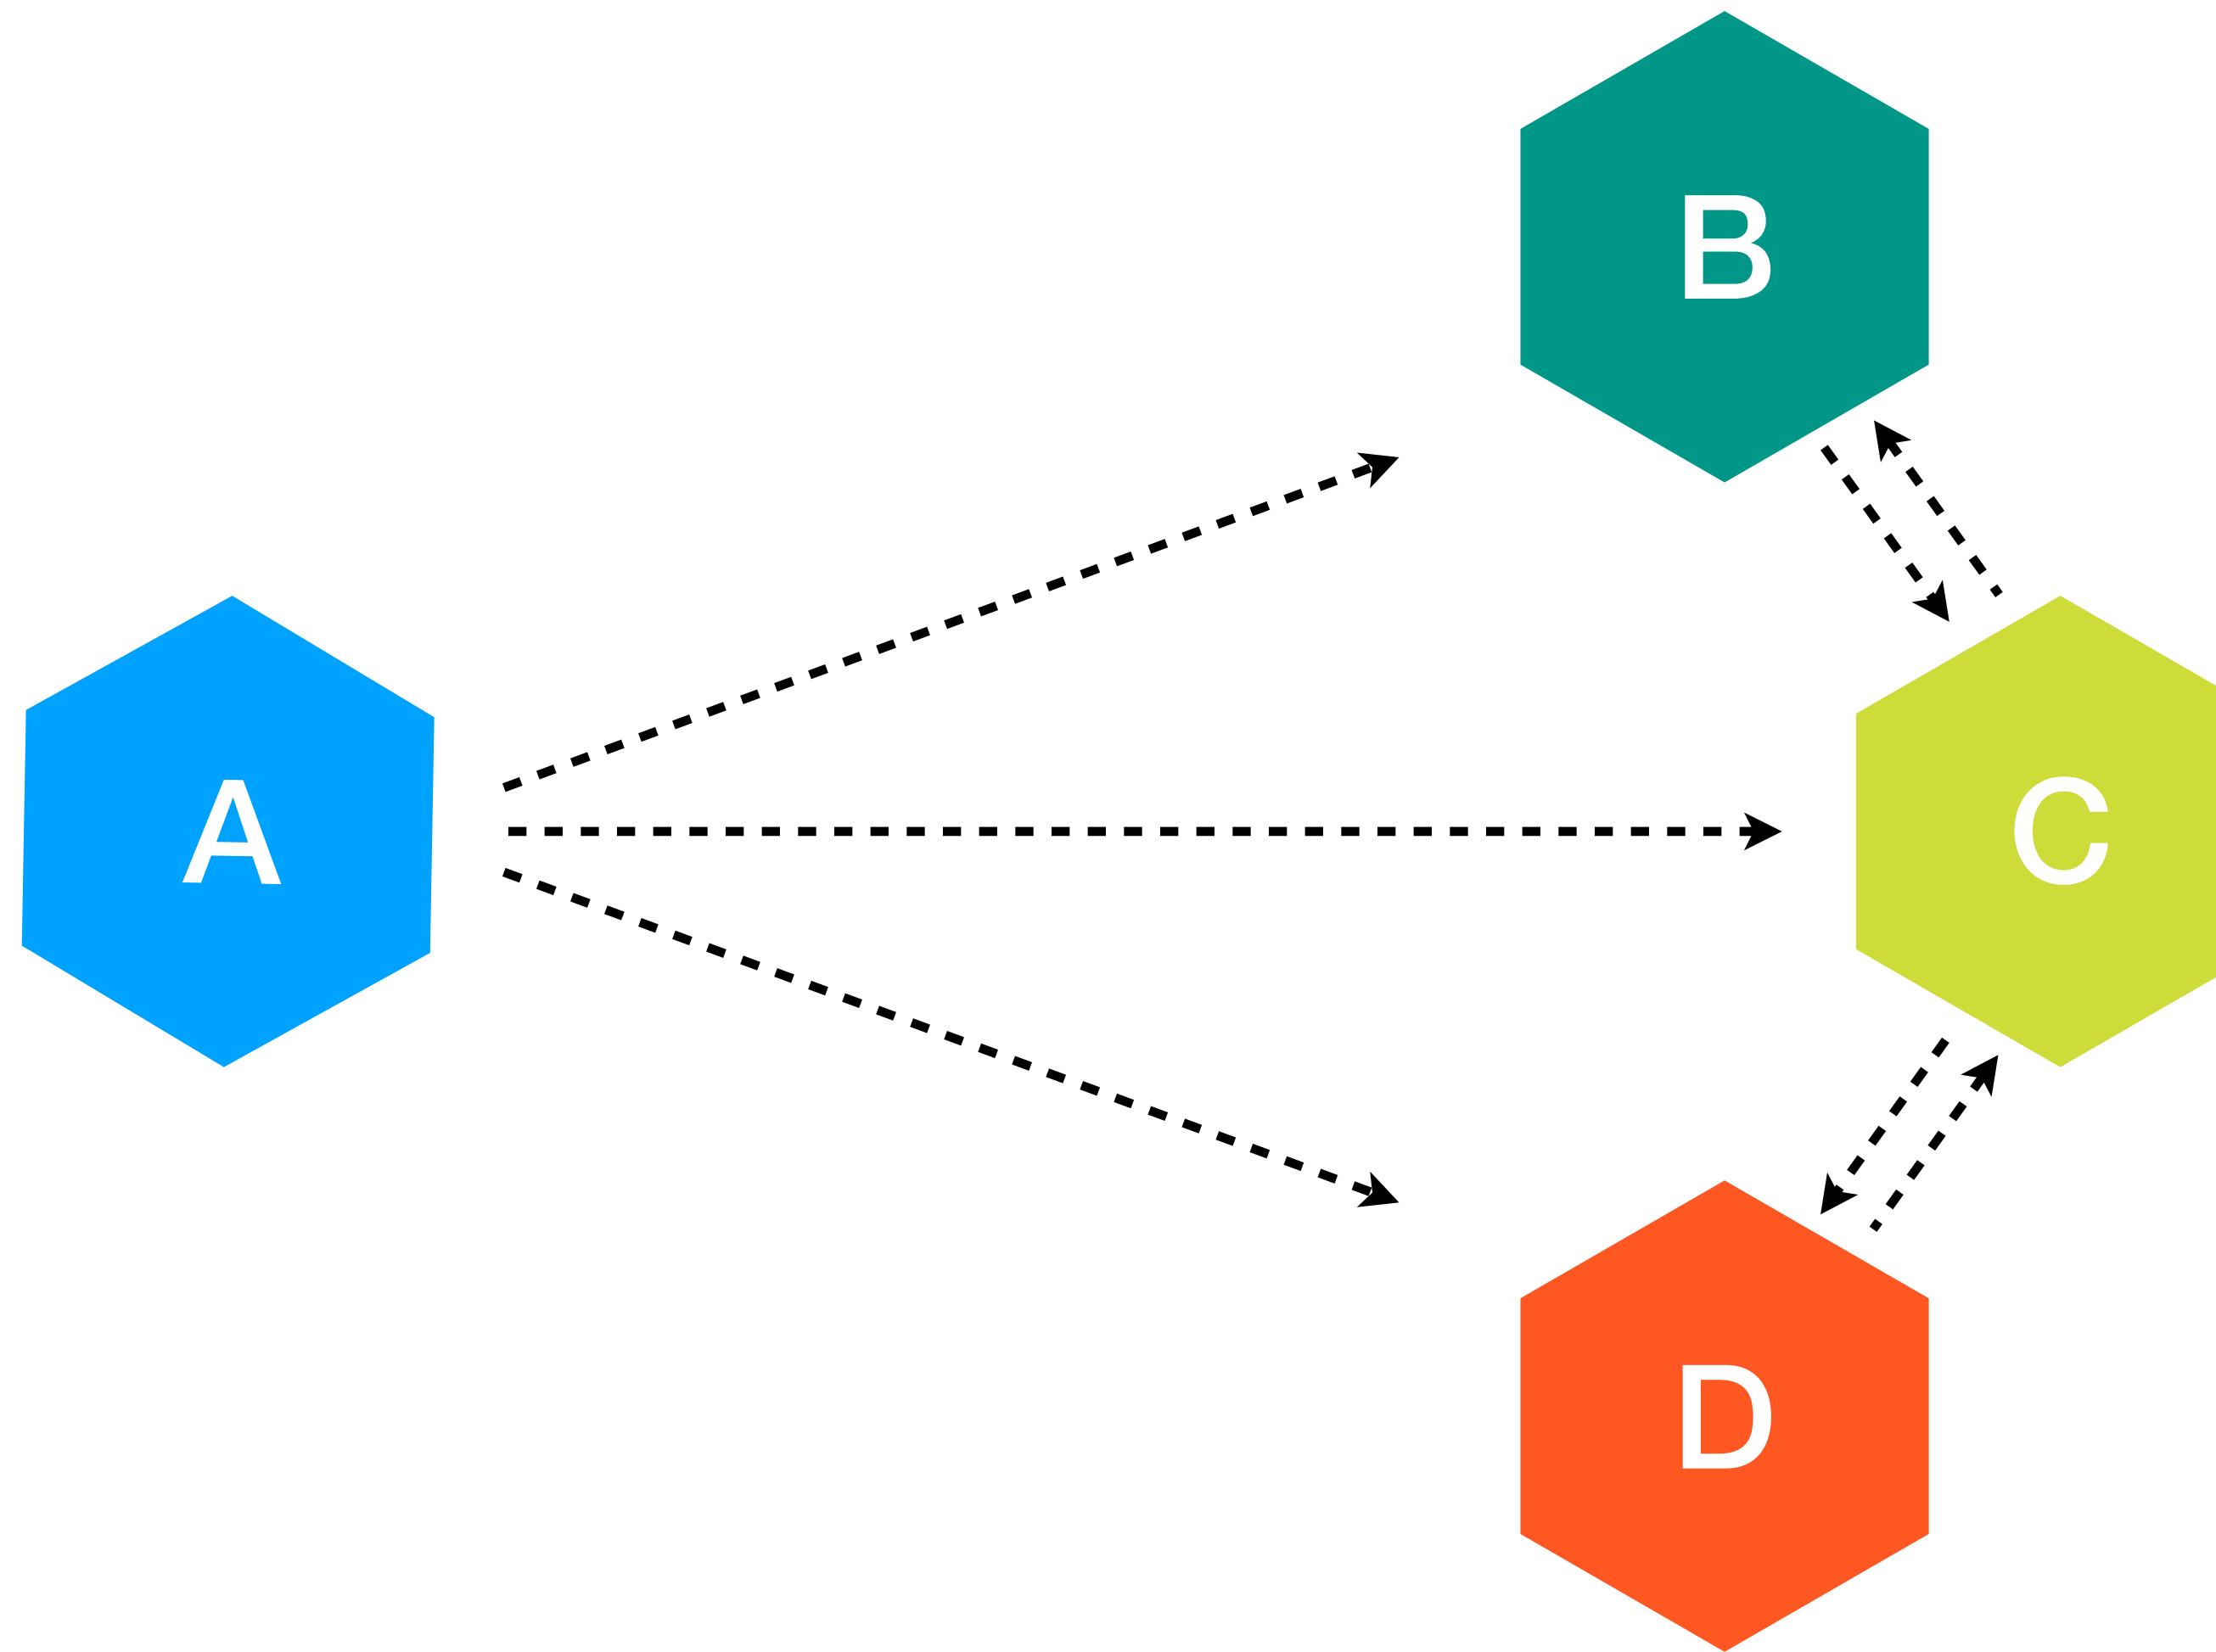
# What's inside?

- Message serialization and deserialization
- Message transport
- Services diversity

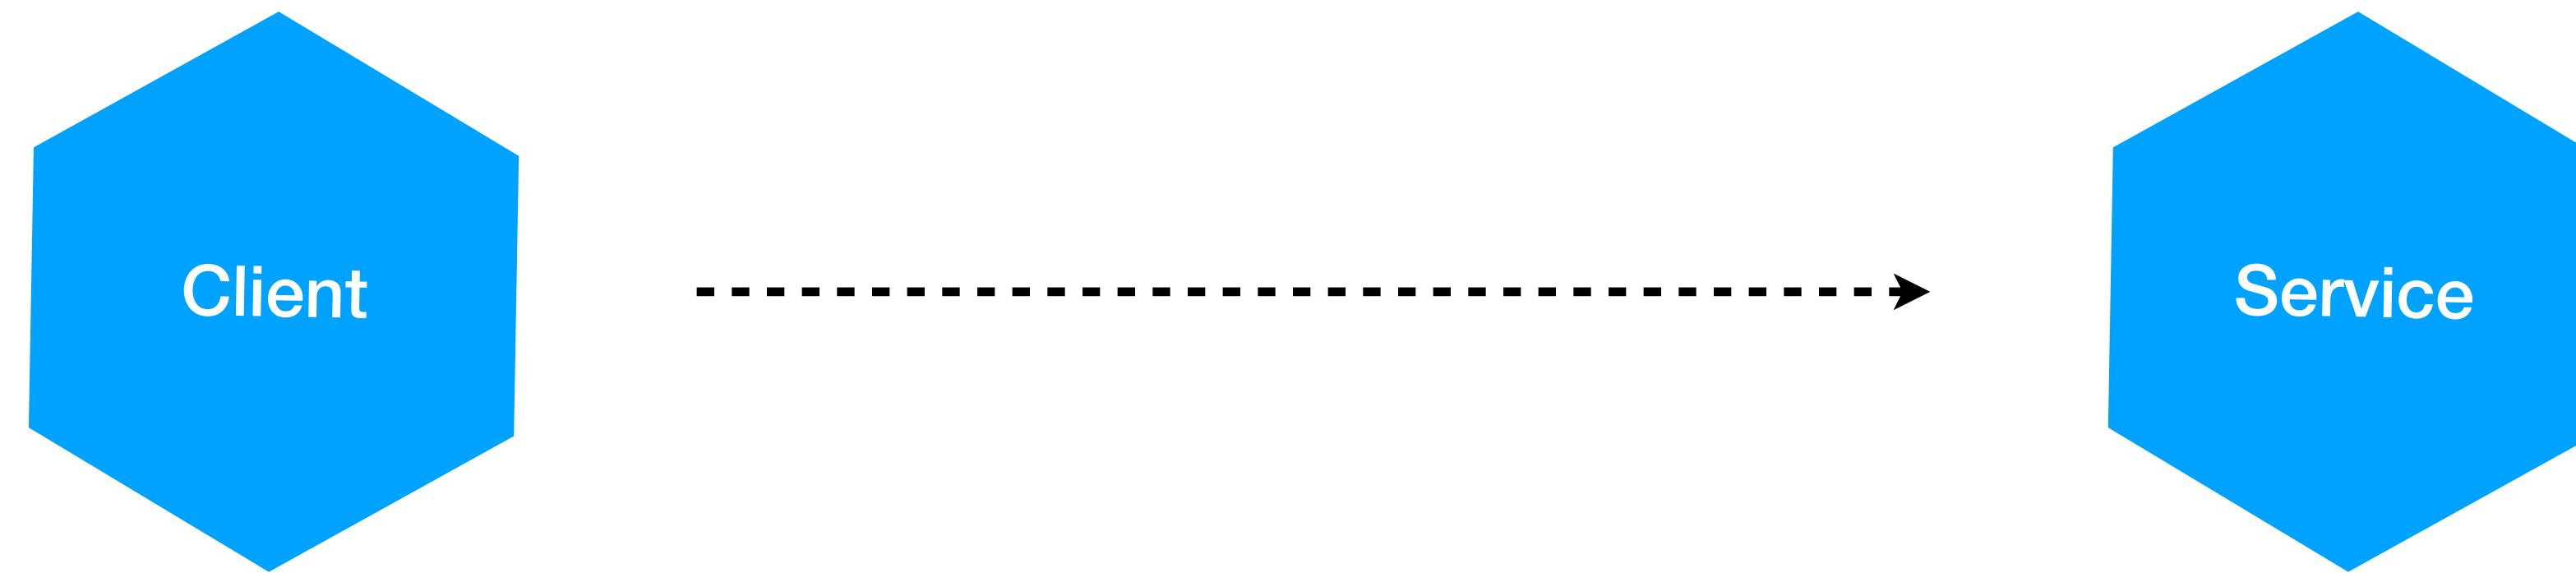
# How it works in real life



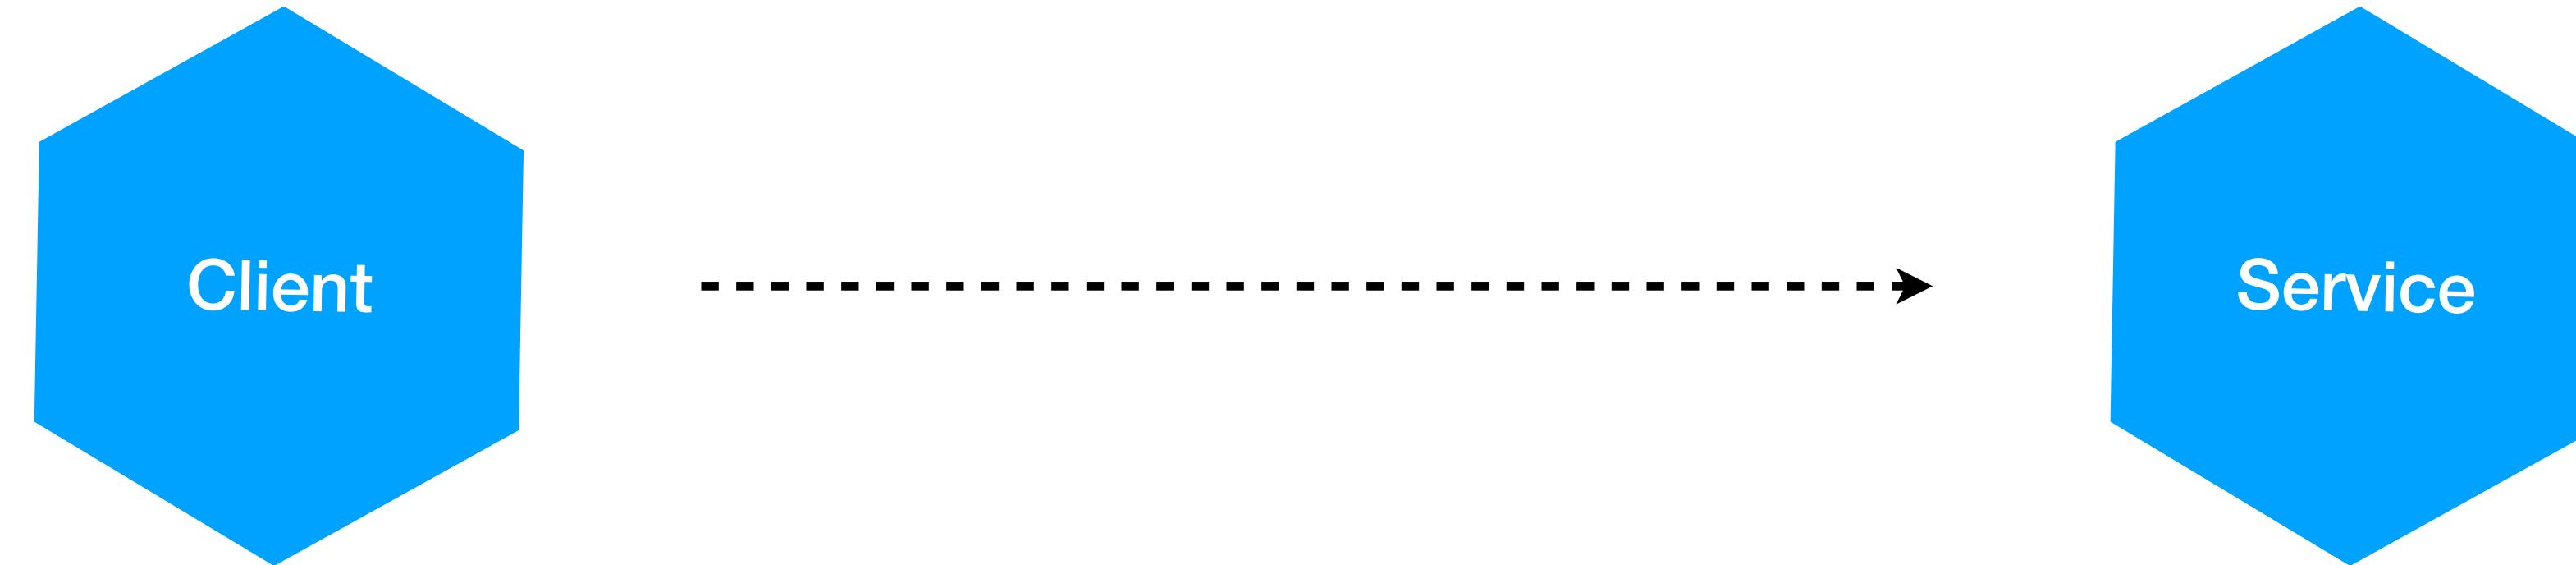
# How it works in real life



# In a nutshell

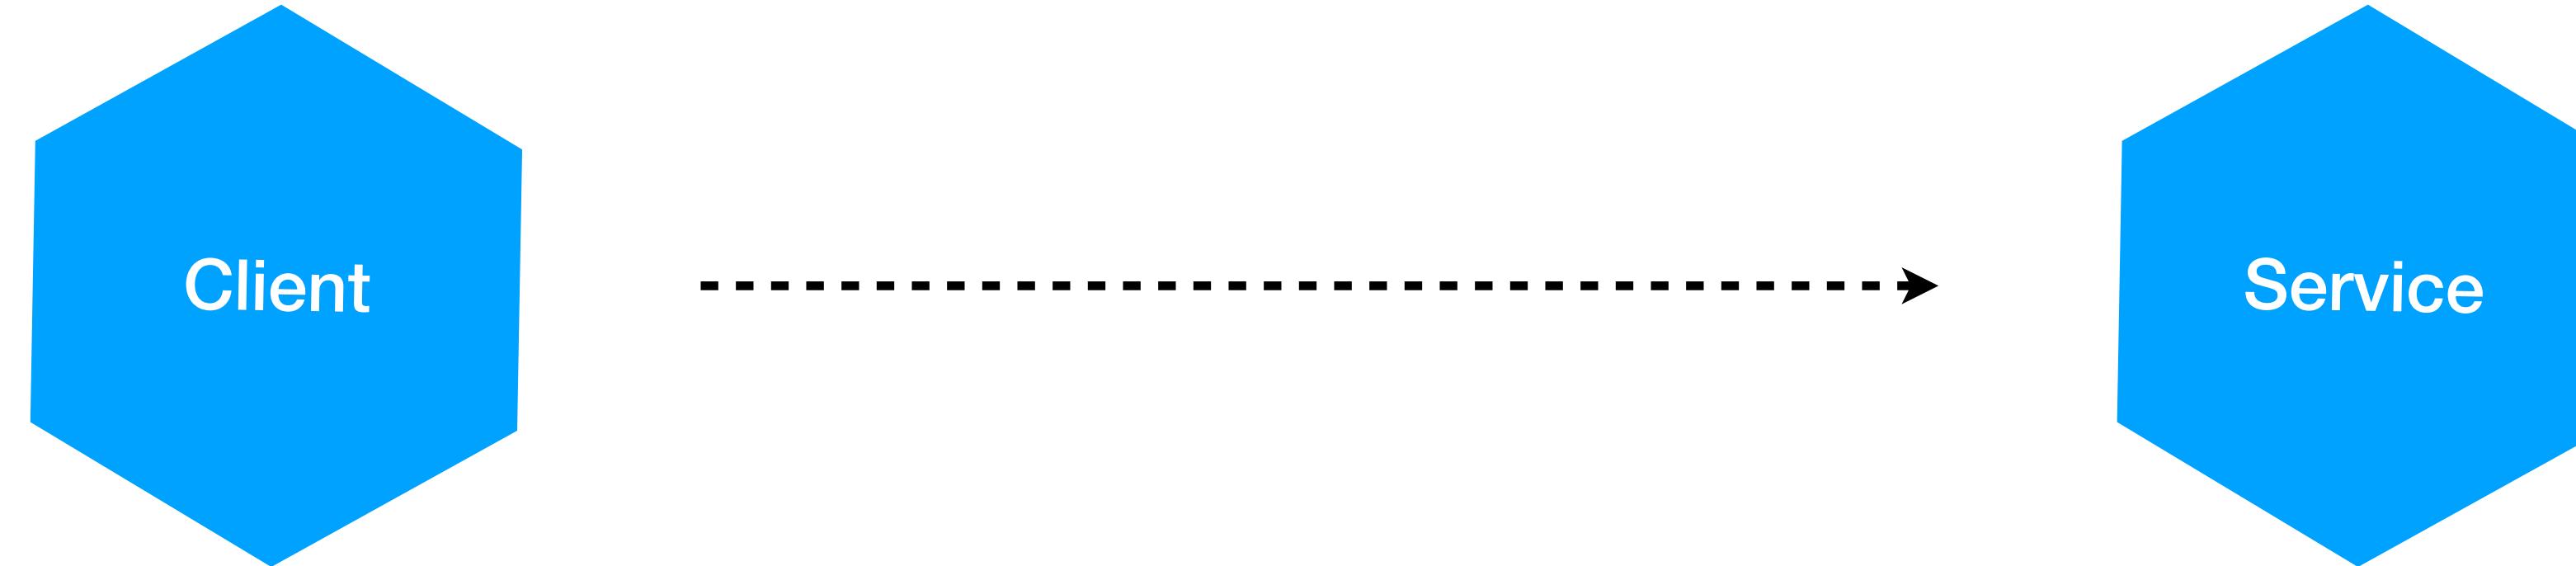


# In a nutshell



- Over HTTP
- Serialized to JSON

# In a nutshell



- Over HTTP
- Serialized to JSON
- Proprietary protocol
- Remote objects

# JSON advantages

# JSON advantages

- Human readable

# JSON advantages

- Human readable
- Schema-less

# JSON advantages

- Human readable
- Schema-less
- Language agnostic

# JSON disadvantages

# JSON disadvantages

- Human readable

# JSON disadvantages

–Human readable    Isn't it a benefit?

# JSON disadvantages

- Human readable      Isn't it a benefit?
- Schema-less

# JSON disadvantages

- Human readable      Isn't it a benefit?
- Schema-less           Isn't it a benefit as well?

# JSON disadvantages

- Human readable      Isn't it a benefit?
- Schema-less          Isn't it a benefit as well?
- Type-less

# Protocol Buffers?

“Protocol buffers are Google's language-neutral, platform-neutral, extensible mechanism for serializing structured data – think XML, but smaller, faster, and simpler.”

- <https://developers.google.com/protocol-buffers/>

# Protocol Buffers?

“Protocol buffers are Google's language-neutral, platform-neutral, extensible mechanism for serializing structured data – think ~~XML~~, but smaller, faster, and simpler.”

- <https://developers.google.com/protocol-buffers/>

# Protocol Buffers?

“Protocol buffers are Google's language-neutral, platform-neutral, extensible mechanism for serializing structured data – think ~~XML~~, but smaller, faster, and simpler.”

**JSON**

- <https://developers.google.com/protocol-buffers/>

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# Protocol Buffers example

```
message Person {  
    reserved 1, 2, 5;  
    reserved "name";  
    int32 id = 3;  
    repeated string aliases = 4;  
}
```

# Why protocol buffers

Binary format

# Why protocol buffers

Enforcing the schema

# Why protocol buffers

Language neutral

# Why protocol buffers

Out-of the box backward compatibility

# Why protocol buffers

Out-of-the-box backward compatibility

# Why protocol buffers

Out-of-the-box backward compatibility

```
if version == 3:  
    ...  
elif version > 4:  
    if (version == 5):  
        ...  
    ...  
    ...
```

# Why protocol buffers

Generally faster

# How to...

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# How to...

```
message Person {  
    string name = 1;           ← Proto message definition  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

# How to...

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```



Proto message definition

```
Person john = Person.newBuilder()  
    .setId(1234)  
    .setName("John Doe")  
    .addAliases("ionel")  
    .build();
```

# How to...

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

Proto message definition

Object creation

```
Person john = Person.newBuilder()  
    .setId(1234)  
    .setName("John Doe")  
    .addAliases("ionel")  
    .build();
```

# How to...

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

Proto message definition

Object creation



```
john = Person()  
john.id = 1234  
john.name = 'John Doe'  
john_aliases.add('ionel')
```

# How to...

```
message Person {  
    string name = 1;  
    int32 id = 2;  
    repeated string aliases = 3;  
}
```

Proto message definition

Object creation →

```
john = Person(id=1234,  
              name='John Doe',  
              aliases=['ionel'])
```

# Communication (REST-ish)

1. URI: <https://api.example.com/person/42>
2. Make an HTTP GET Request
3. Receive a plaintext JSON
4. Parse it
5. ...

# Request

GET /person/42 HTTP/1.1

Accept: \*/\*

Accept-Encoding: gzip, deflate

Connection: keep-alive

Host: api.example.com

...

# Response headers

HTTP/1.1 200 OK

Access-Control-Allow-Credentials: true

Cache-Control: public, max-age=14400

Content-Encoding: gzip

Content-Type: application/json;

charset=utf-8

...

# Response body

```
{  
  "name": "John Doe",  
  "id": 42,  
  "aliases": [  
    "ionel",  
    "honzík"  
  ]  
}
```

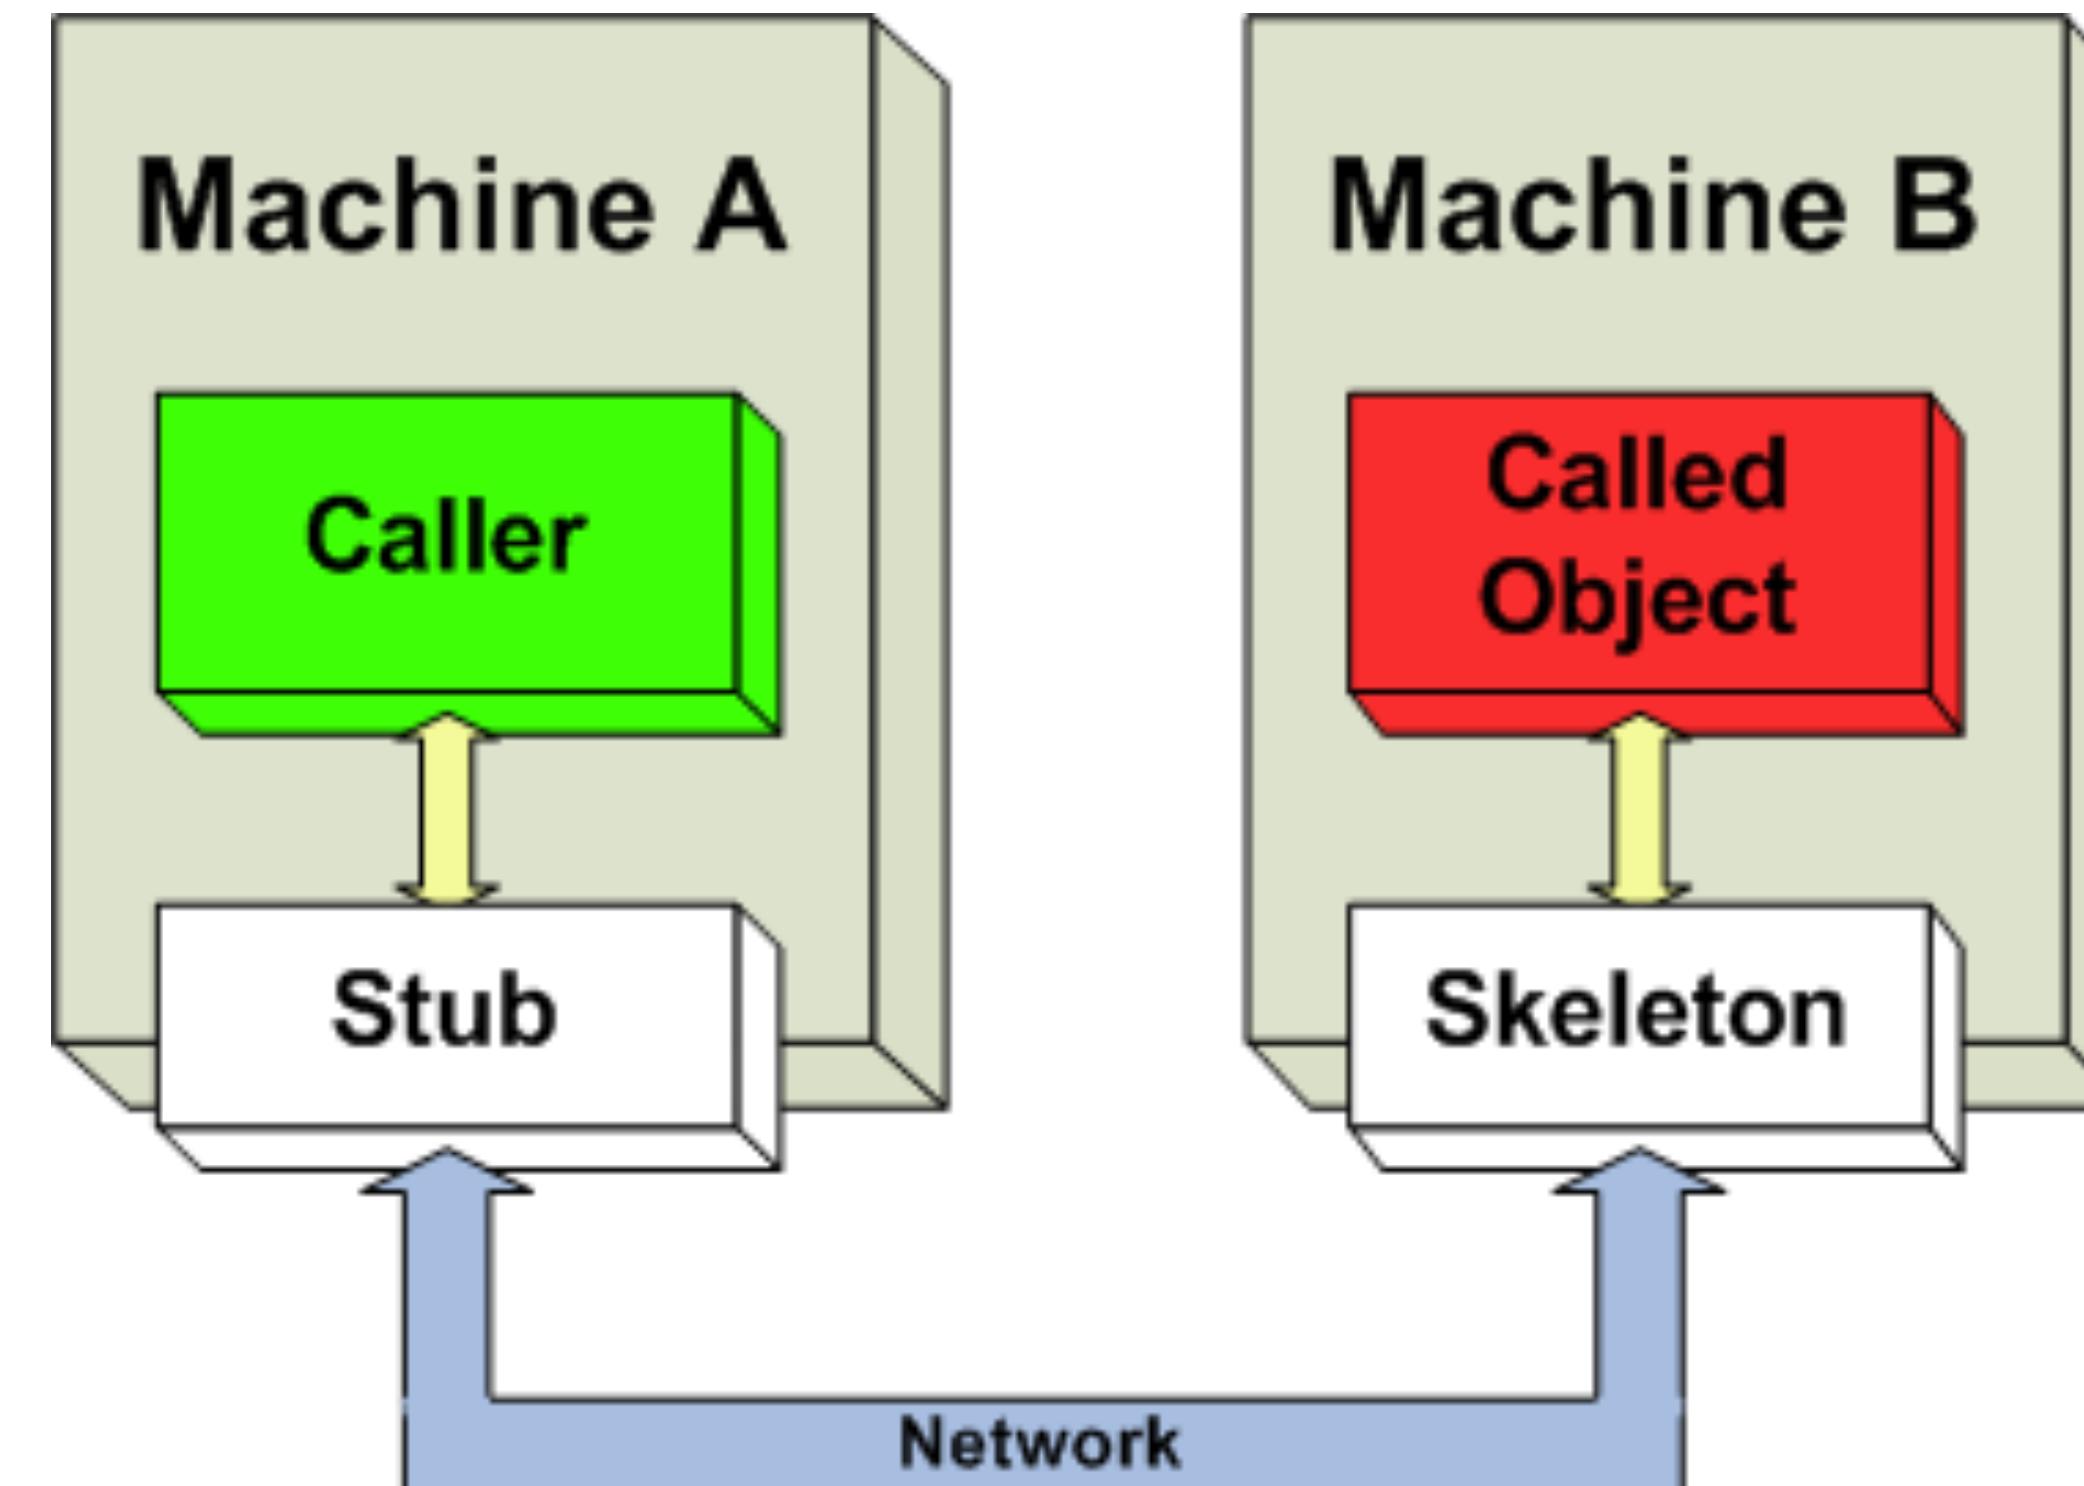
# HTTP

## HISTORY & PERFORMANCE



<https://www.youtube.com/watch?v=DtTKF5OcpsU>

# Distributed objects



[https://en.wikipedia.org/wiki/Distributed\\_object](https://en.wikipedia.org/wiki/Distributed_object)

# Distributed objects



**“First Law of Distributed Object Design:  
don't distribute your objects”.**

— Martin Fowler

<https://martinfowler.com/articles/distributed-objects-microservices.html>

# The 8 Fallacies of distributed computing

1. The network is reliable.
2. Latency is zero.
3. Bandwidth is infinite.
4. The network is secure.
5. Topology doesn't change.
6. There is one administrator.
7. Transport cost is zero.
8. The network is homogeneous.

<http://www.rgoarchitects.com/Files/fallacies.pdf>

# Keep in mind...

“Anything that can go wrong will go wrong.”

—Murphy’s Law

# So, what's next?

↑GRPC↓



**Services not Objects, Messages not References**

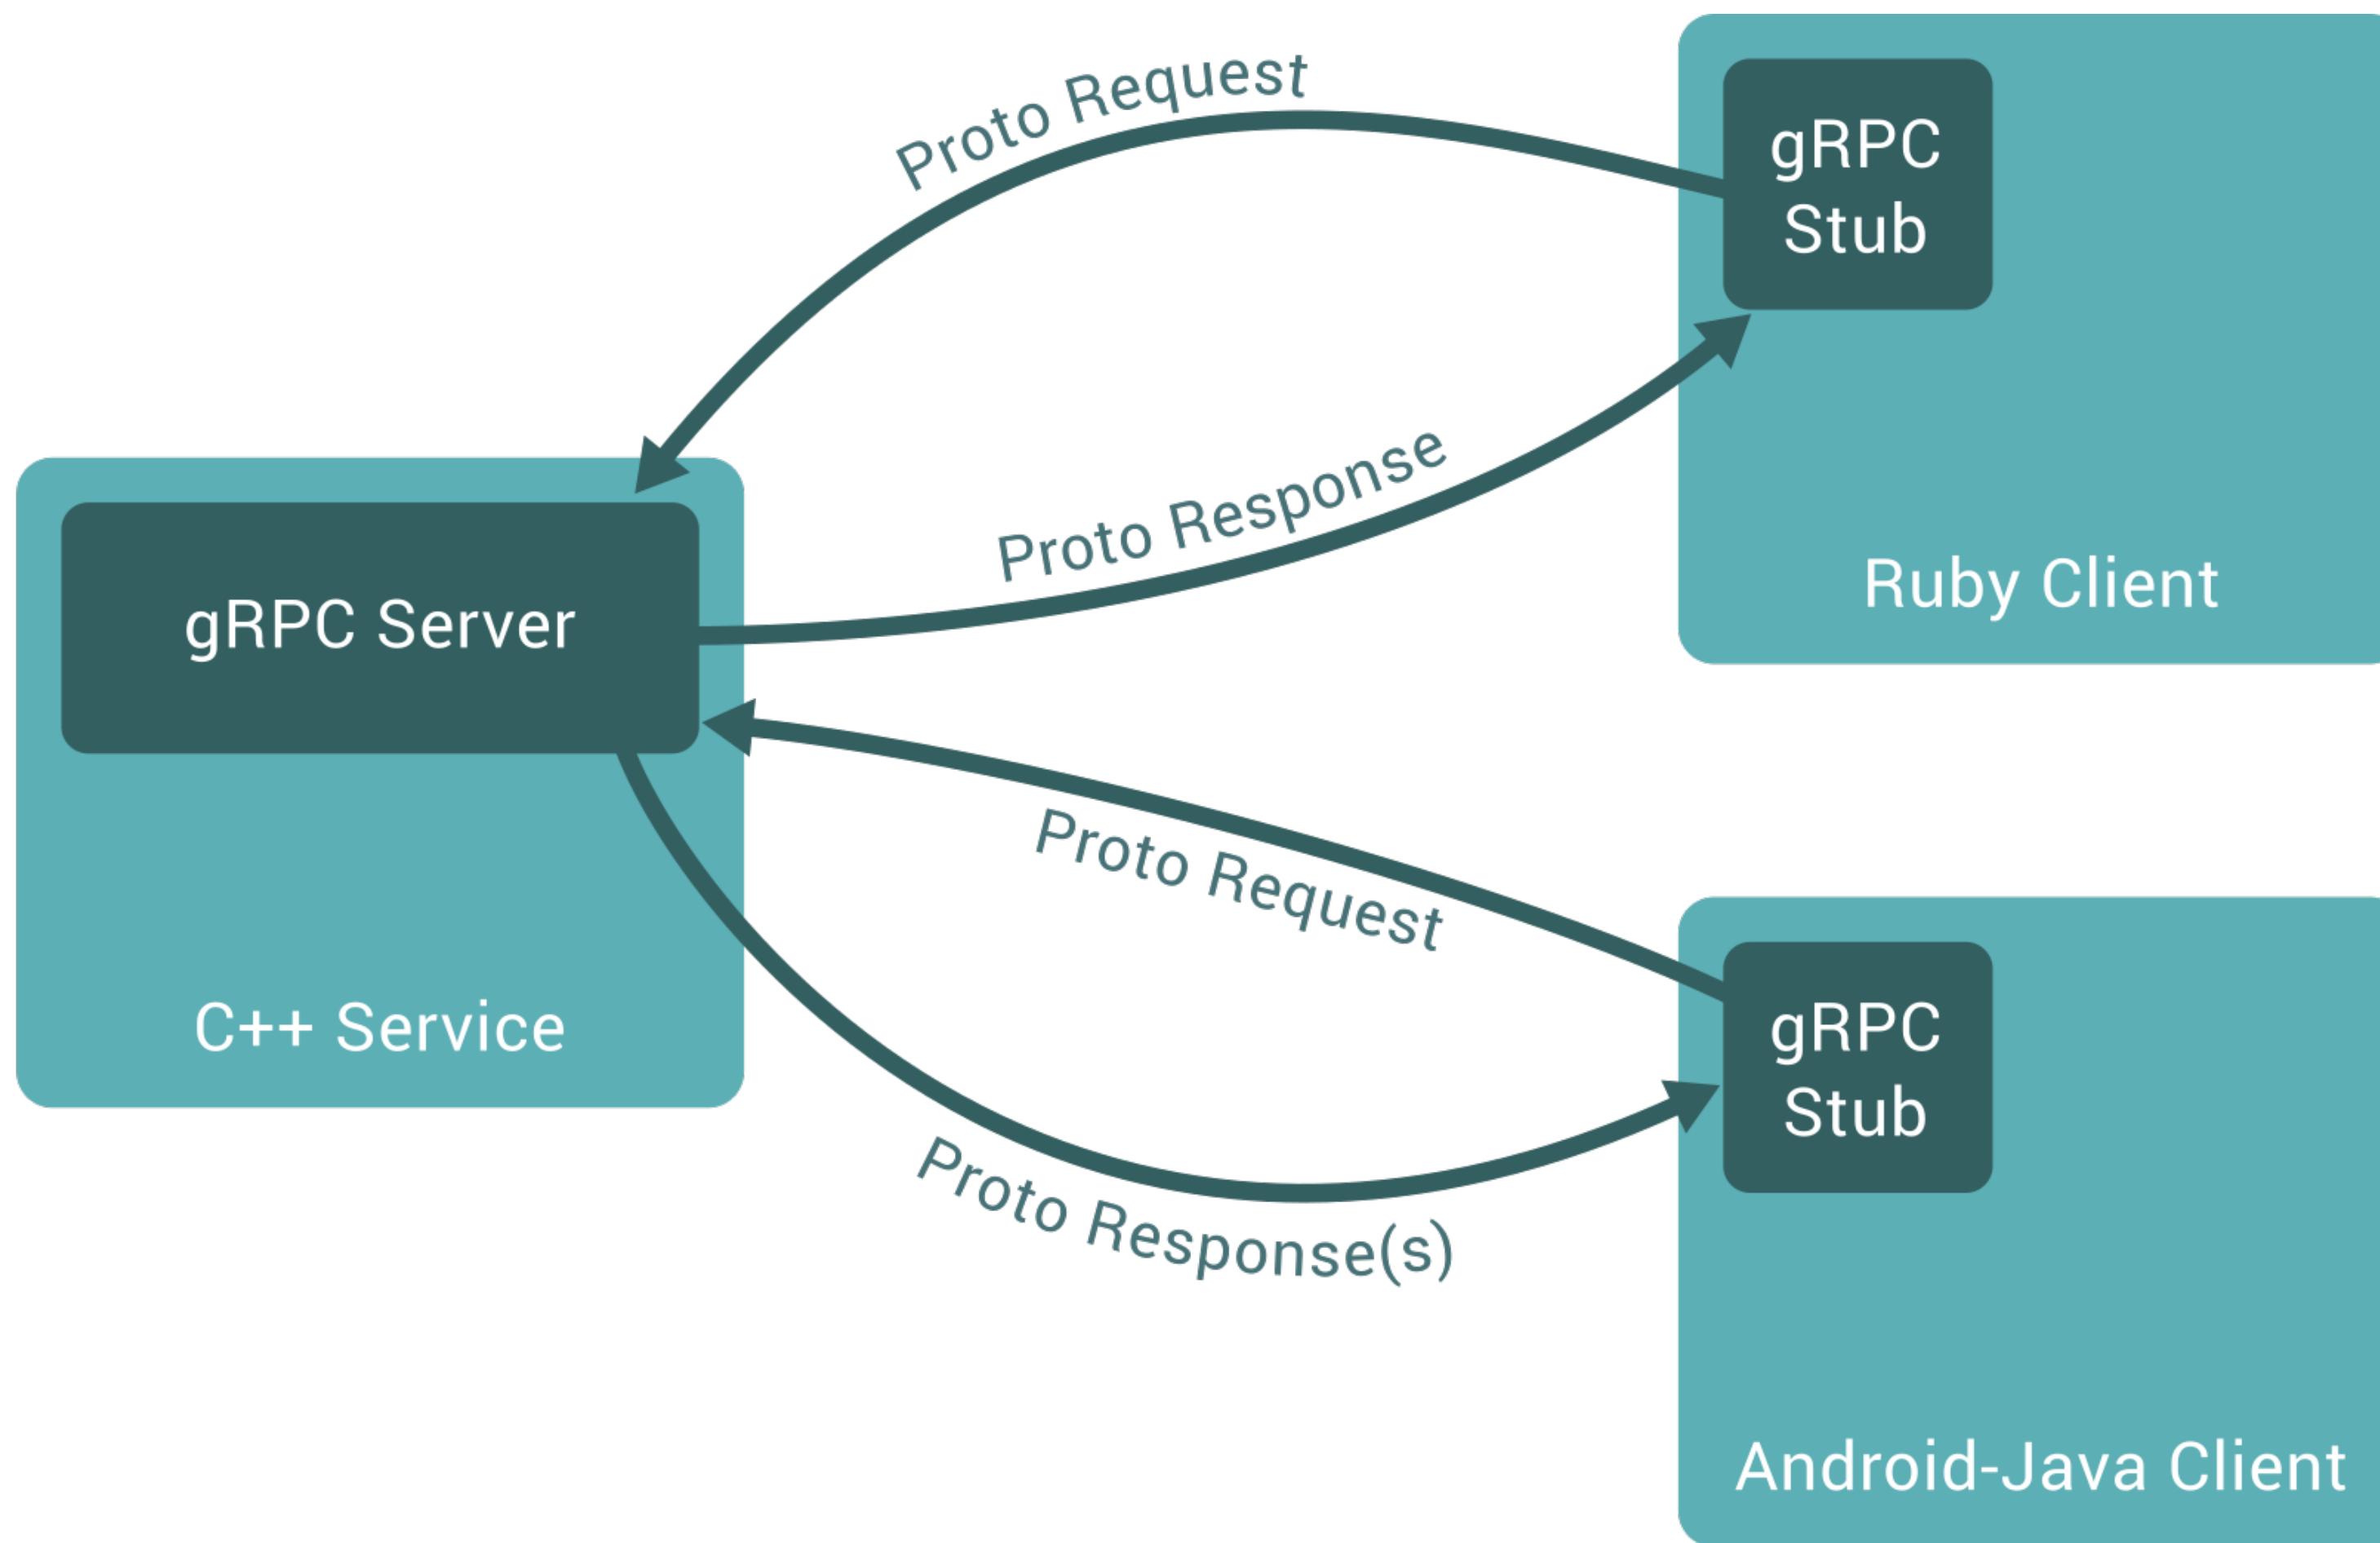
<http://www.grpc.io/blog/principles>



# Coverage & Simplicity

<http://www.grpc.io/blog/principles>

# How does it work



# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition

```
message GetRoutesRequest {  
    Location origin = 1;  
    Location destination = 2;  
}
```

```
message GetRoutesResponse {  
    repeated Route routes = 1;  
}
```

# Generate server code

```
$ python -m grpc_tools.protoc \
--proto_path=../protos \
--python_out=. \
--grpc_python_out=. \
../protos/route_planner.proto
```

# Generate server code

```
$ python -m grpc_tools.protoc \
    --proto_path=../protos \
    --python_out=. \
    --grpc_python_out=. \
    ../protos/route_planner.proto
```

# Generate server code

```
$ python -m grpc_tools.protoc \
    --proto_path=../protos \
    --python_out=. \
    --grpc_python_out=. \
    ../protos/route_planner.proto
```

# Generate server code

```
$ python -m grpc_tools.protoc \
--proto_path=../protos \
--python_out=. \
--grpc_python_out=. \
../protos/route_planner.proto
```

# Generate server code

```
$ python -m grpc_tools.protoc \
--proto_path=../protos \
--python_out=. \
--grpc_python_out=. \
../protos/route_planner.proto
```

# Generate server code

```
$ python -m grpc_tools.protoc \
--proto_path=../protos \
--python_out=. \
--grpc_python_out=. \
../protos/route_planner.proto
```

# Generate server code

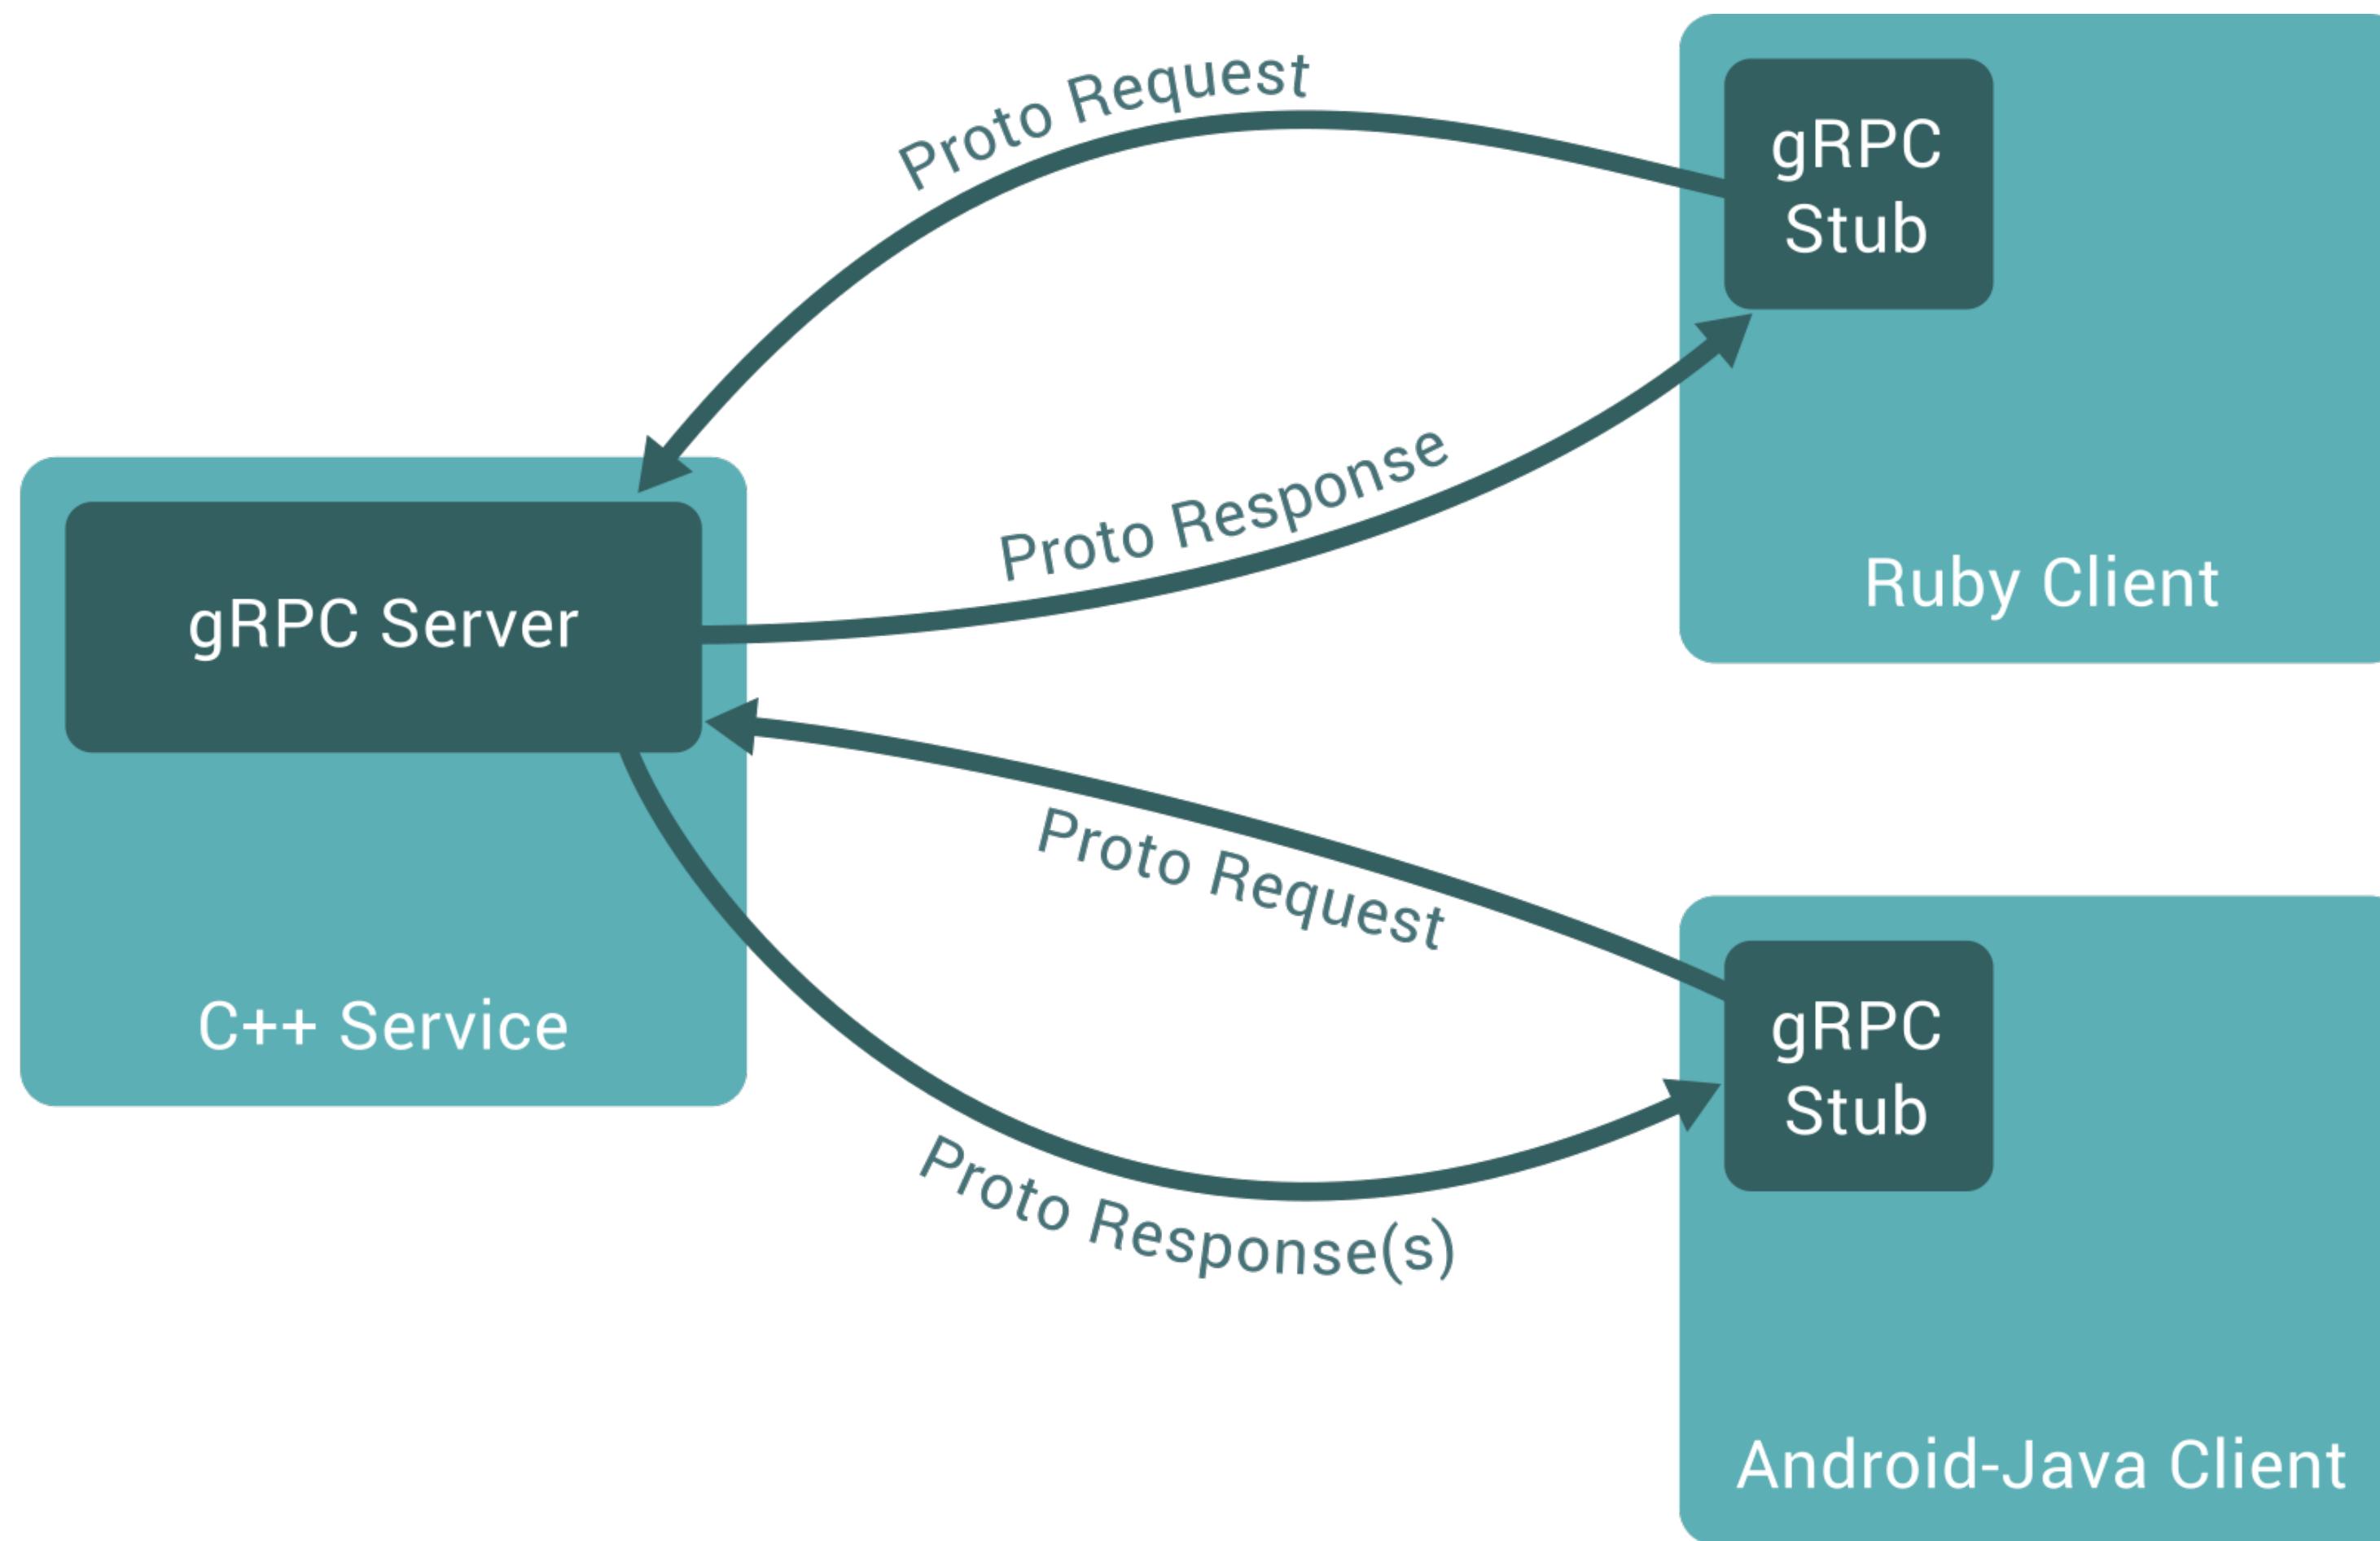
```
$ python -m grpc_tools.protoc \
--proto_path=../protos \
--python_out=. \
--grpc_python_out=. \
../protos/route_planner.proto
```

# Generated code

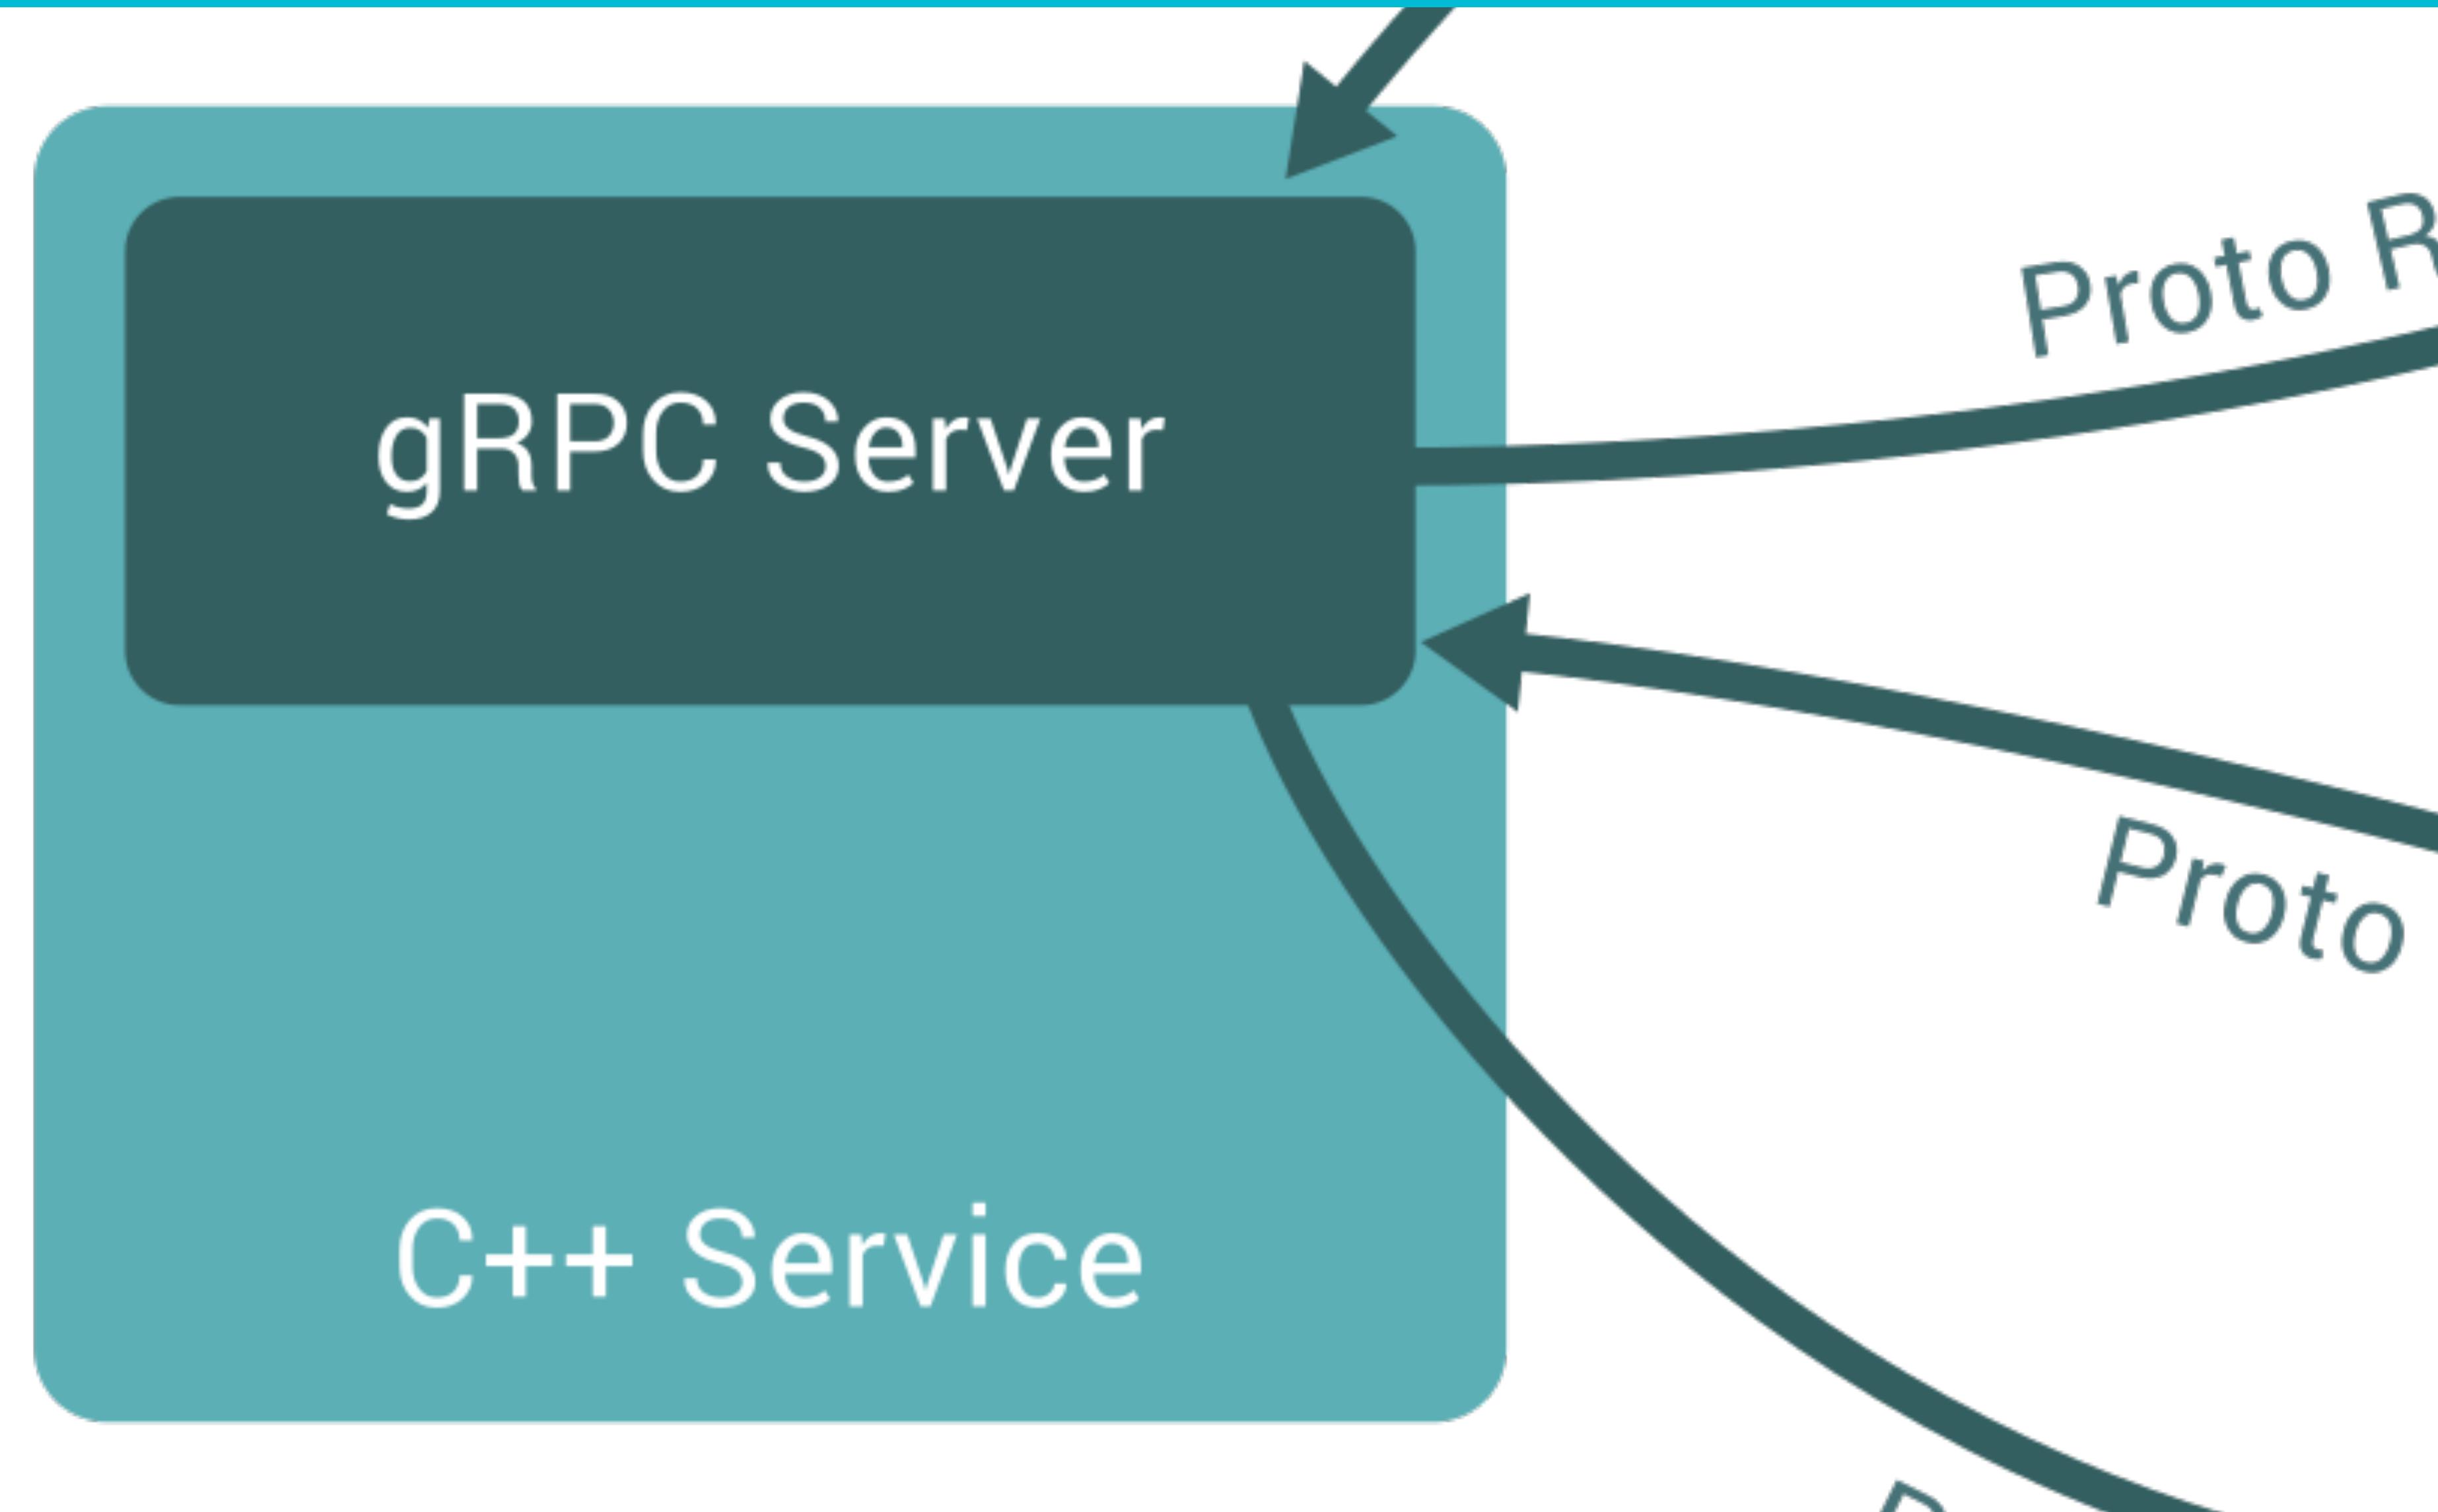
```
$ tree  
.  
└── route_planner_pb2.py  
    └── route_planner_pb2_grpc.py
```

0 directories, 2 files

# Implementing the service



# Implementing the service



# Service code

```
class Servicer(route_planner_pb2_grpc.RoutePlannerServicer):  
    """Service implementation."""  
  
    def GetRoutes(self, request, context):  
        return process_magically_the_request(request)
```

# Service code

```
server = grpc.server(  
    futures.ThreadPoolExecutor(max_workers=10))  
  
route_planner_pb2_grpc.add_RoutePlannerServicer_to_server(  
    Servicer(), server)  
  
server.add_insecure_port('[:]':12345')  
server.start()
```

# Service code

```
server = grpc.server(  
    futures.ThreadPoolExecutor(max_workers=10))  
  
route_planner_pb2_grpc.add_RoutePlannerServicer_to_server(  
    Servicer(), server)  
  
server.add_insecure_port('[:]':12345')  
server.start()
```

# Service code

```
server = grpc.server(  
    futures.ThreadPoolExecutor(max_workers=10))  
  
route_planner_pb2_grpc.add_RoutePlannerServicer_to_server(  
    Servicer(), server)  
  
server.add_insecure_port('[:]':12345')  
server.start()
```

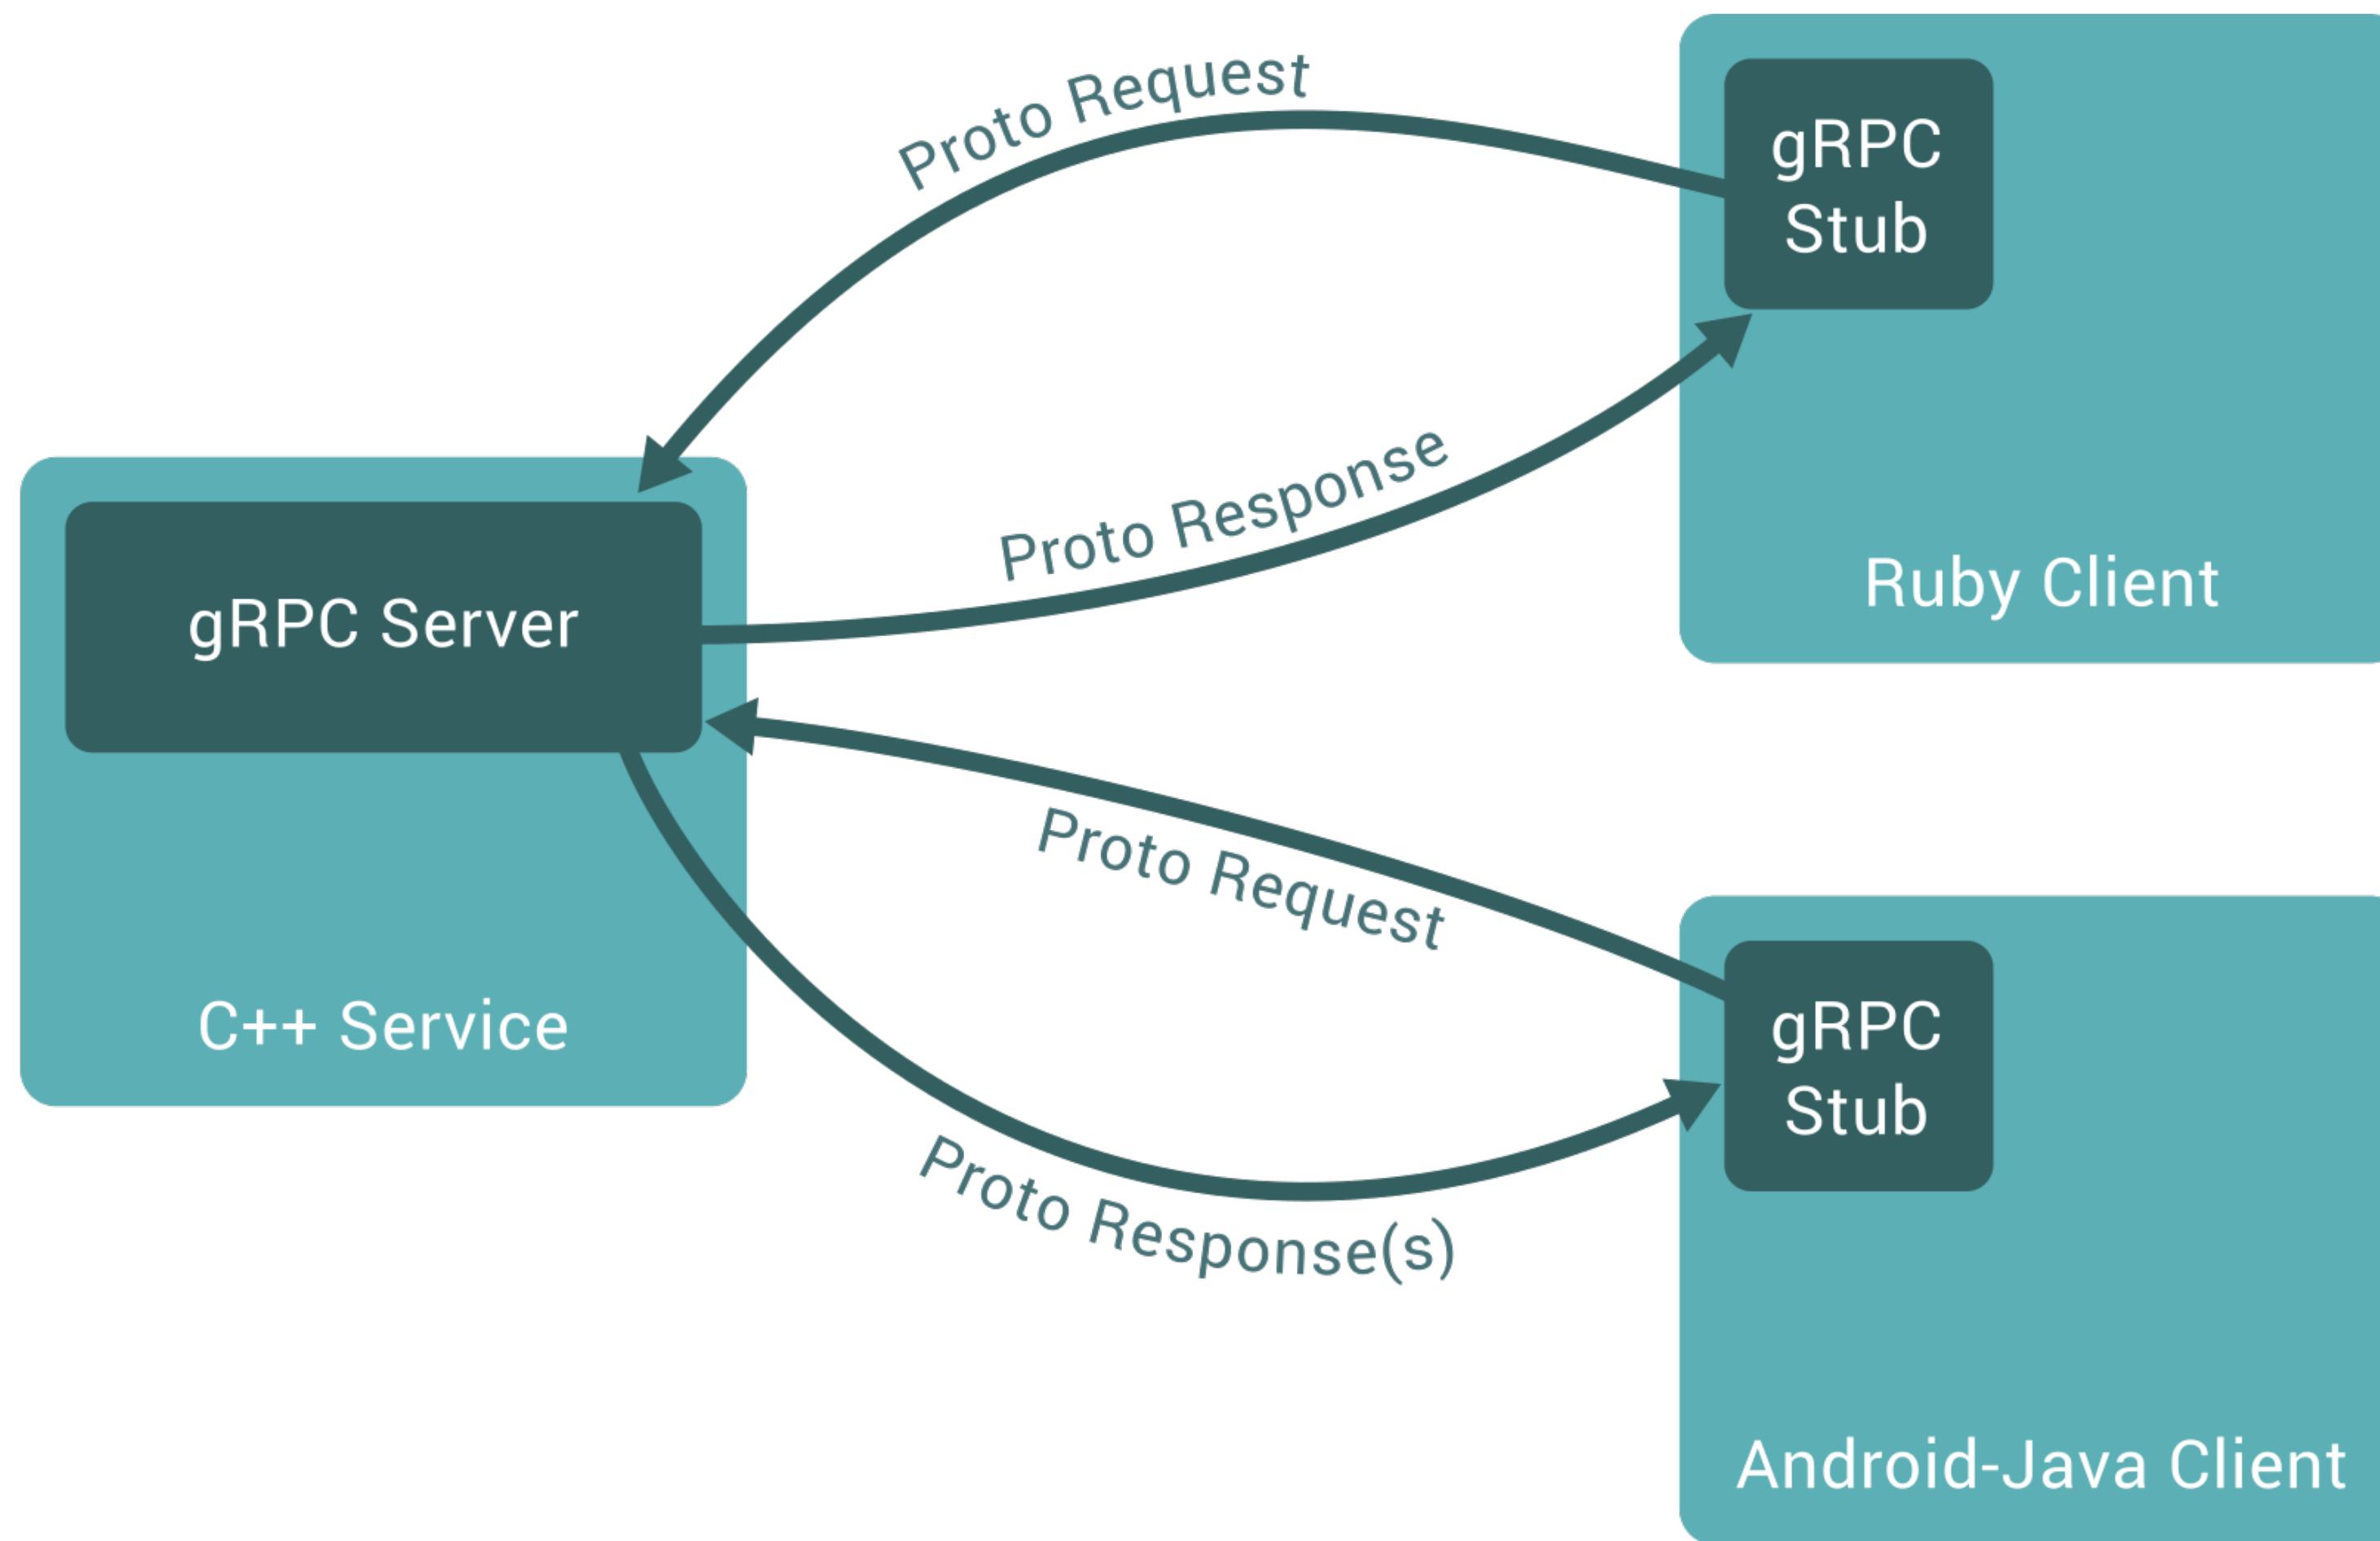
# Service code

```
server = grpc.server(  
    futures.ThreadPoolExecutor(max_workers=10))  
  
route_planner_pb2_grpc.add_RoutePlannerServicer_to_server(  
    Servicer(), server)  
  
server.add_insecure_port('[:]':12345')  
server.start()
```

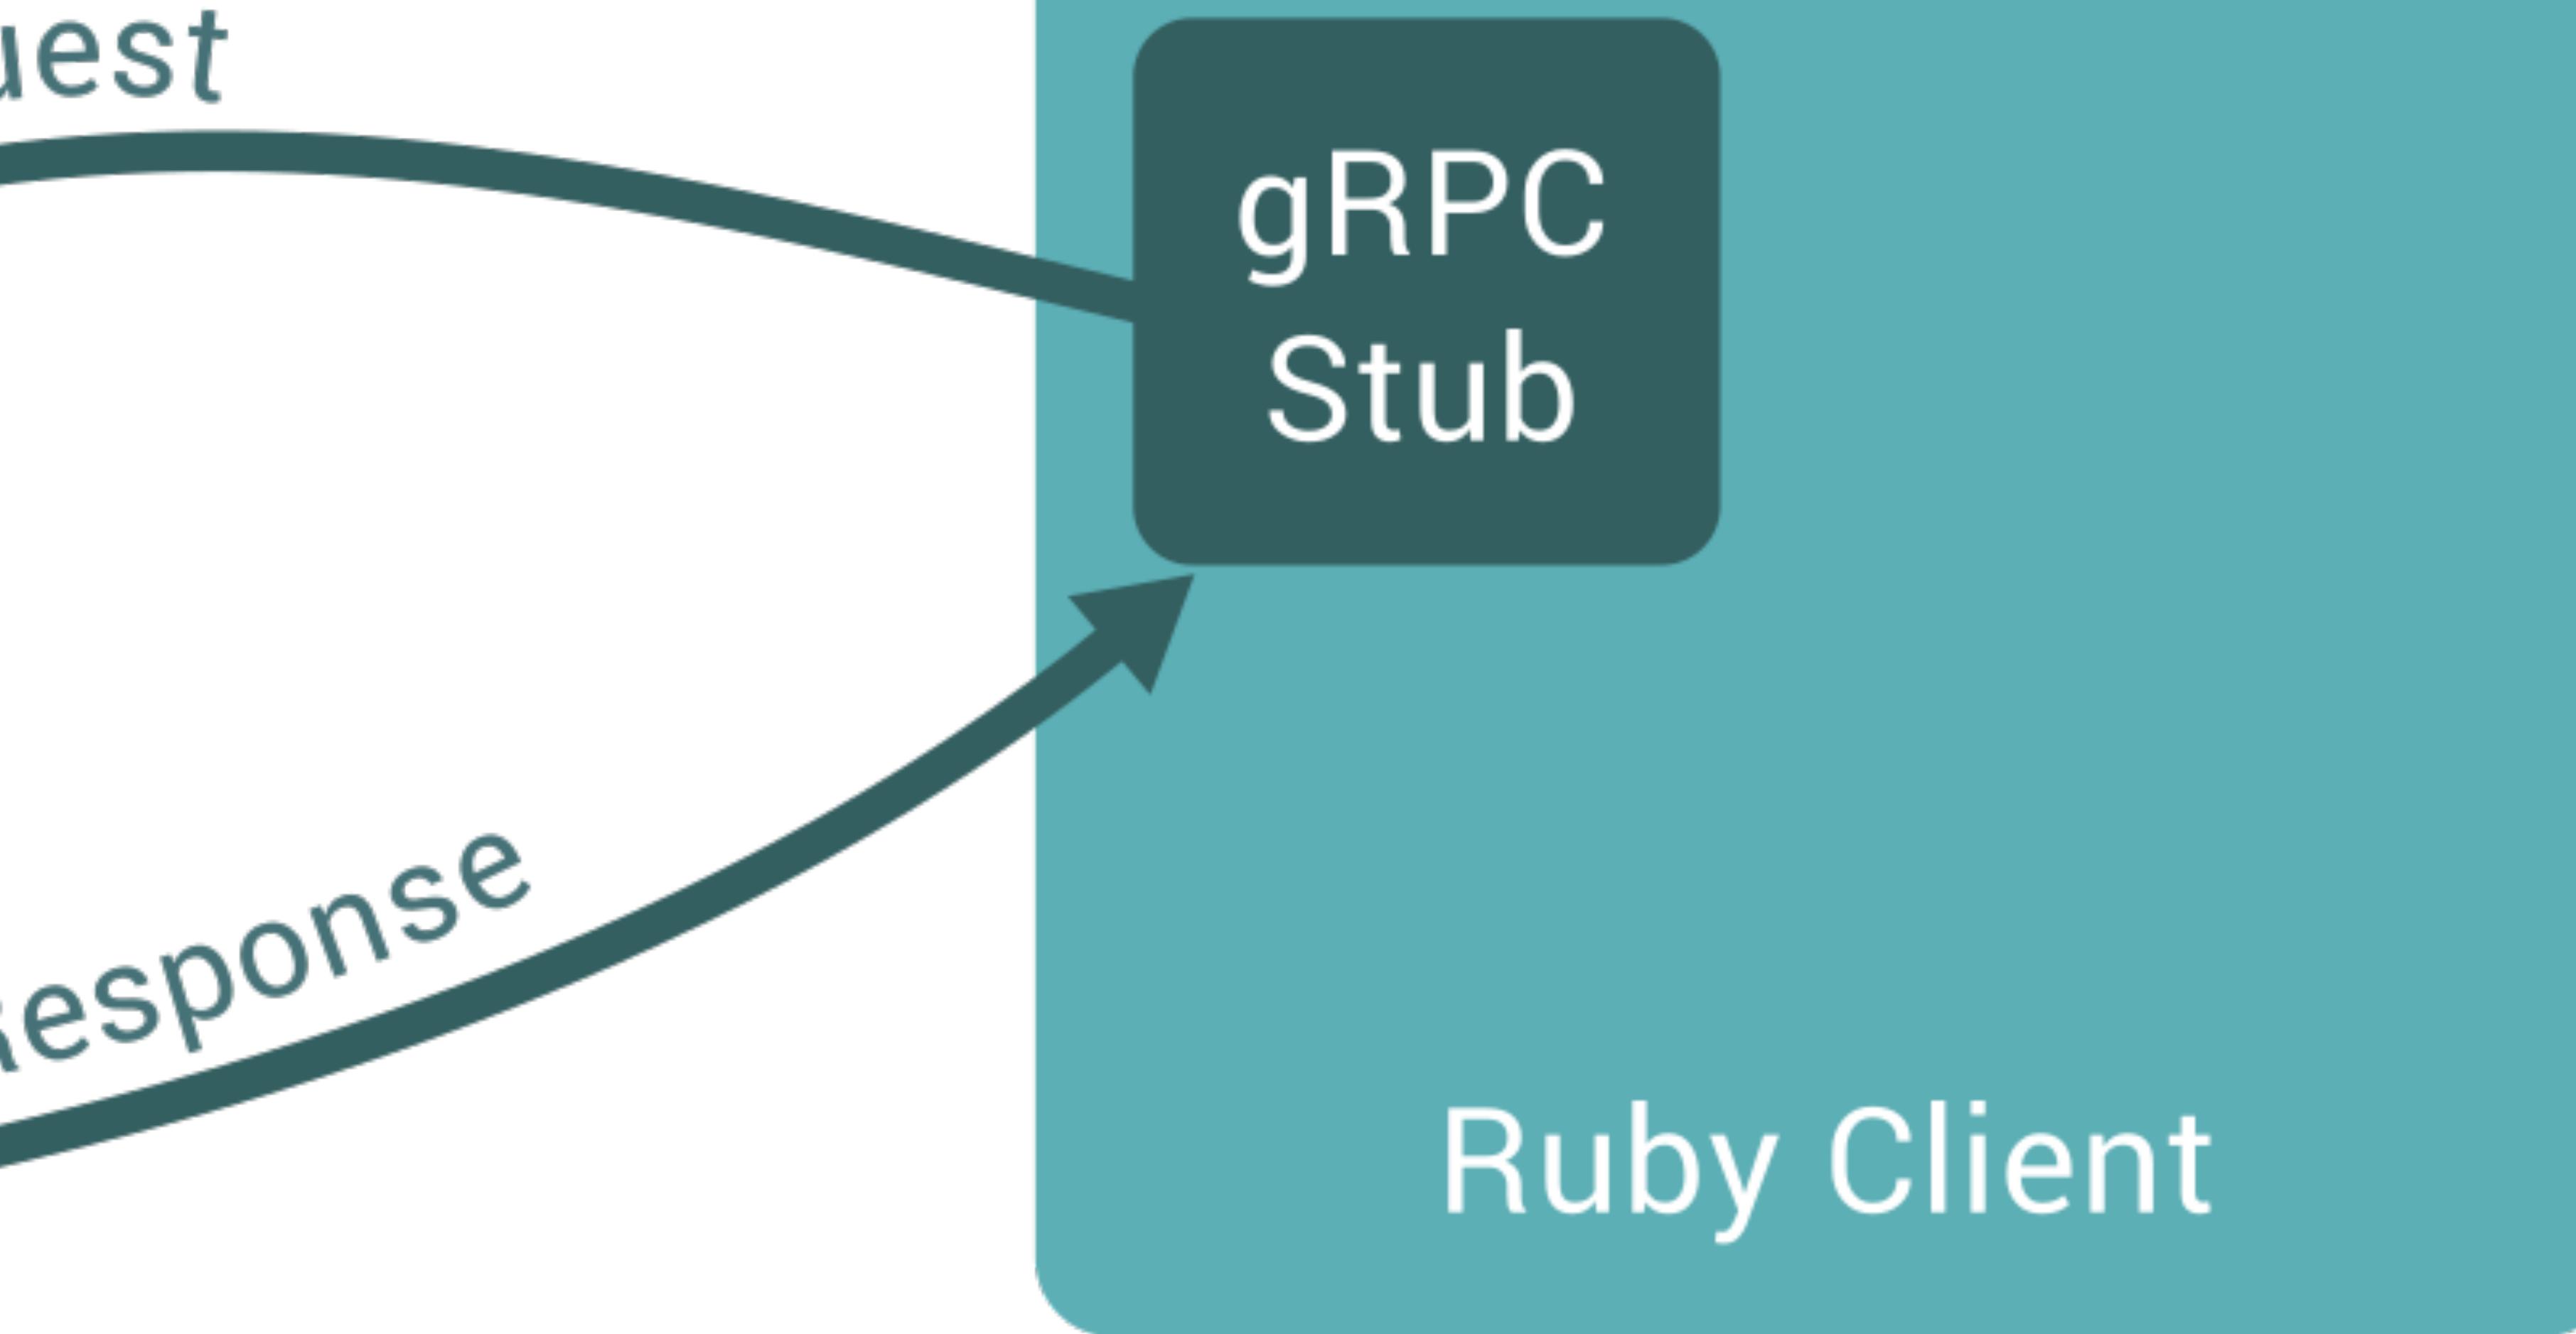
# Service code

```
server = grpc.server(  
    futures.ThreadPoolExecutor(max_workers=10))  
  
route_planner_pb2_grpc.add_RoutePlannerServicer_to_server(  
    Servicer(), server)  
  
server.add_insecure_port('[:]':12345')  
server.start()
```

# Implementing the client



# Implementing the client



# Client code

```
channel = grpc.insecure_channel('localhost:12345')

stub = route_planner_pb2_grpc.RoutePlannerStub(channel)

request = route_planner_pb2.GetRoutesRequest(
    origin=CURRENT_LOCATION,
    destination=DESTINATION_COORDS)

response = stub.GetRoutes(request)
```

# Client code

```
channel = grpc.insecure_channel('localhost:12345')

stub = route_planner_pb2_grpc.RoutePlannerStub(channel)

request = route_planner_pb2.GetRoutesRequest(
    origin=CURRENT_LOCATION,
    destination=DESTINATION_COORDS)

response = stub.GetRoutes(request)
```

# Client code

```
channel = grpc.insecure_channel('localhost:12345')
```

```
stub = route_planner_pb2_grpc.RoutePlannerStub(channel)
```

```
request = route_planner_pb2.GetRoutesRequest(  
    origin=CURRENT_LOCATION,  
    destination=DESTINATION_COORDS)
```

```
response = stub.GetRoutes(request)
```

# Client code

```
channel = grpc.insecure_channel('localhost:12345')

stub = route_planner_pb2_grpc.RoutePlannerStub(channel)

request = route_planner_pb2.GetRoutesRequest(
    origin=CURRENT_LOCATION,
    destination=DESTINATION_COORDS)

response = stub.GetRoutes(request)
```

# Client code

```
channel = grpc.insecure_channel('localhost:12345')

stub = route_planner_pb2_grpc.RoutePlannerStub(channel)

request = route_planner_pb2.GetRoutesRequest(
    origin=CURRENT_LOCATION,
    destination=DESTINATION_COORDS)

response = stub.GetRoutes(request)
```

# Client code

```
channel = grpc.insecure_channel('localhost:12345')

stub = route_planner_pb2_grpc.RoutePlannerStub(channel)

request = route_planner_pb2.GetRoutesRequest(
    origin=CURRENT_LOCATION,
    destination=DESTINATION_COORDS)

response = stub.GetRoutes(request)
```

# Client code (async)

```
response_future = stub.GetRoutes.future(request)  
  
response_future.result()
```

# grpc\_cli

```
$ grpc_cli call localhost:12345 \
  RoutePlanner.GetRoutes \
  <<- PROTO
    origin: <long: 0.0 lat: 0.0>
    destination: <long: 1.1 lat: 1.1>
PROTO
```

# grpc\_cli

```
$ grpc_cli call localhost:12345 \
  RoutePlanner.GetRoutes \
  <<- PROTO
    origin: <long: 0.0 lat: 0.0>
    destination: <long: 1.1 lat: 1.1>
PROTO
```

# grpc\_cli

```
$ grpc_cli call localhost:12345 \
    RoutePlanner.GetRoutes \
    <<- PROTO
        origin: <long: 0.0 lat: 0.0>
        destination: <long: 1.1 lat: 1.1>
PROTO
```

# grpc\_cli

```
$ grpc_cli call localhost:12345 \
  RoutePlanner.GetRoutes \
<<- PROTO
  origin: <long: 0.0 lat: 0.0>
  destination: <long: 1.1 lat: 1.1>
PROTO
```

# grpc\_cli

```
$ grpc_cli call localhost:12345 \
  RoutePlanner.GetRoutes \
  <<- PROTO
    origin: <long: 0.0 lat: 0.0>
    destination: <long: 1.1 lat: 1.1>
PROTO
```

# grpc\_cli

Rpc succeeded with OK status

Response:

routes: <...>

routes: <...>

routes: <...>

routes: <...>

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition - response streaming

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (stream GetRoutesResponse) {}  
}
```

# Service definition

```
service RoutePlanner {  
    rpc GetRoutes (GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Service definition - request streaming

```
service RoutePlanner {  
    rpc GetRoutes (stream GetRoutesRequest)  
        returns (GetRoutesResponse) {}  
}
```

# Request streaming? Response streaming?

# Request streaming? Response streaming?



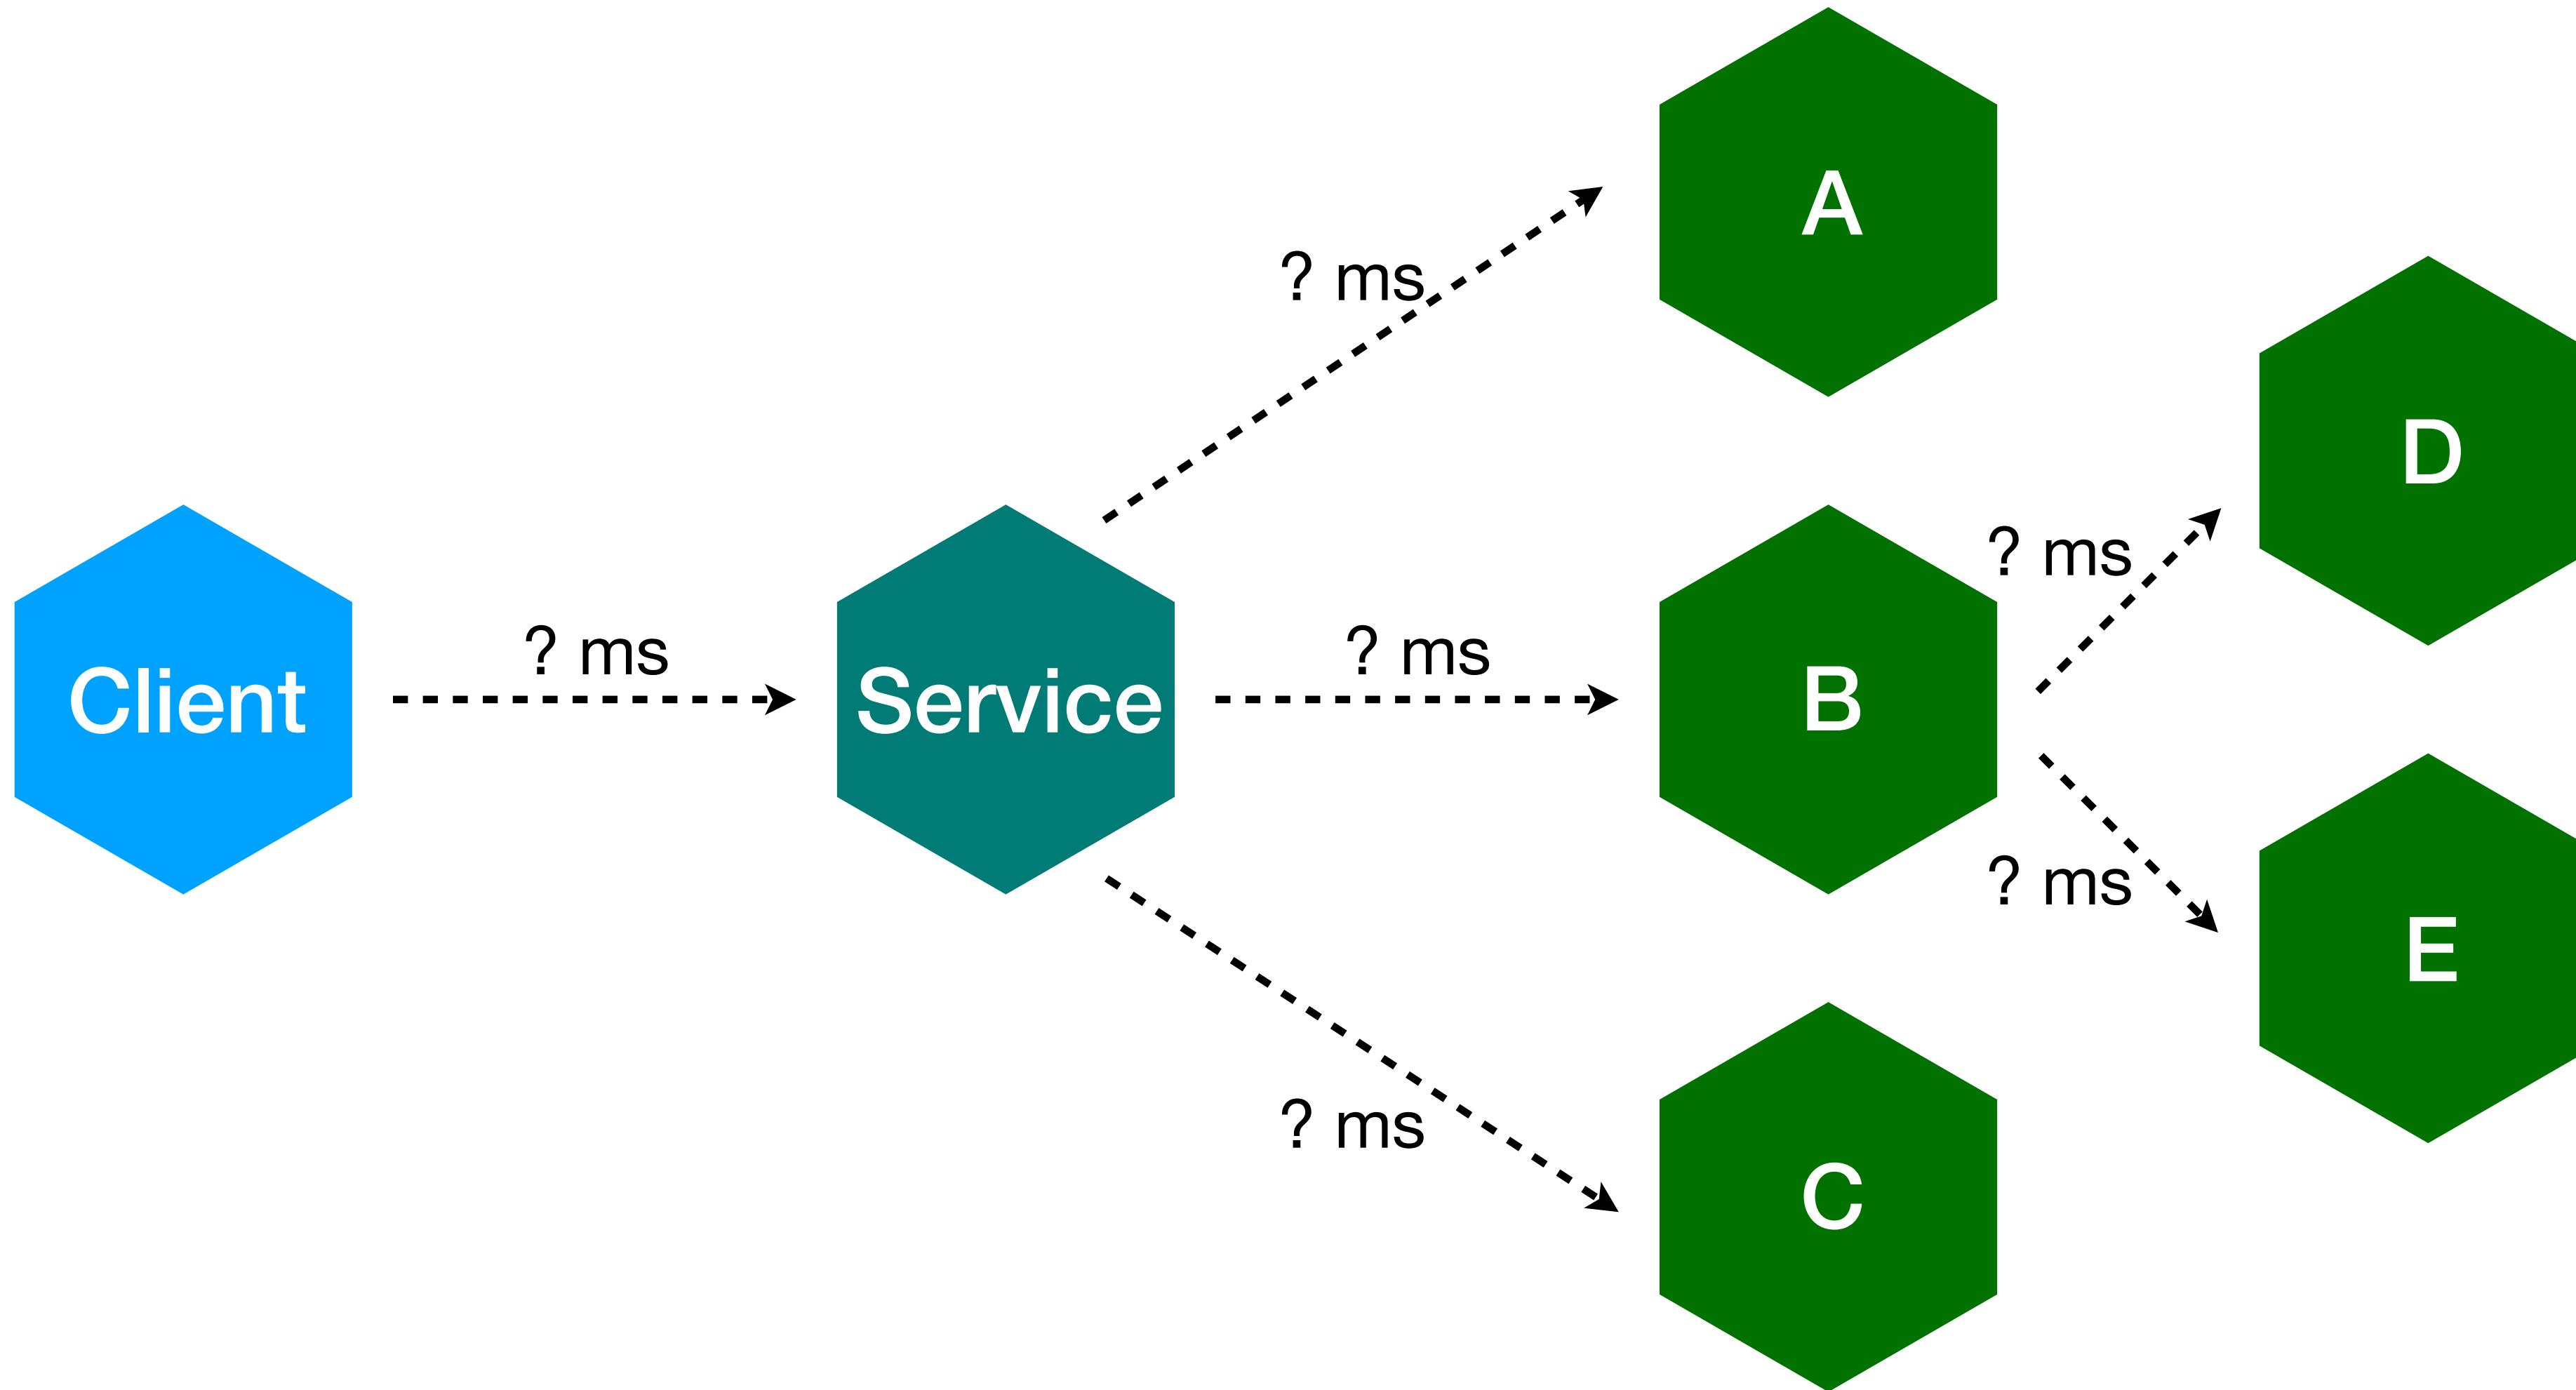
# Service definition - bidirectional streaming

```
service RoutePlanner {  
    rpc GetRoutes (stream GetRoutesRequest)  
        returns (stream GetRoutesResponse) {}  
}
```

# Keep in mind...

Things will go wrong

# Timeouts



# Uniform timeout

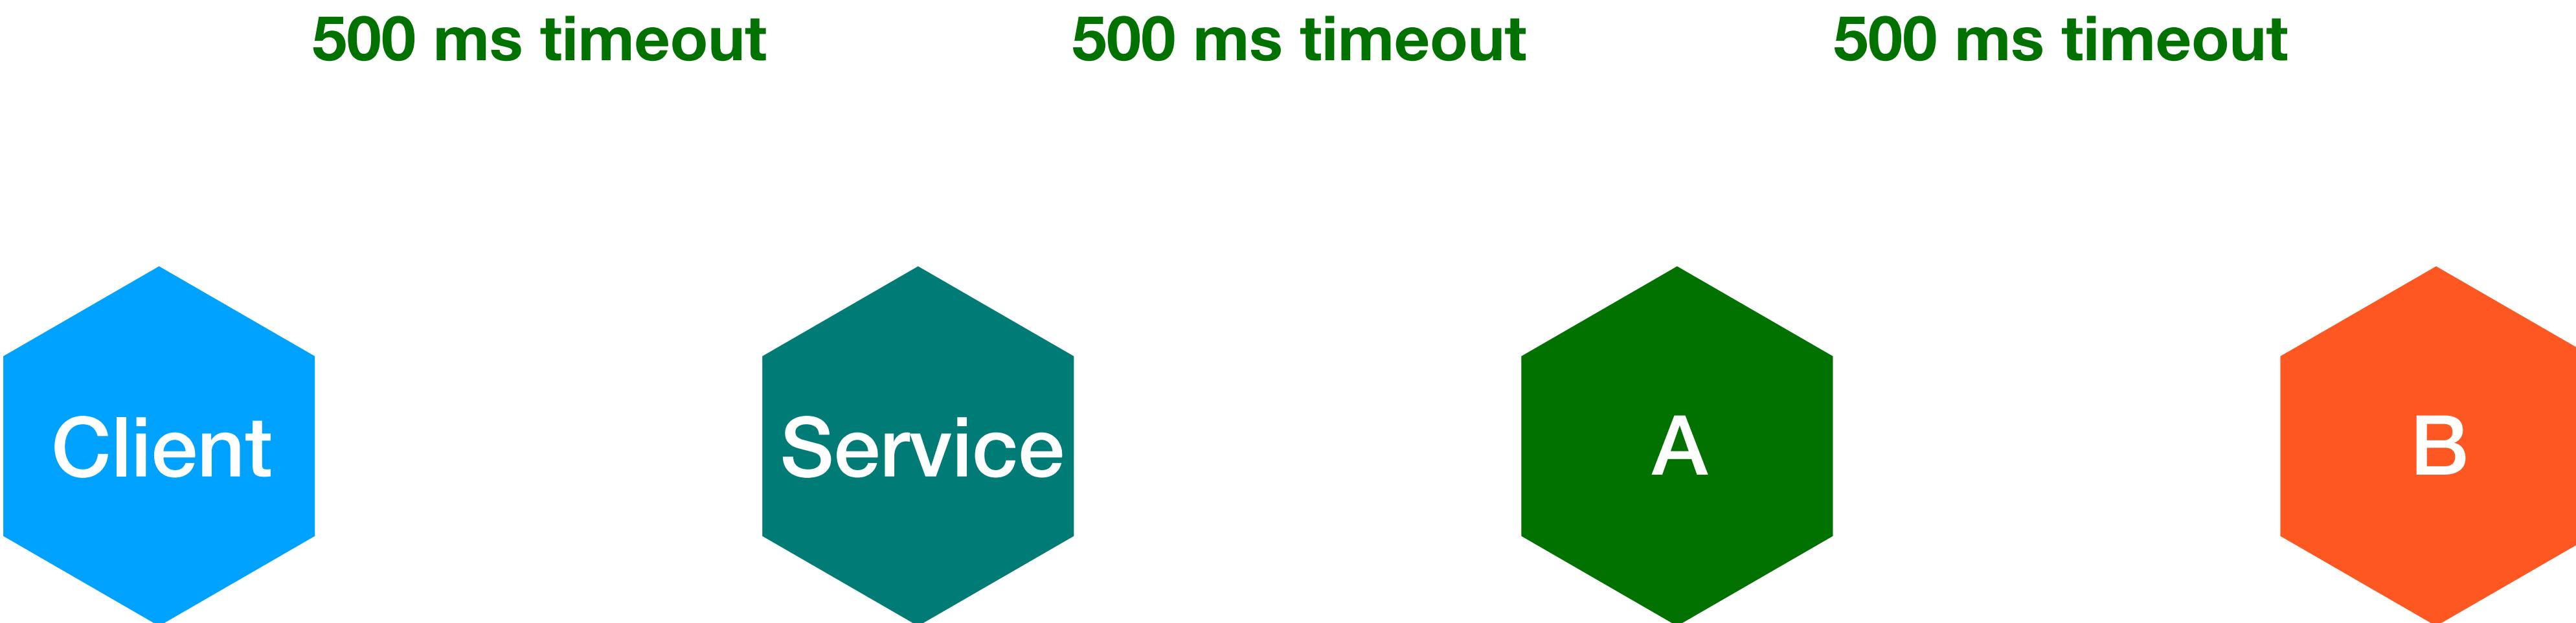
Client

Service

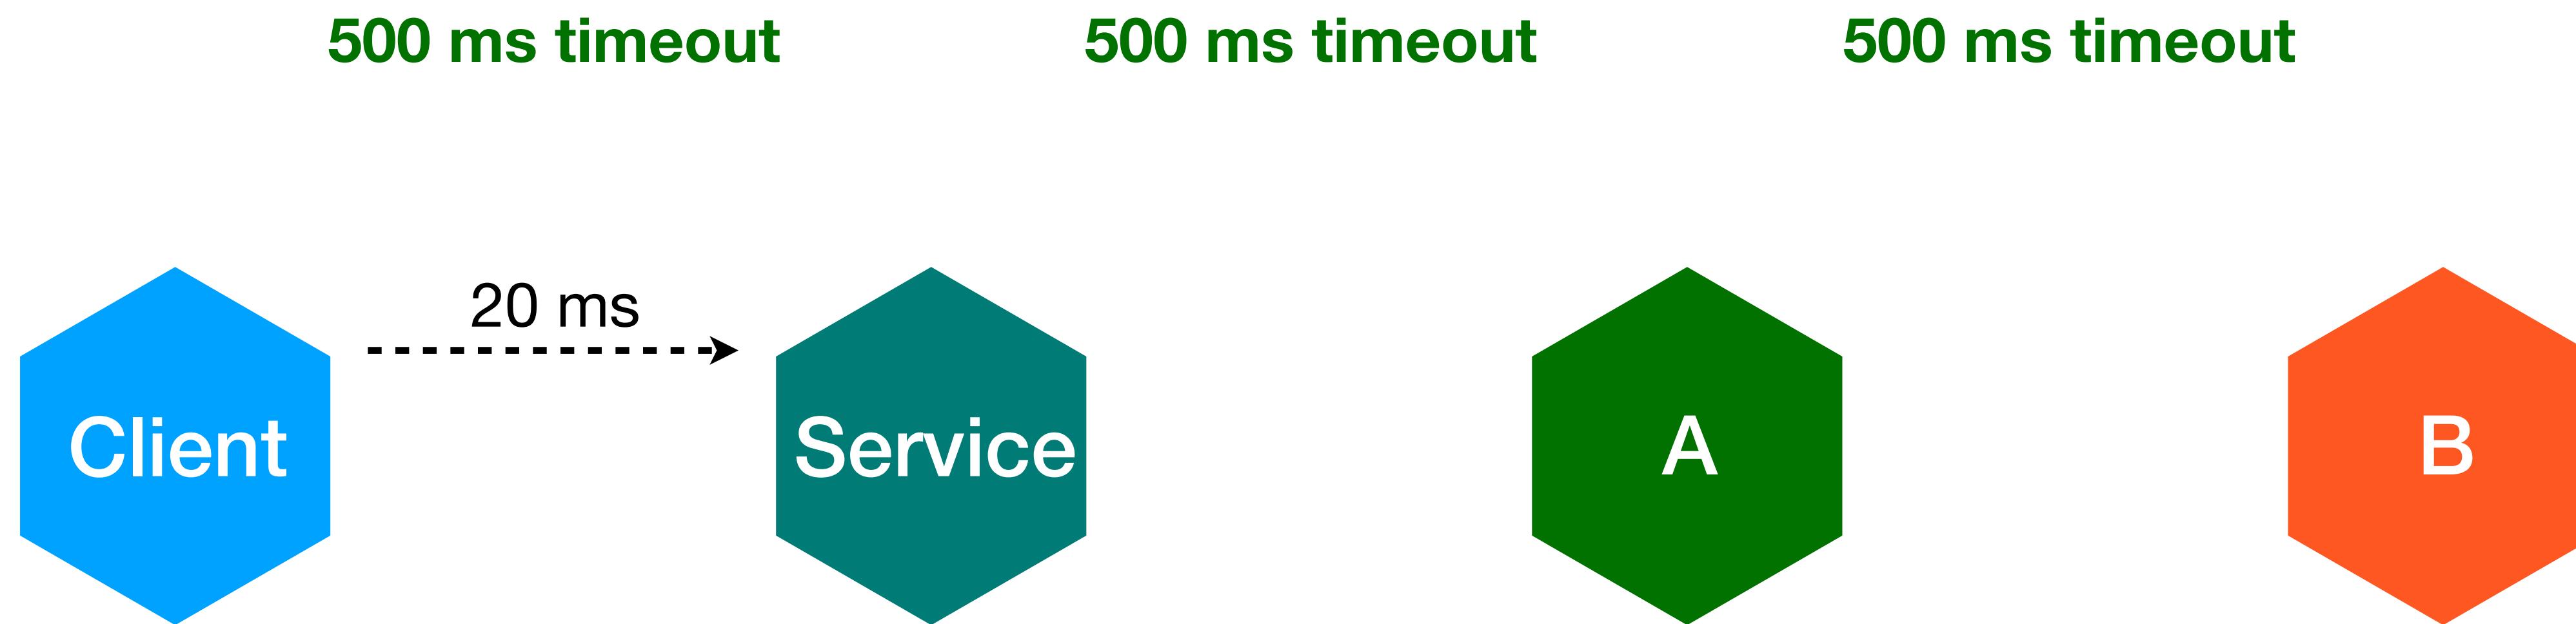
A

B

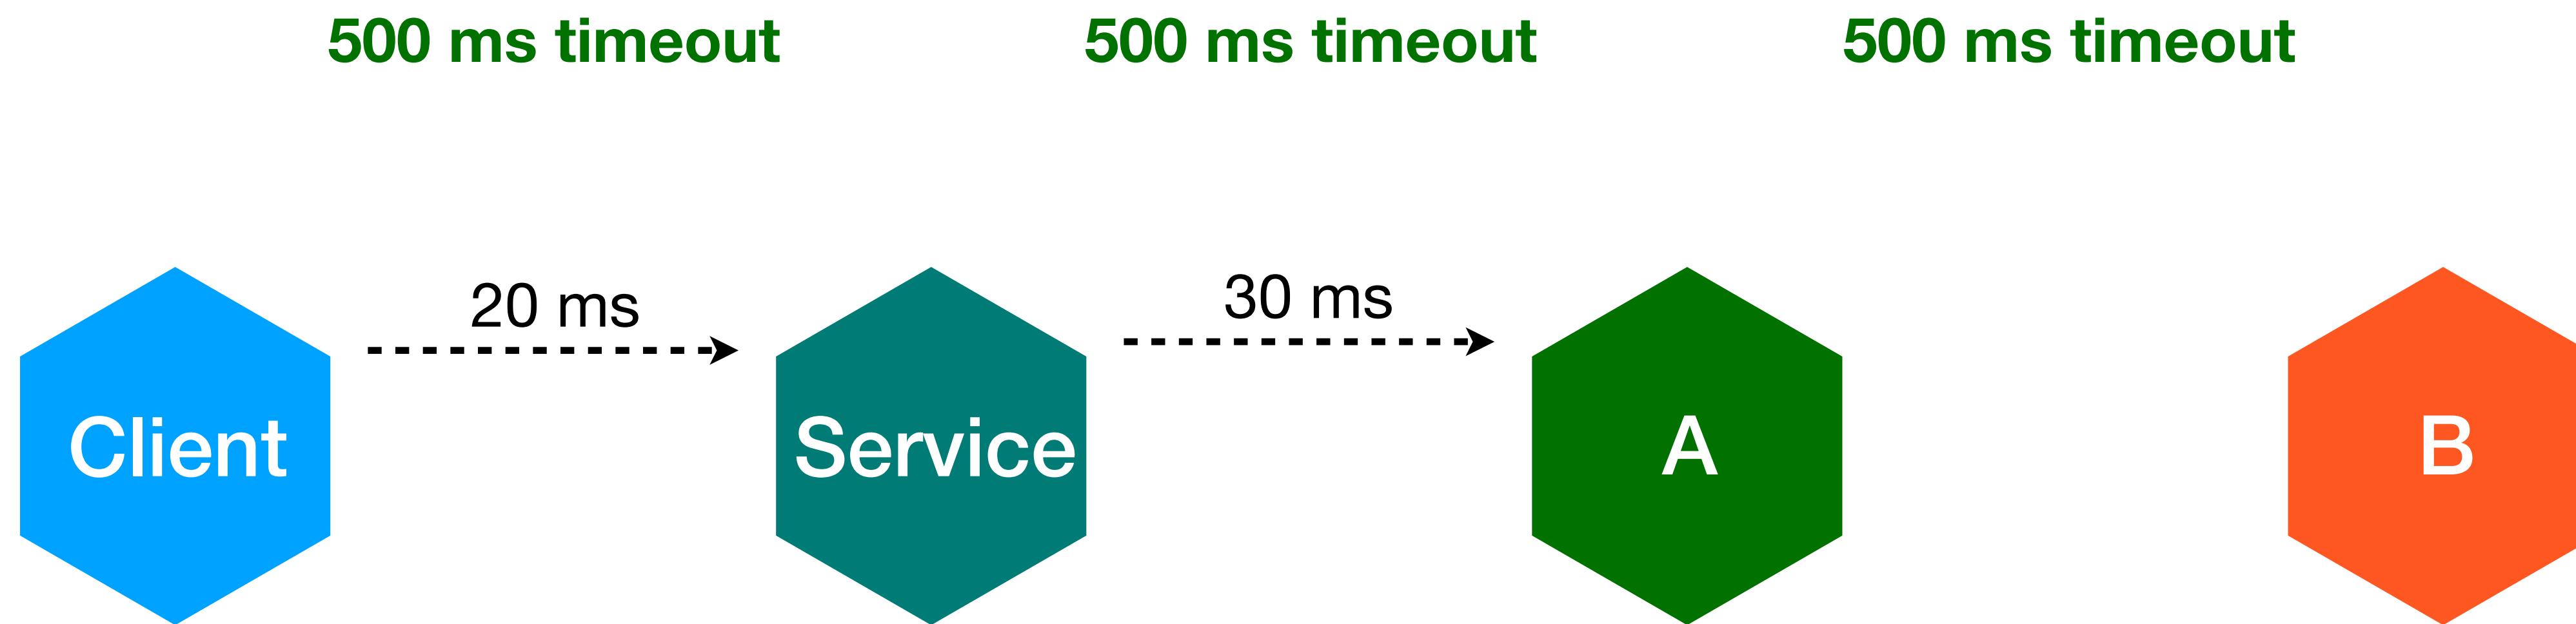
# Uniform timeout



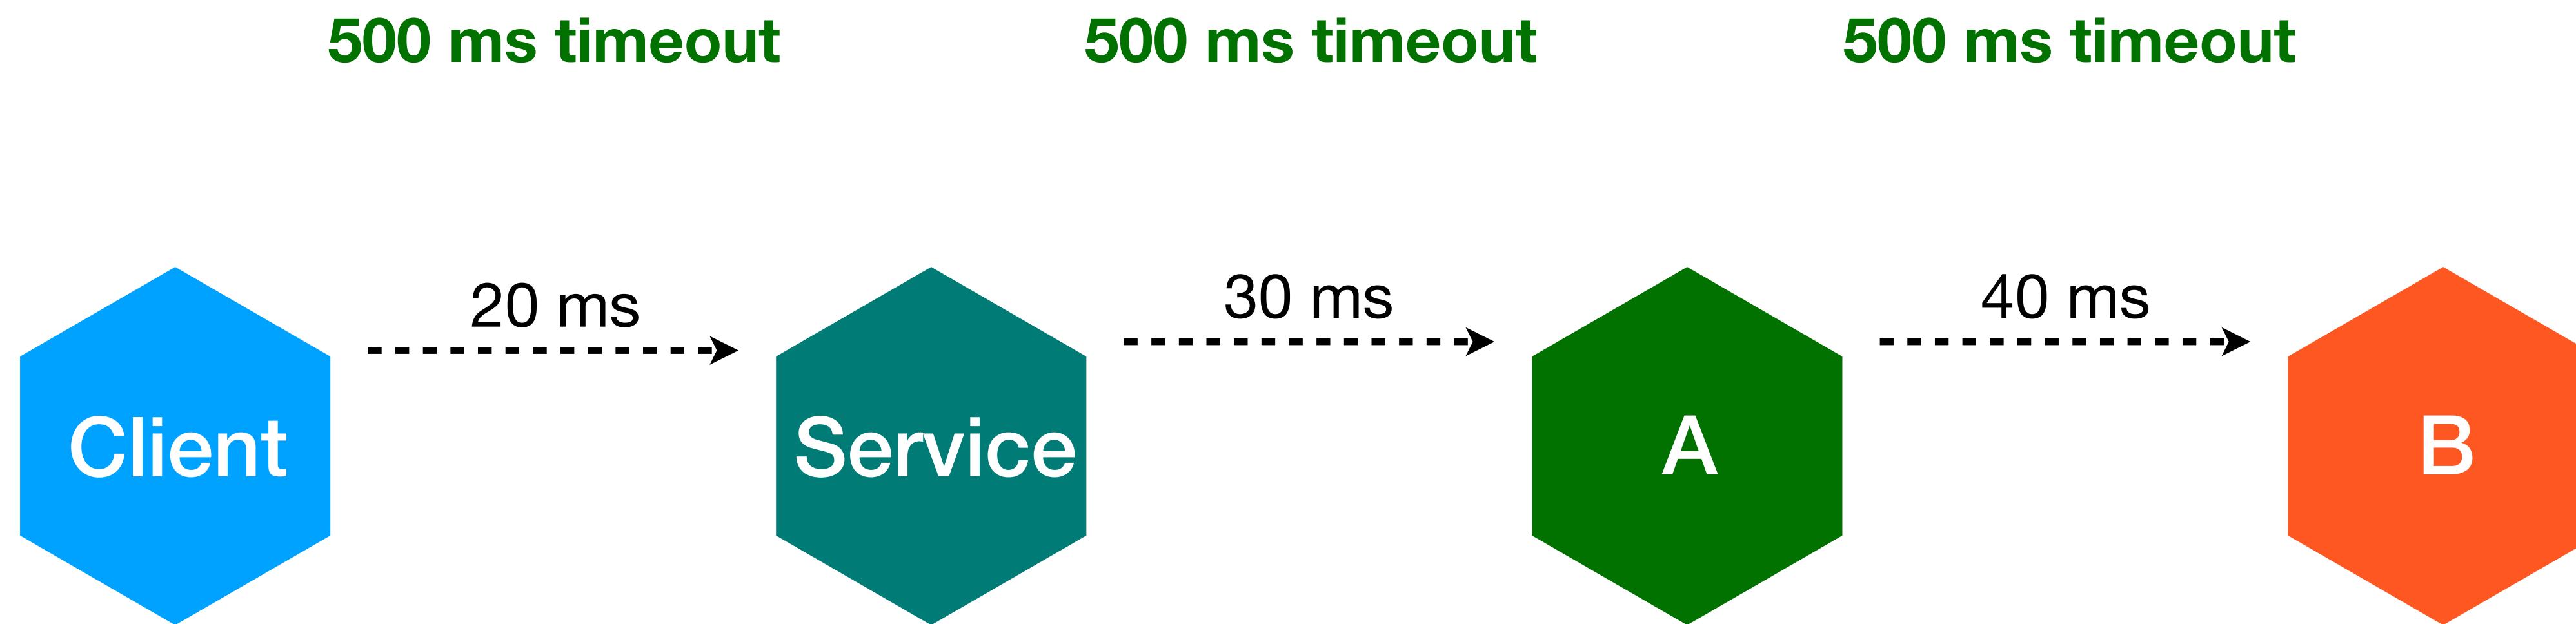
# Uniform timeout



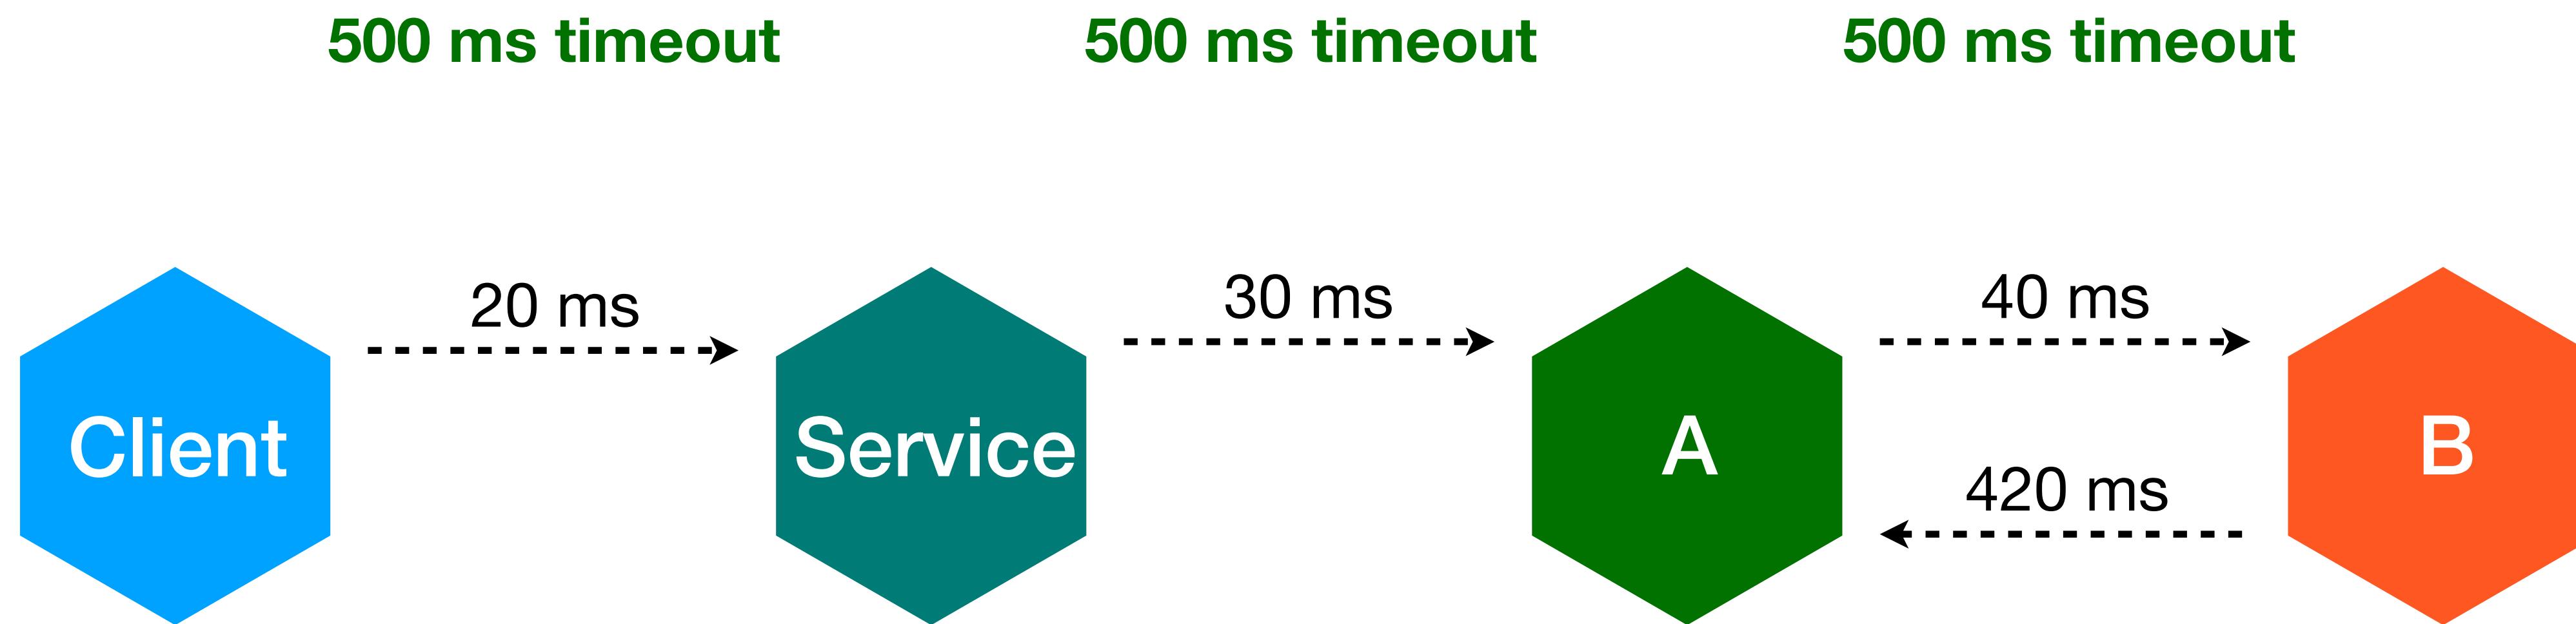
# Uniform timeout



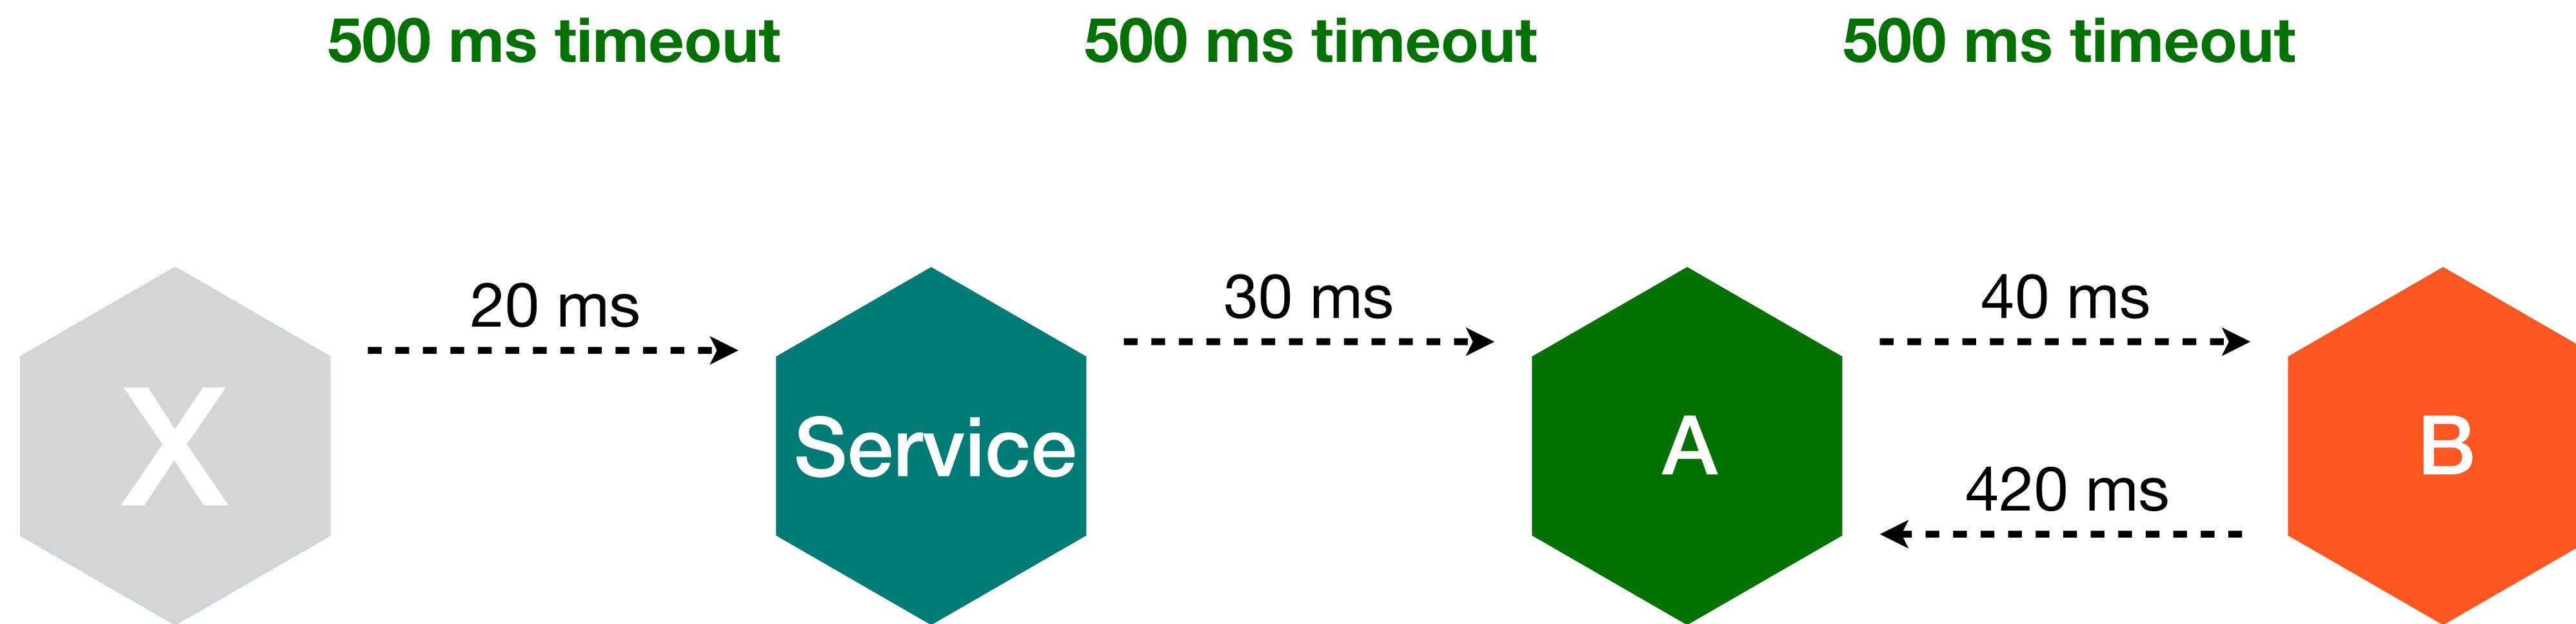
# Uniform timeout



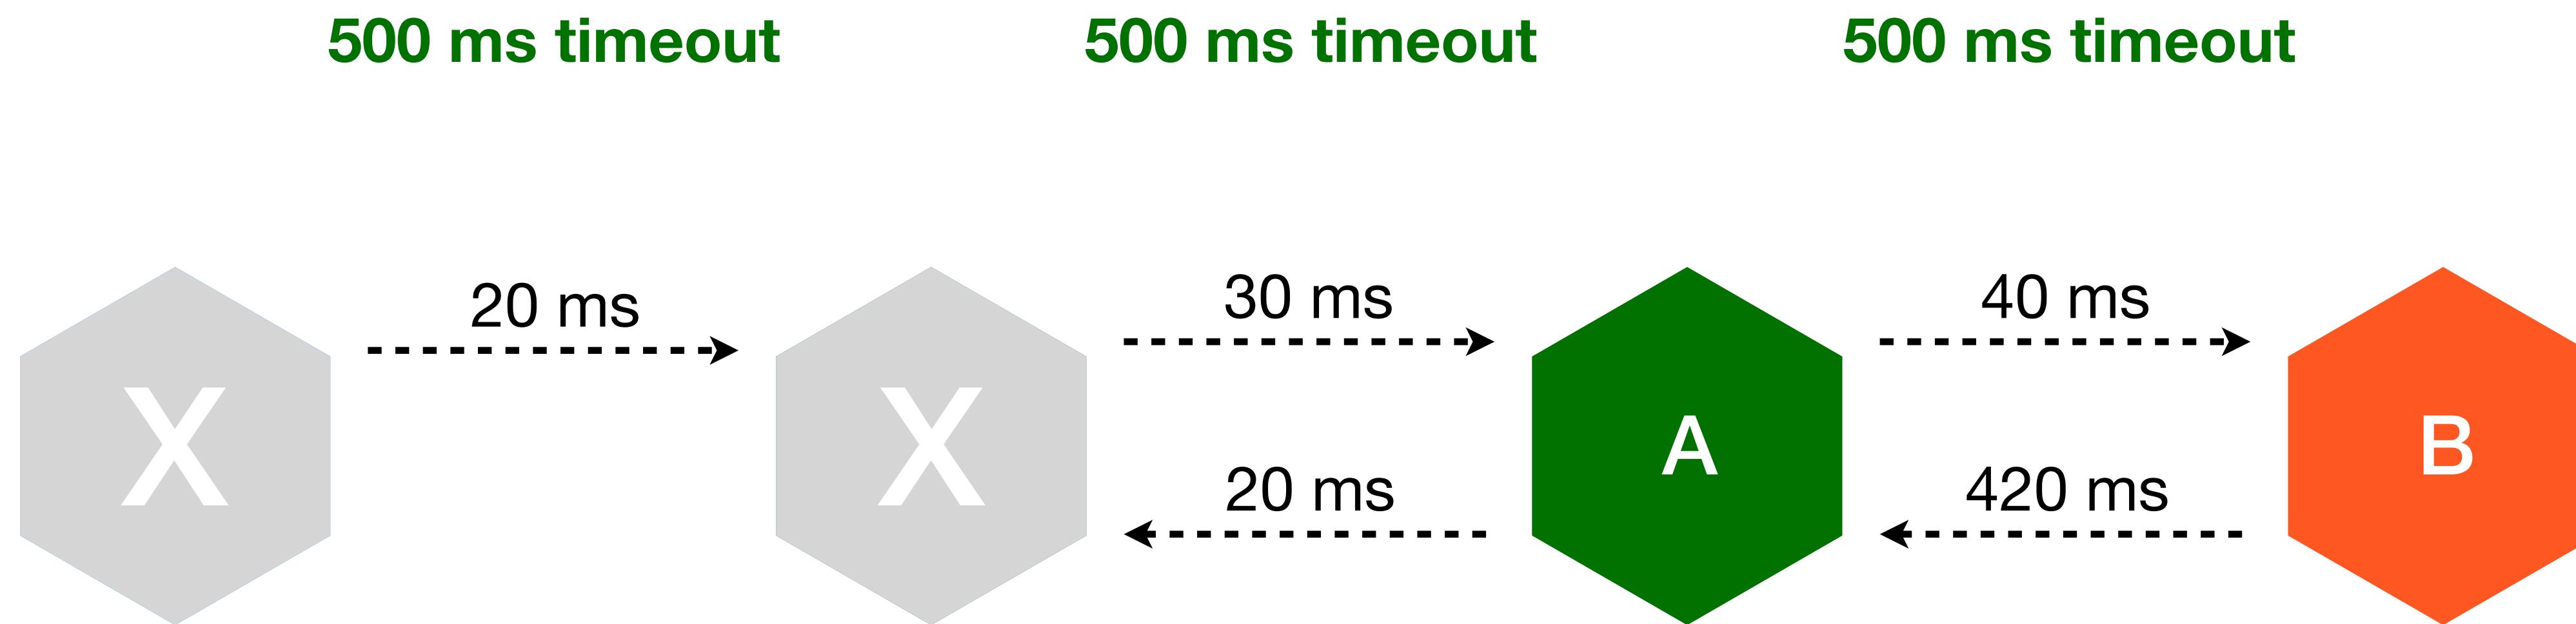
# Uniform timeout



# Uniform timeout



# Uniform timeout



# Fine-tuned timeout

Client

Service

A

B

# Fine-tuned timeout

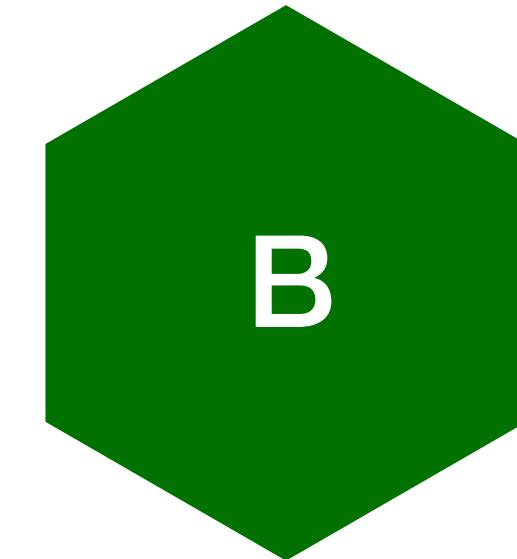
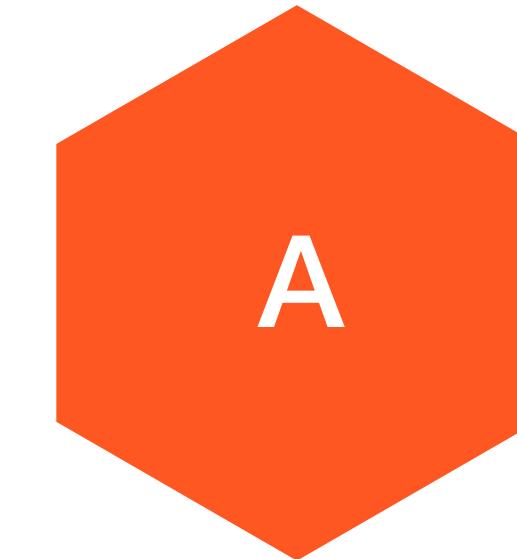
**300 ms timeout**



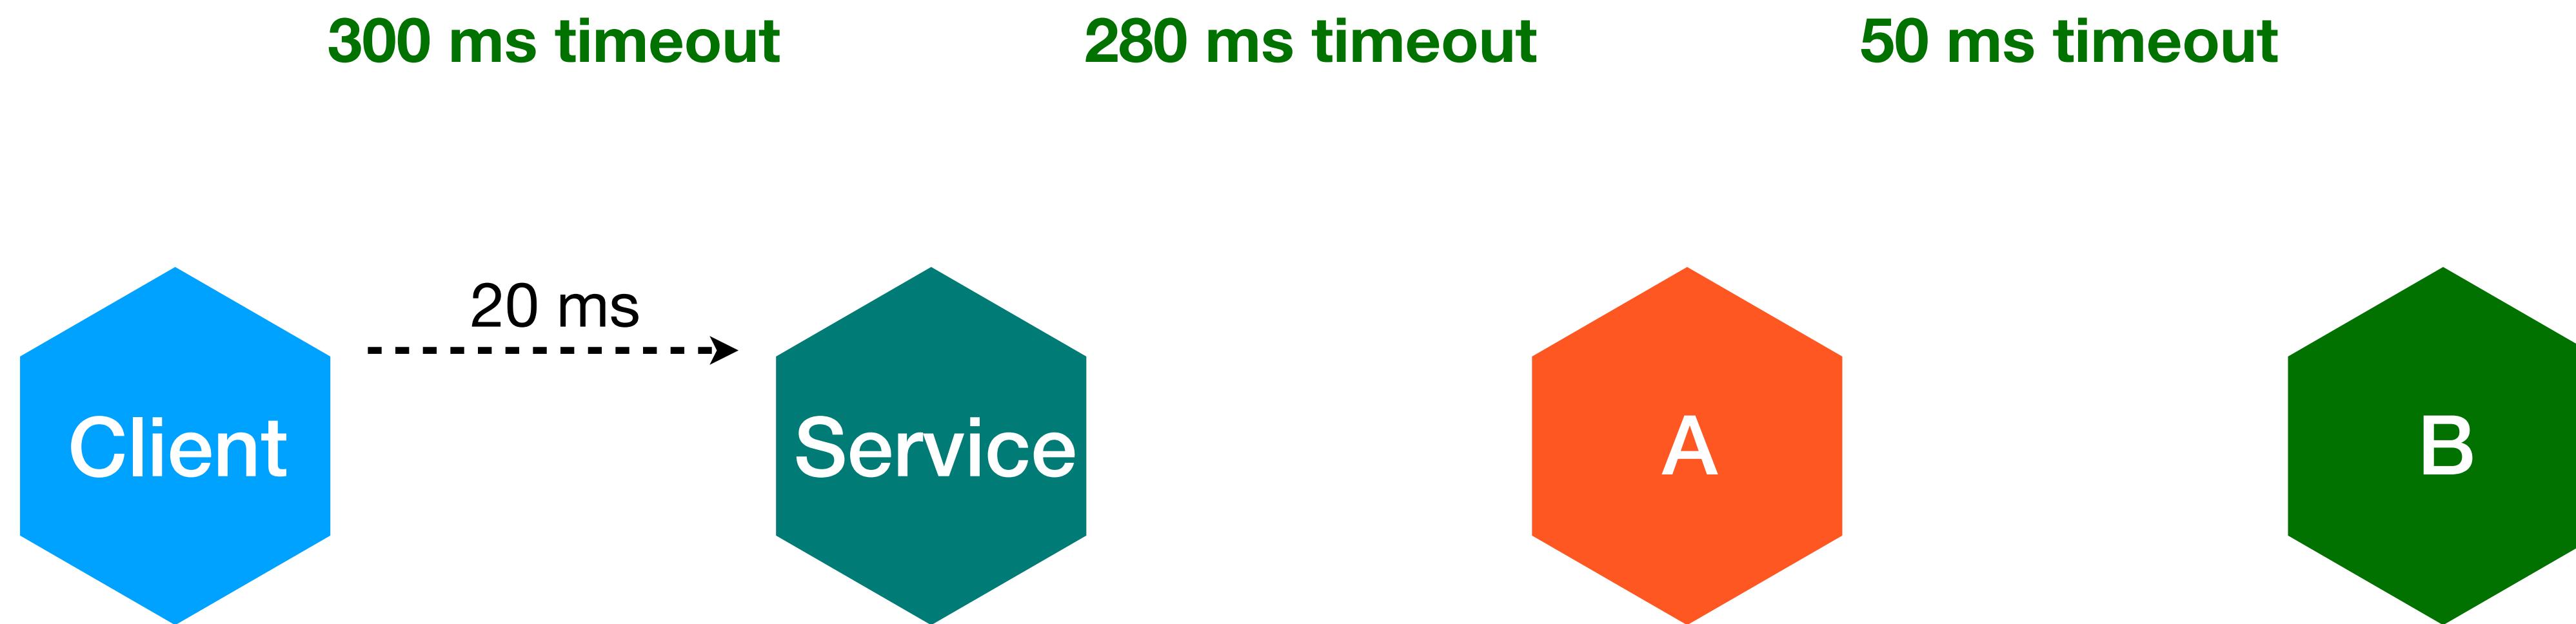
**280 ms timeout**



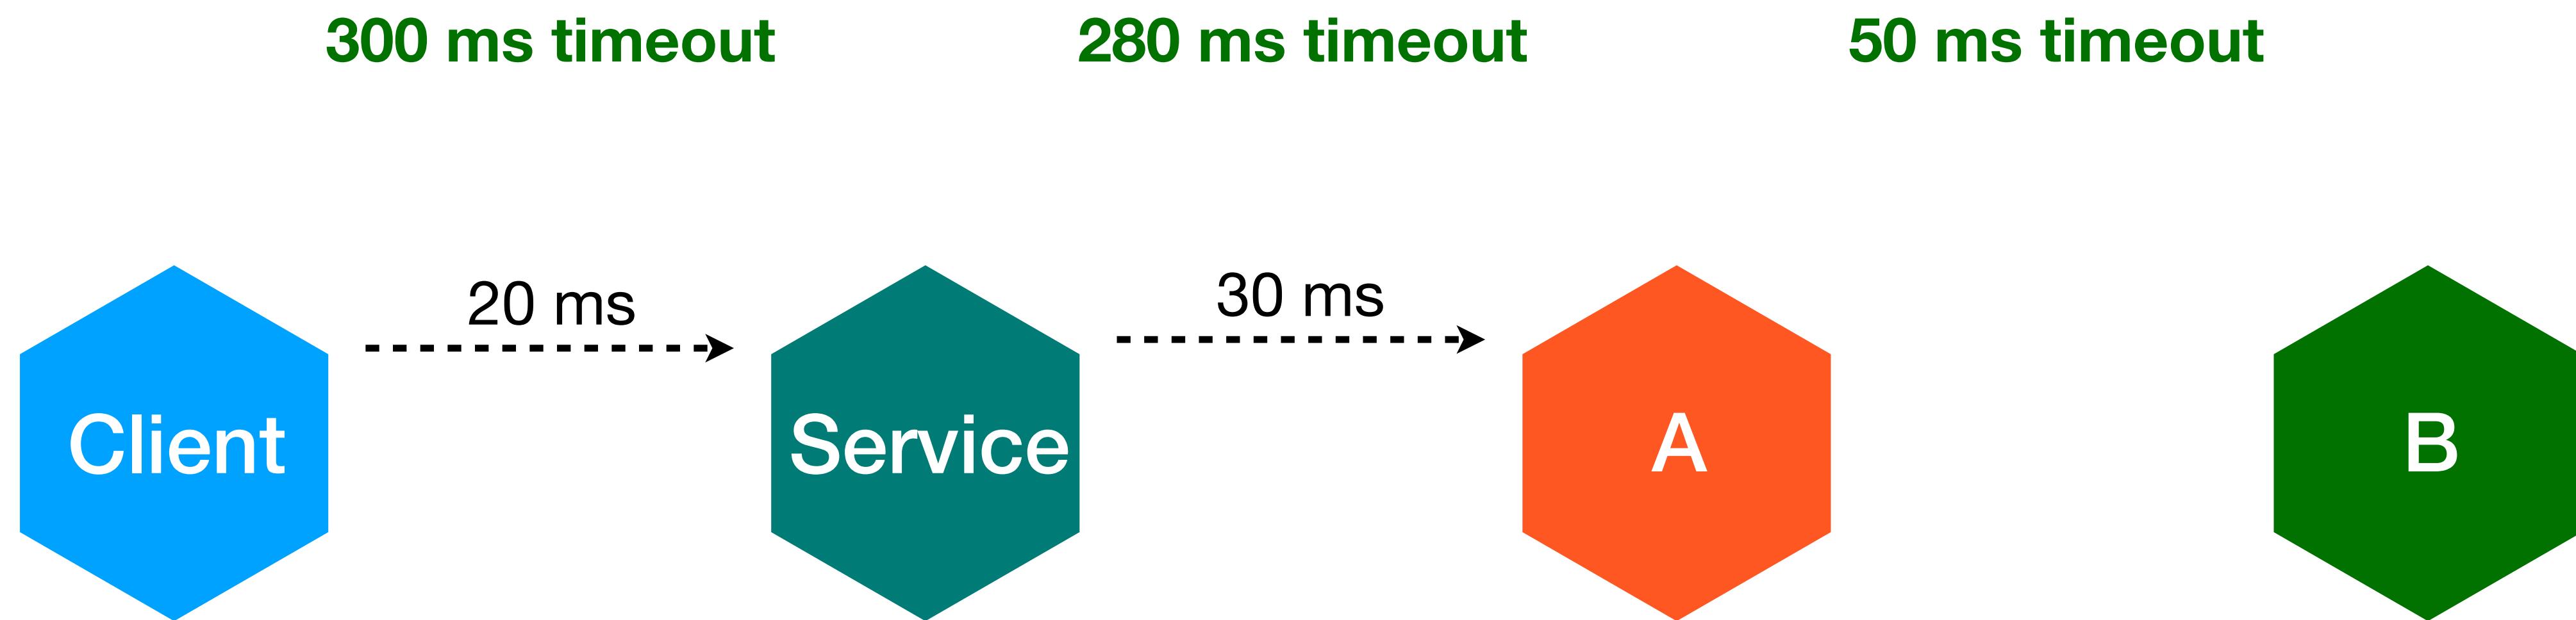
**50 ms timeout**



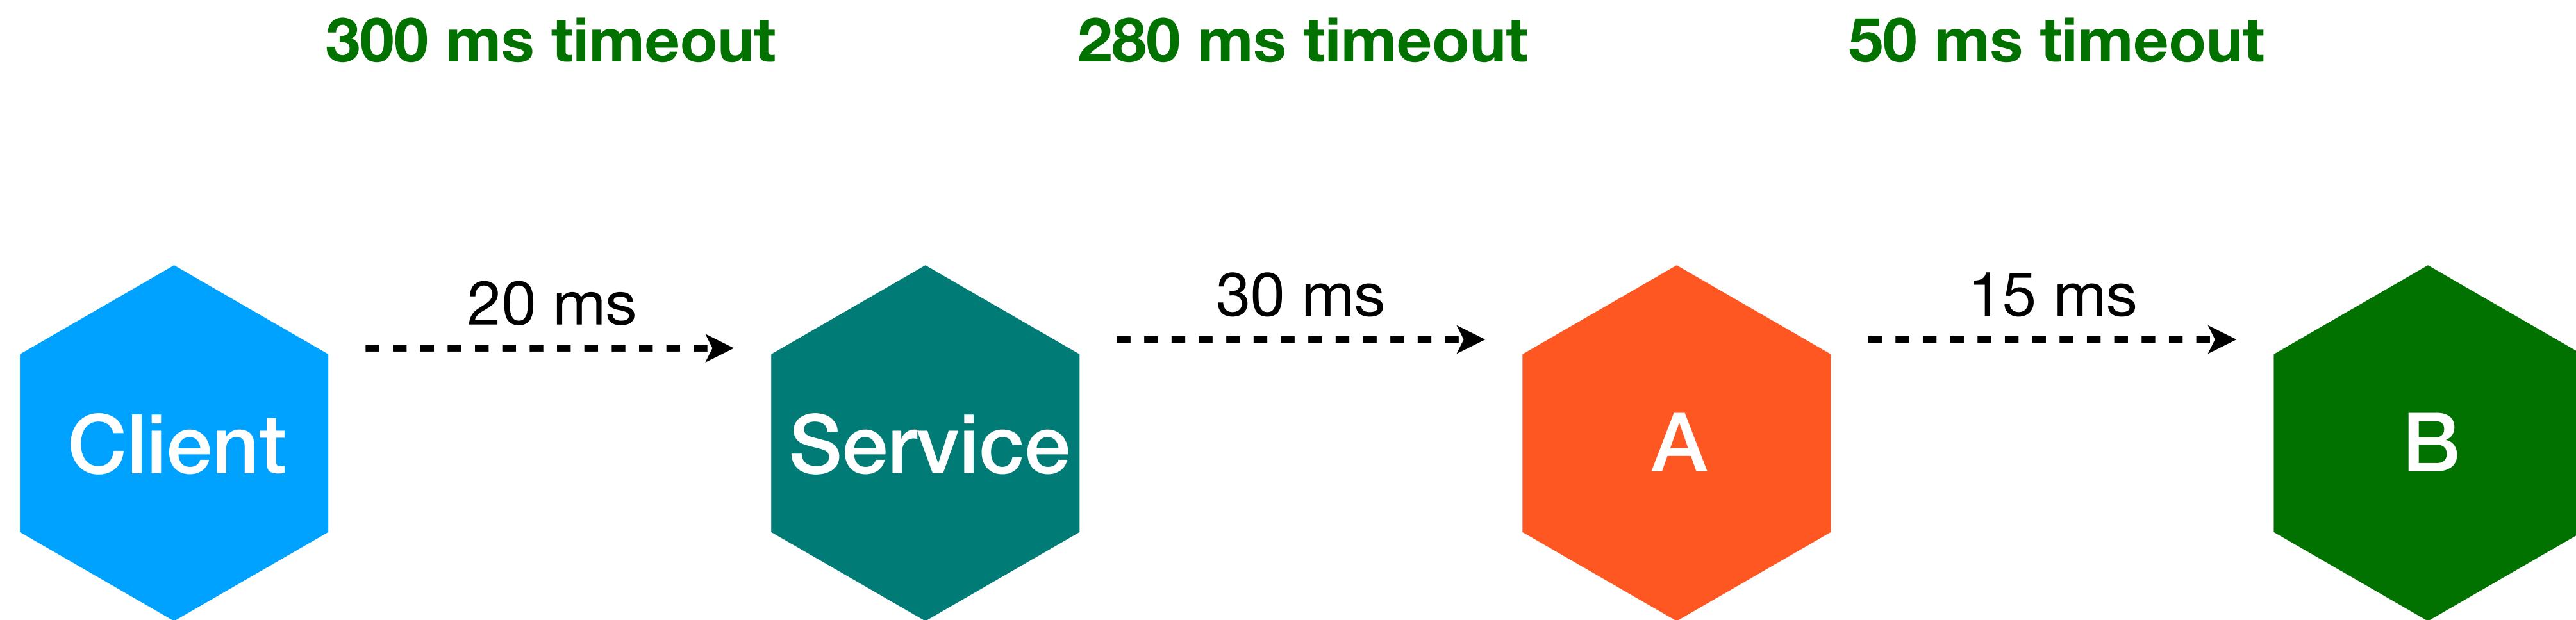
# Fine-tuned timeout



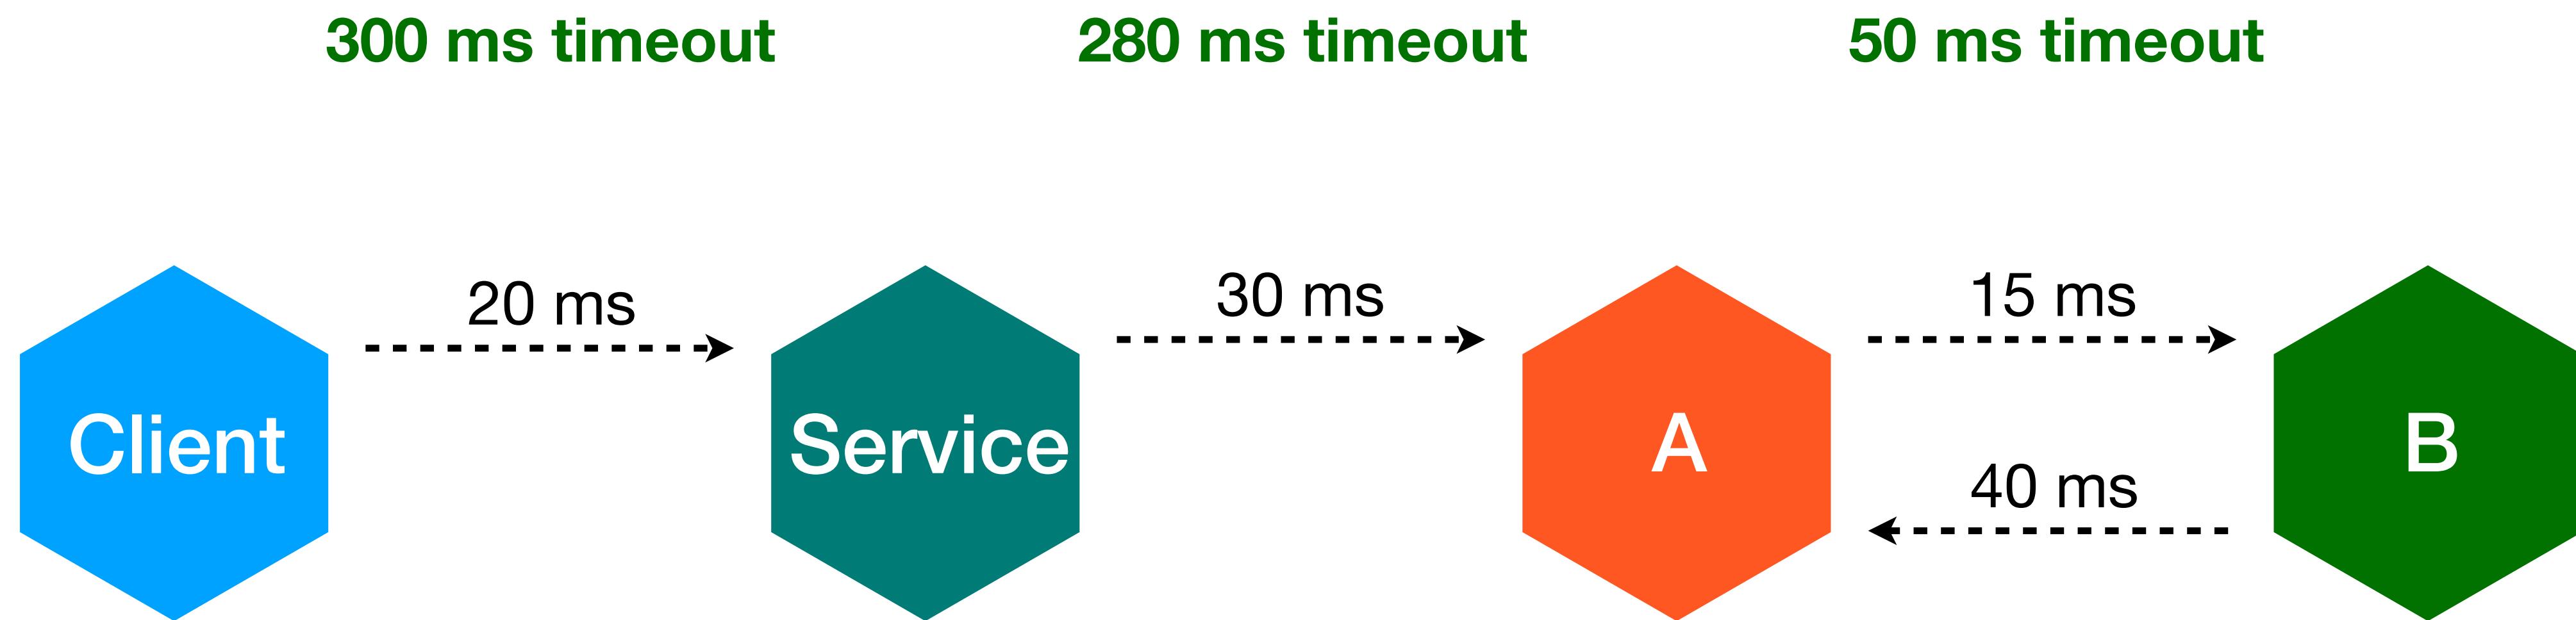
# Fine-tuned timeout



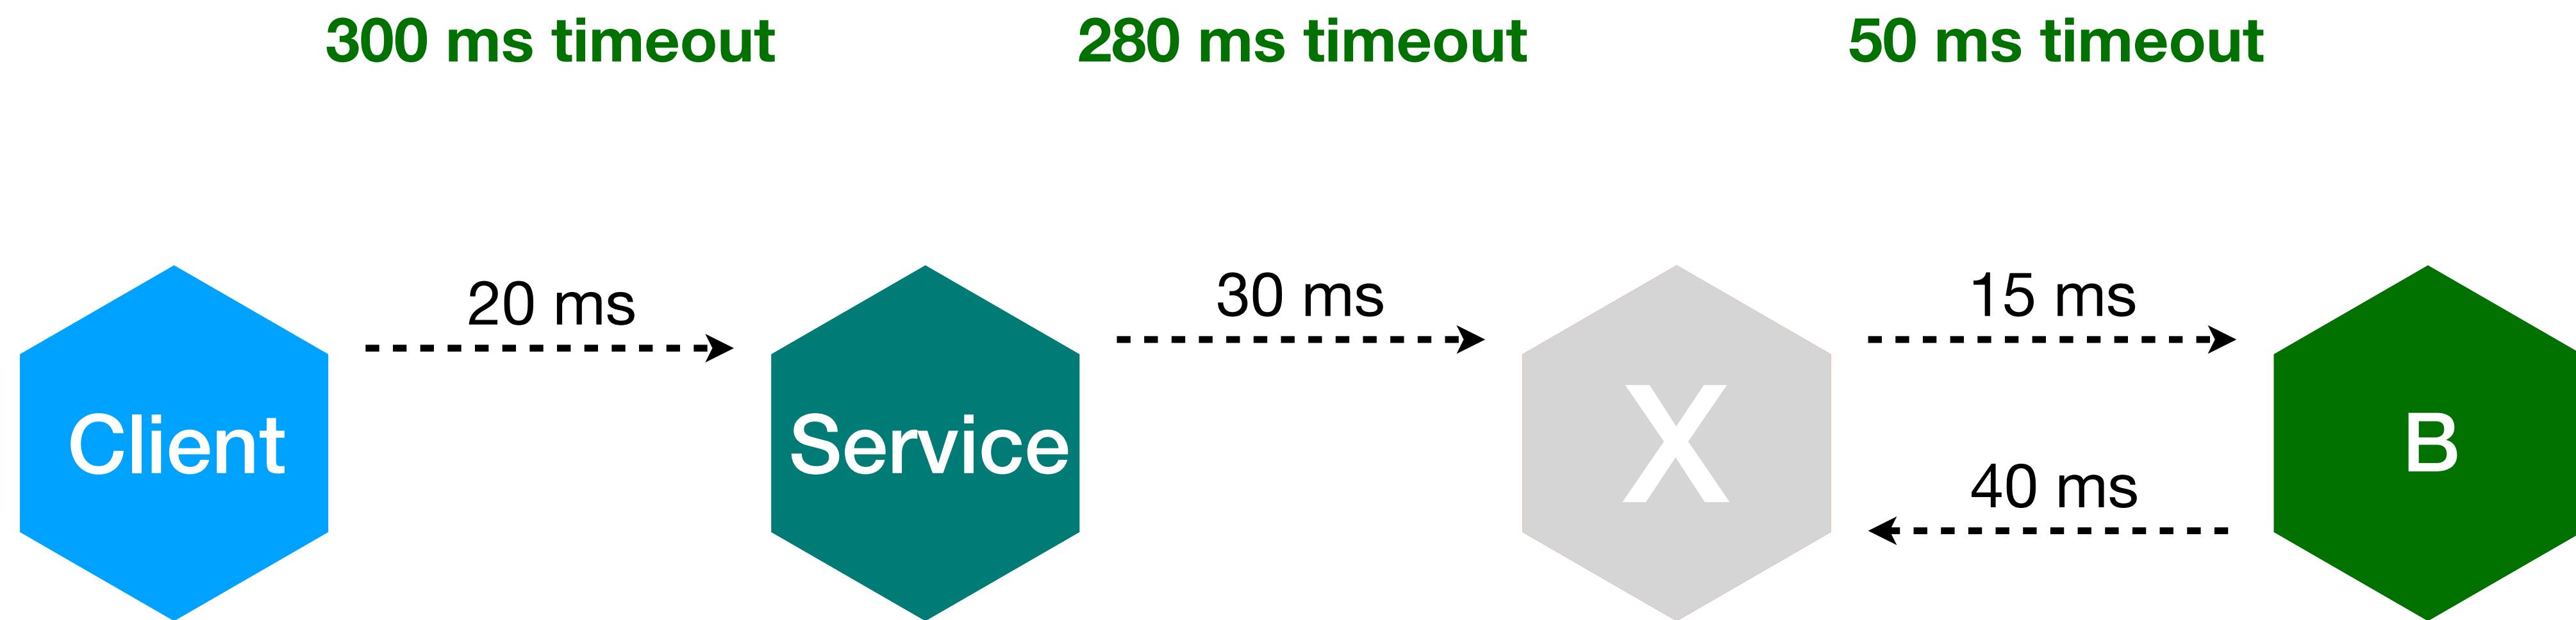
# Fine-tuned timeout



# Fine-tuned timeout



# Fine-tuned timeout



# Adaptive timeout

Client

Service

A

B

# Adaptive timeout

**200 ms timeout**



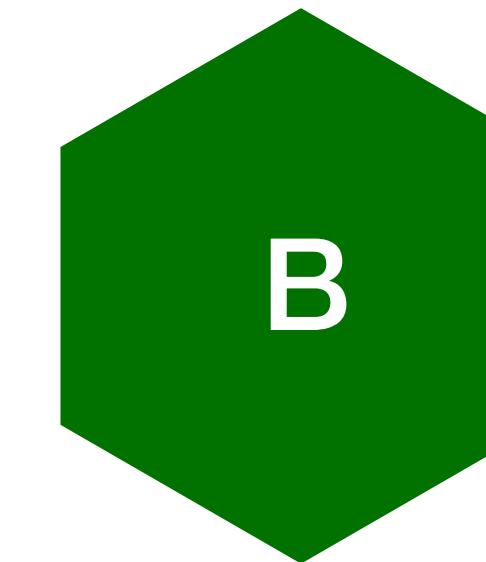
Client



Service



A



B

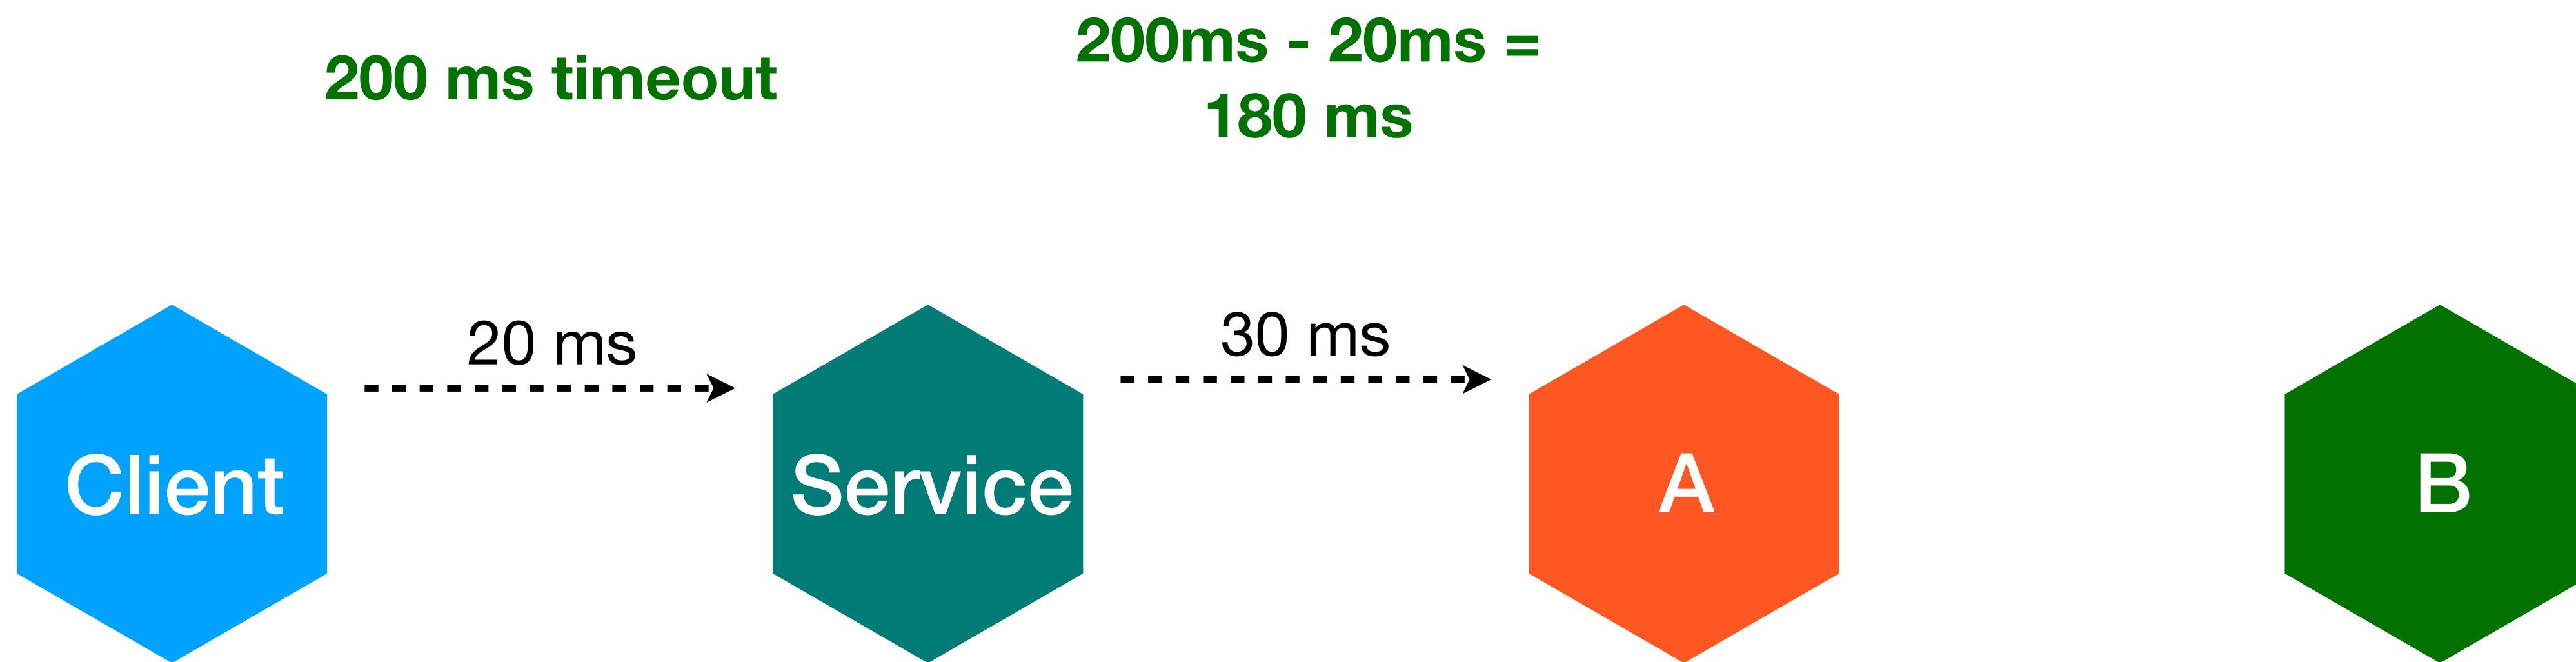
# Adaptive timeout



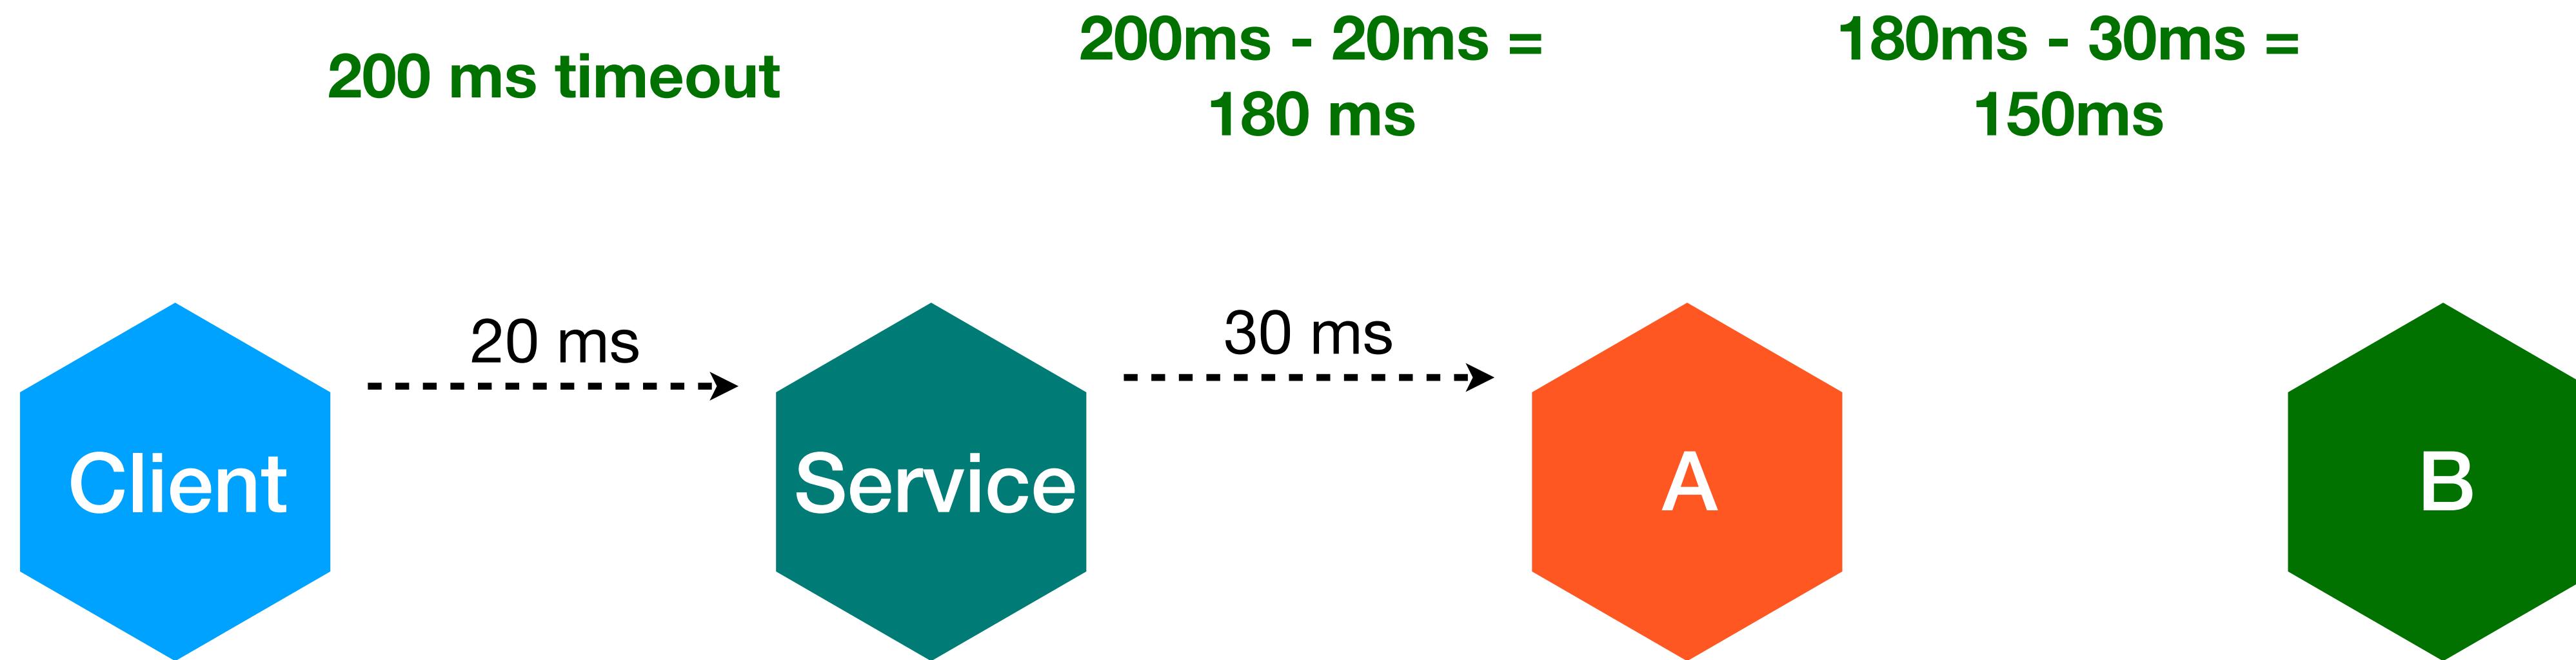
# Adaptive timeout



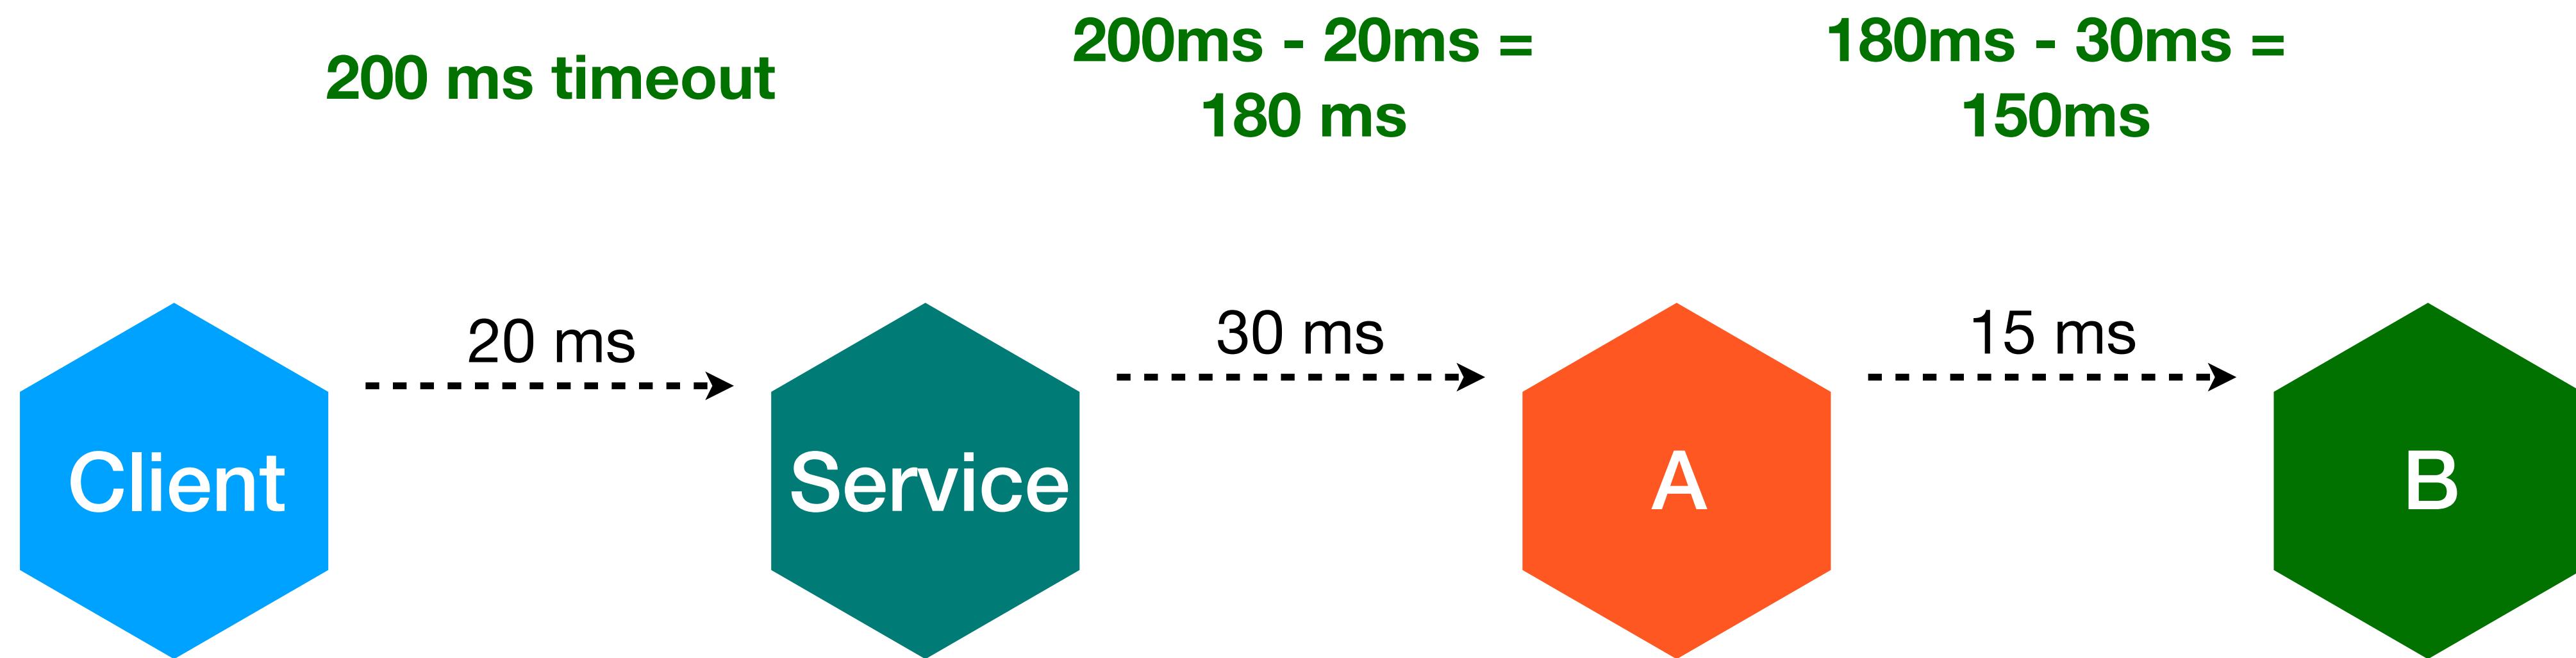
# Adaptive timeout



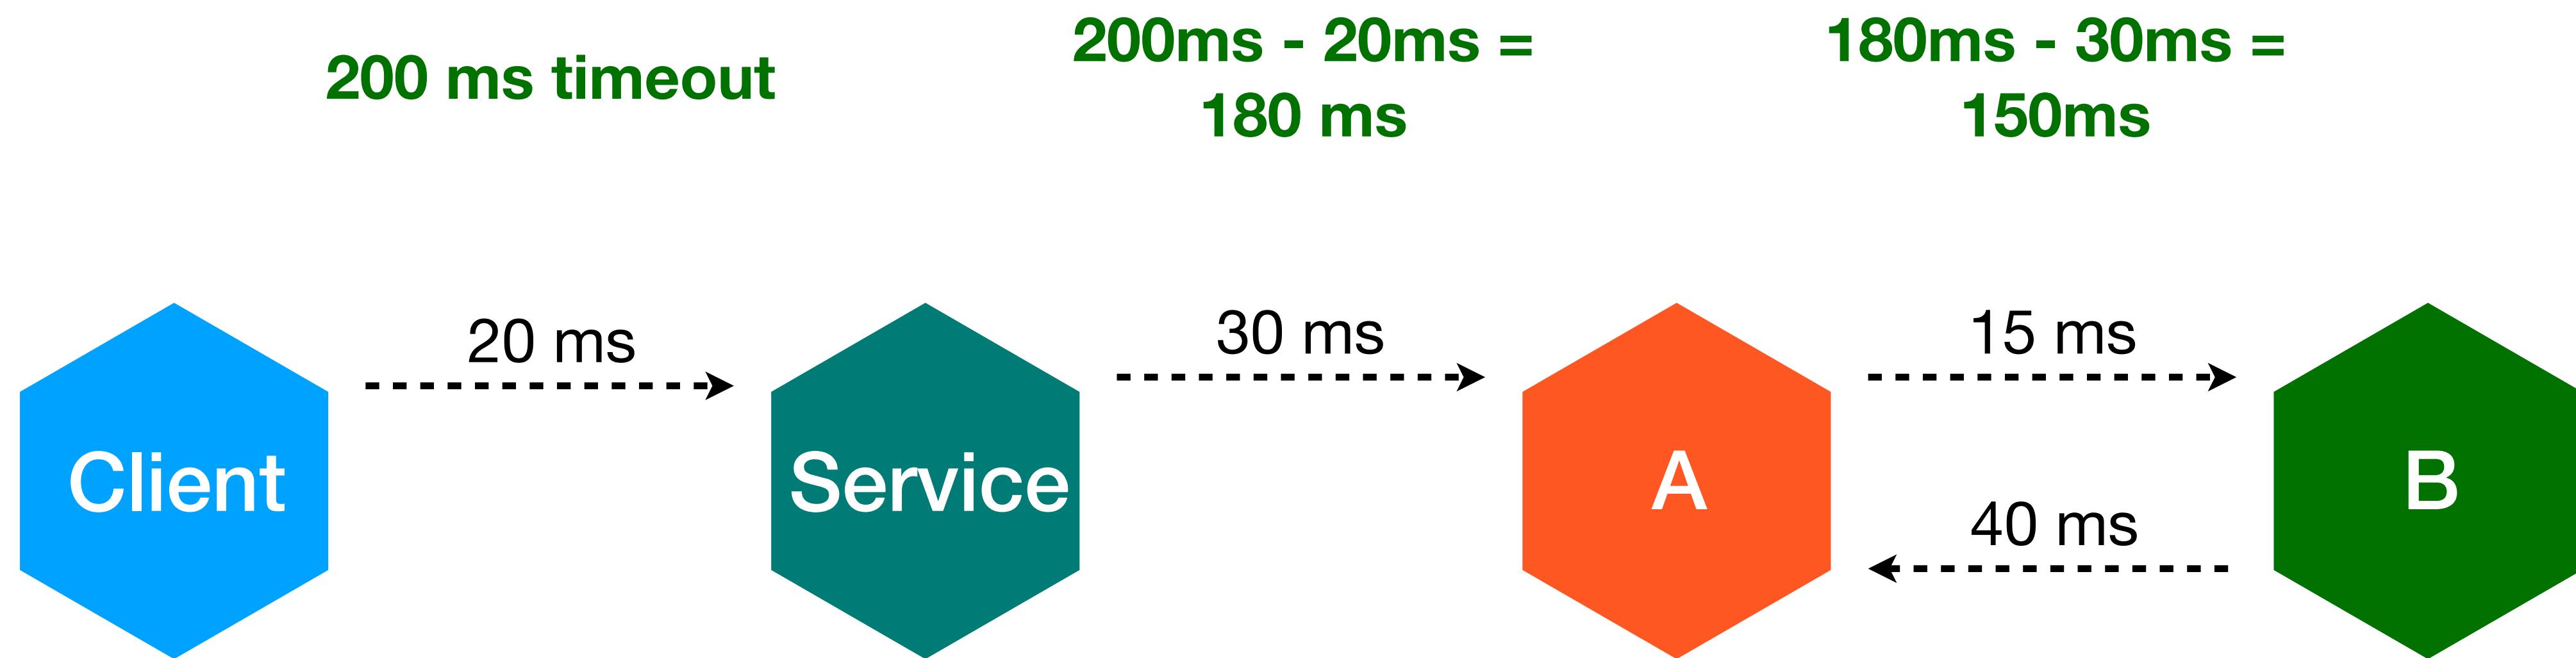
# Adaptive timeout



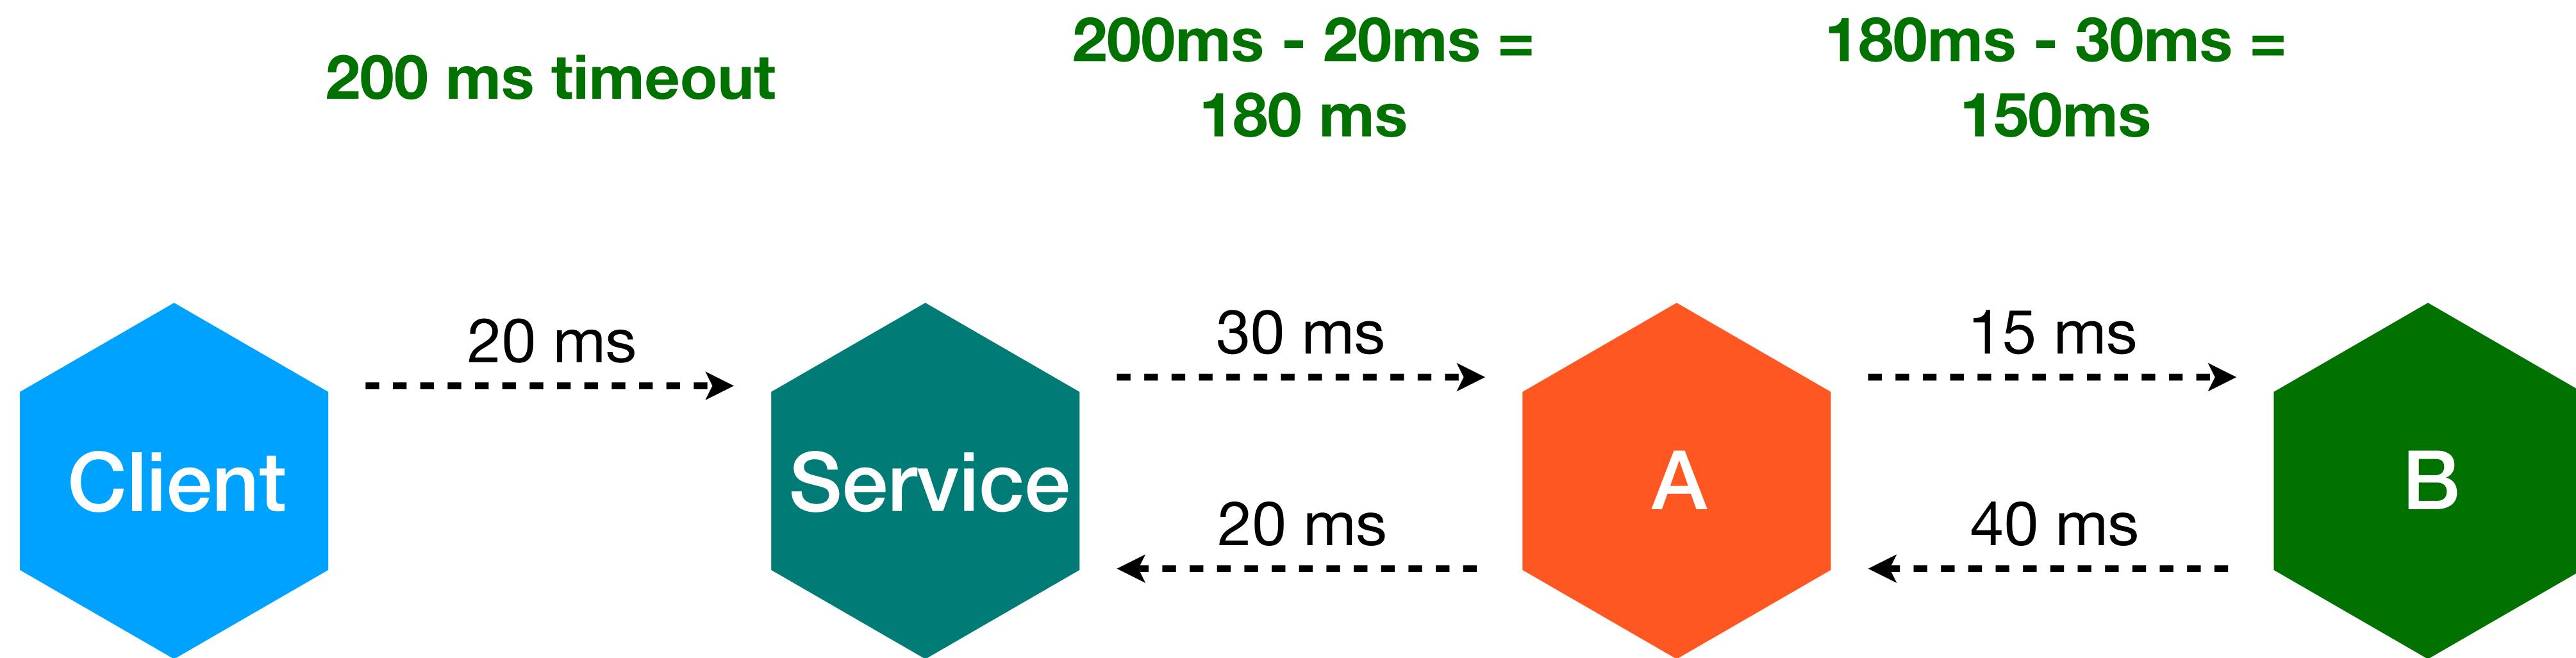
# Adaptive timeout



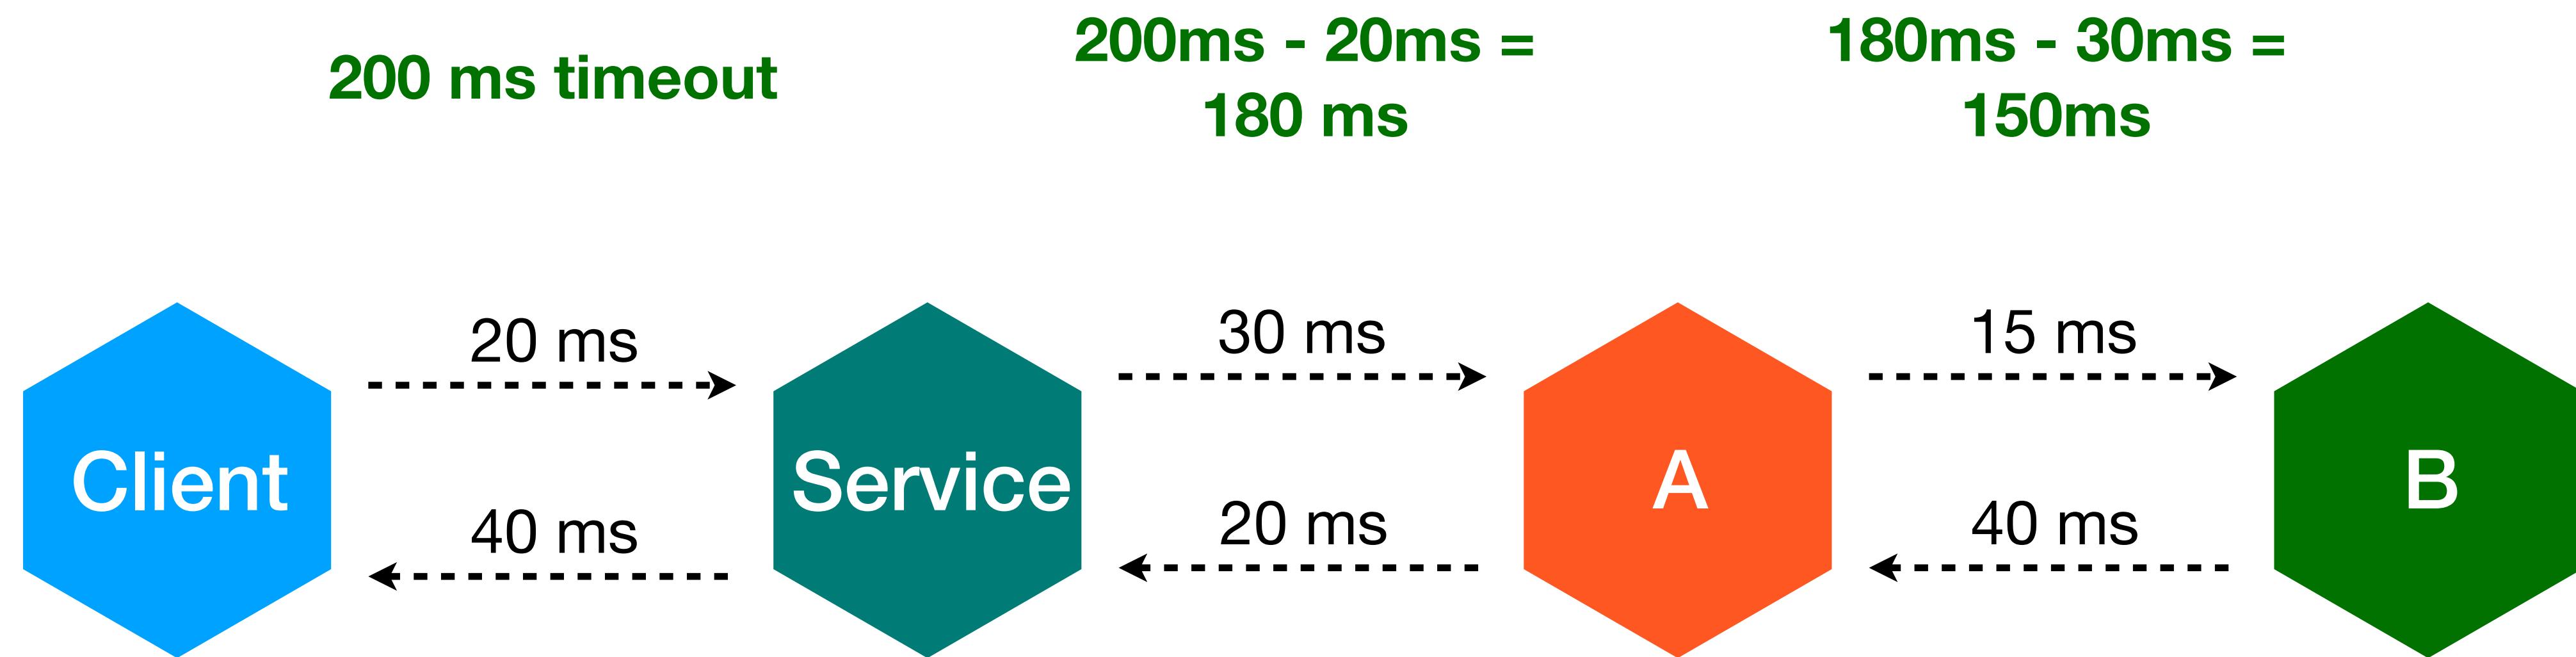
# Adaptive timeout



# Adaptive timeout



# Adaptive timeout



# gRPC: Deadlines

# gRPC: Deadlines

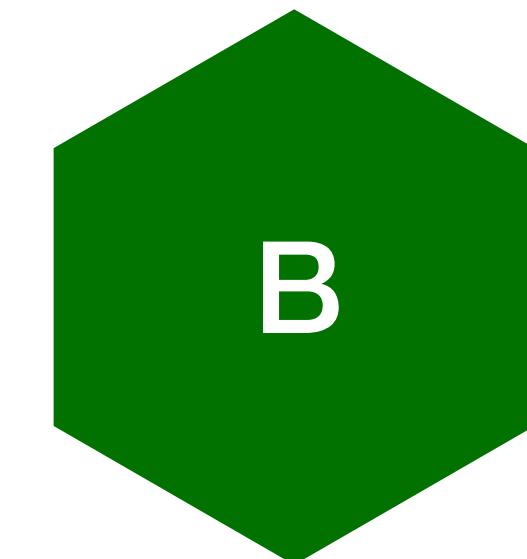
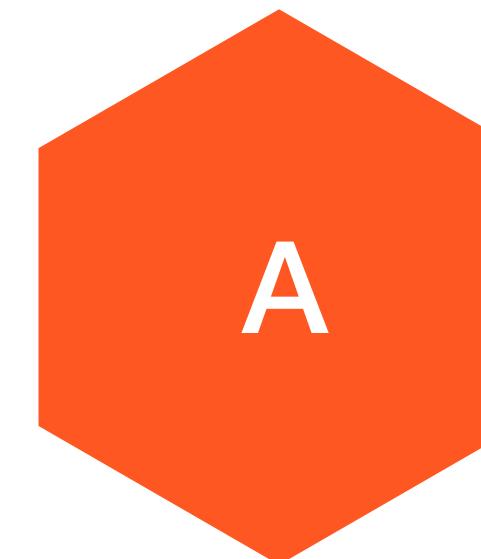
- Timeout is relative

# gRPC: Deadlines

- Timeout is relative
- Deadline is absolute

# Deadline propagation

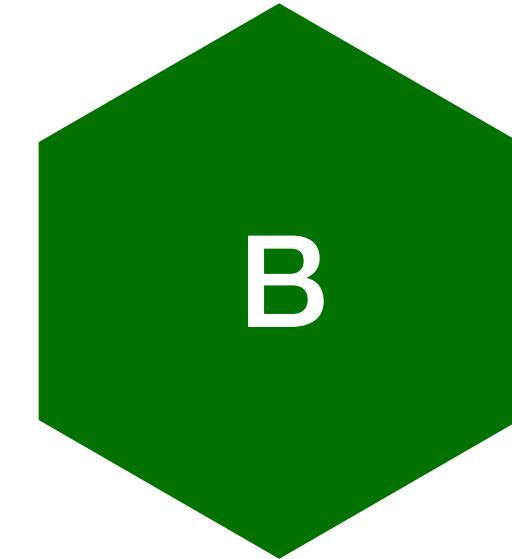
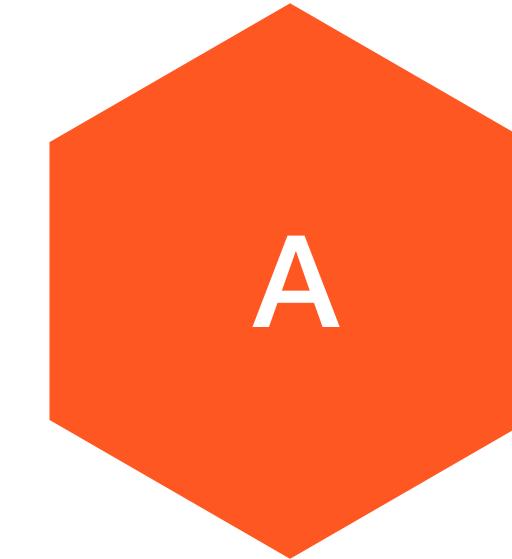
Start TS: **3600000**  
Timeout: 200  
Deadline: **3600200**



# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

TS: 3600000



# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

TS: 3600000

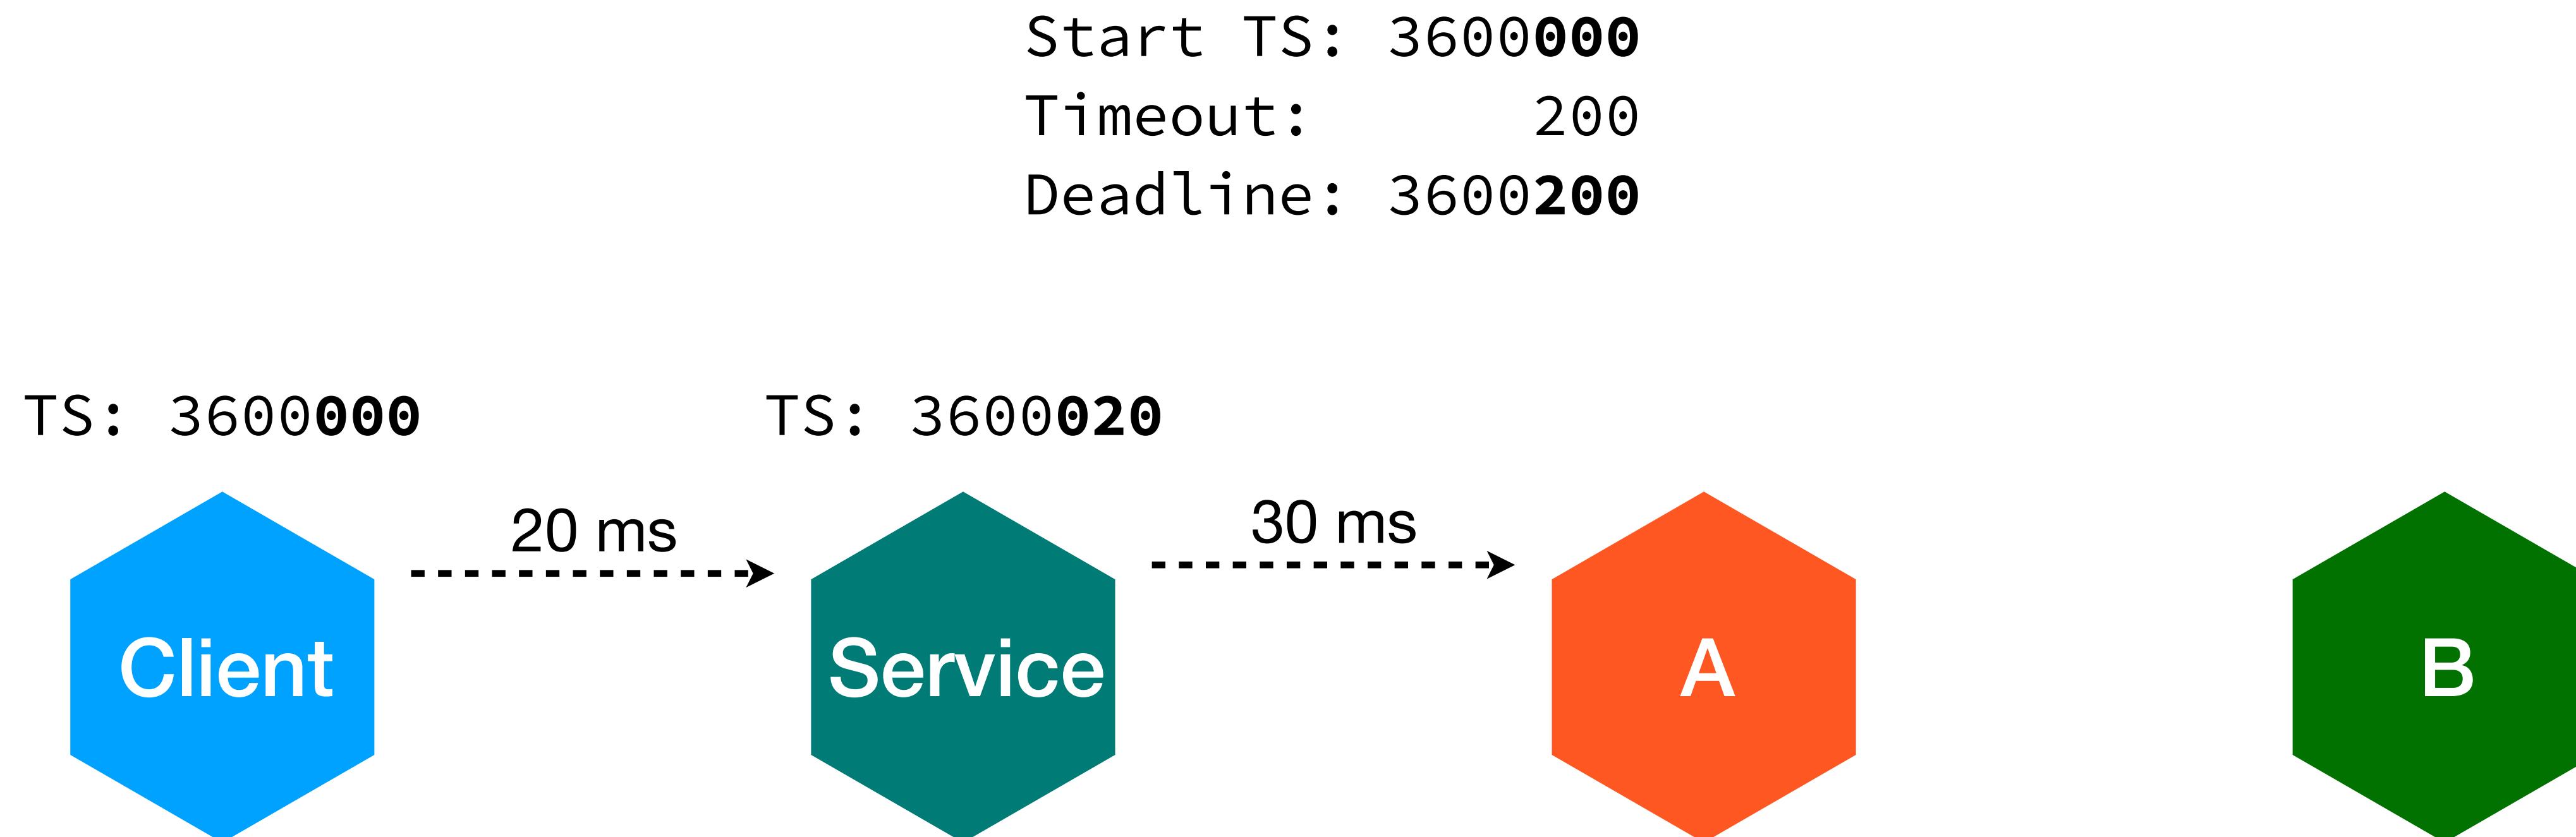


# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

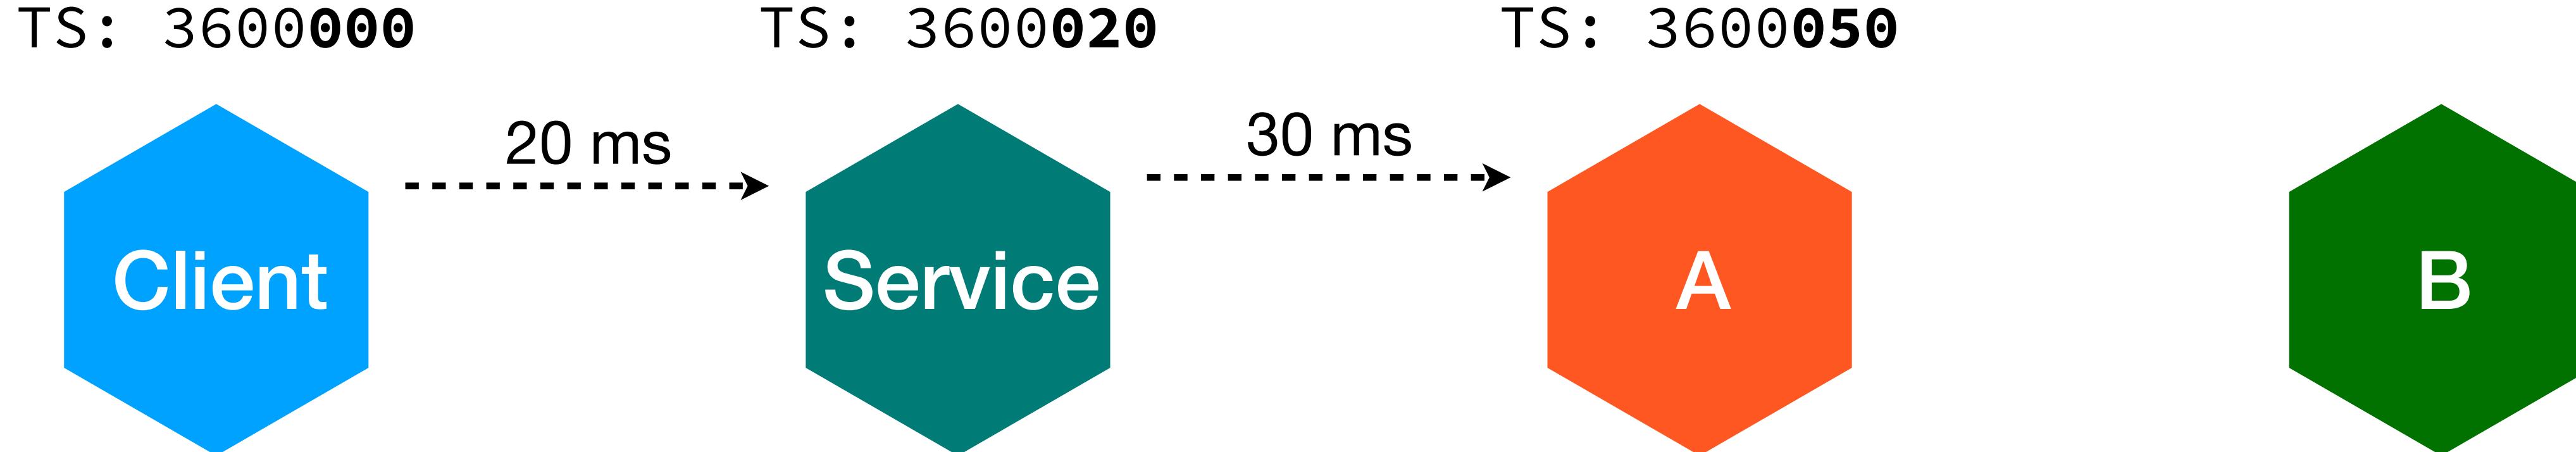


# Deadline propagation

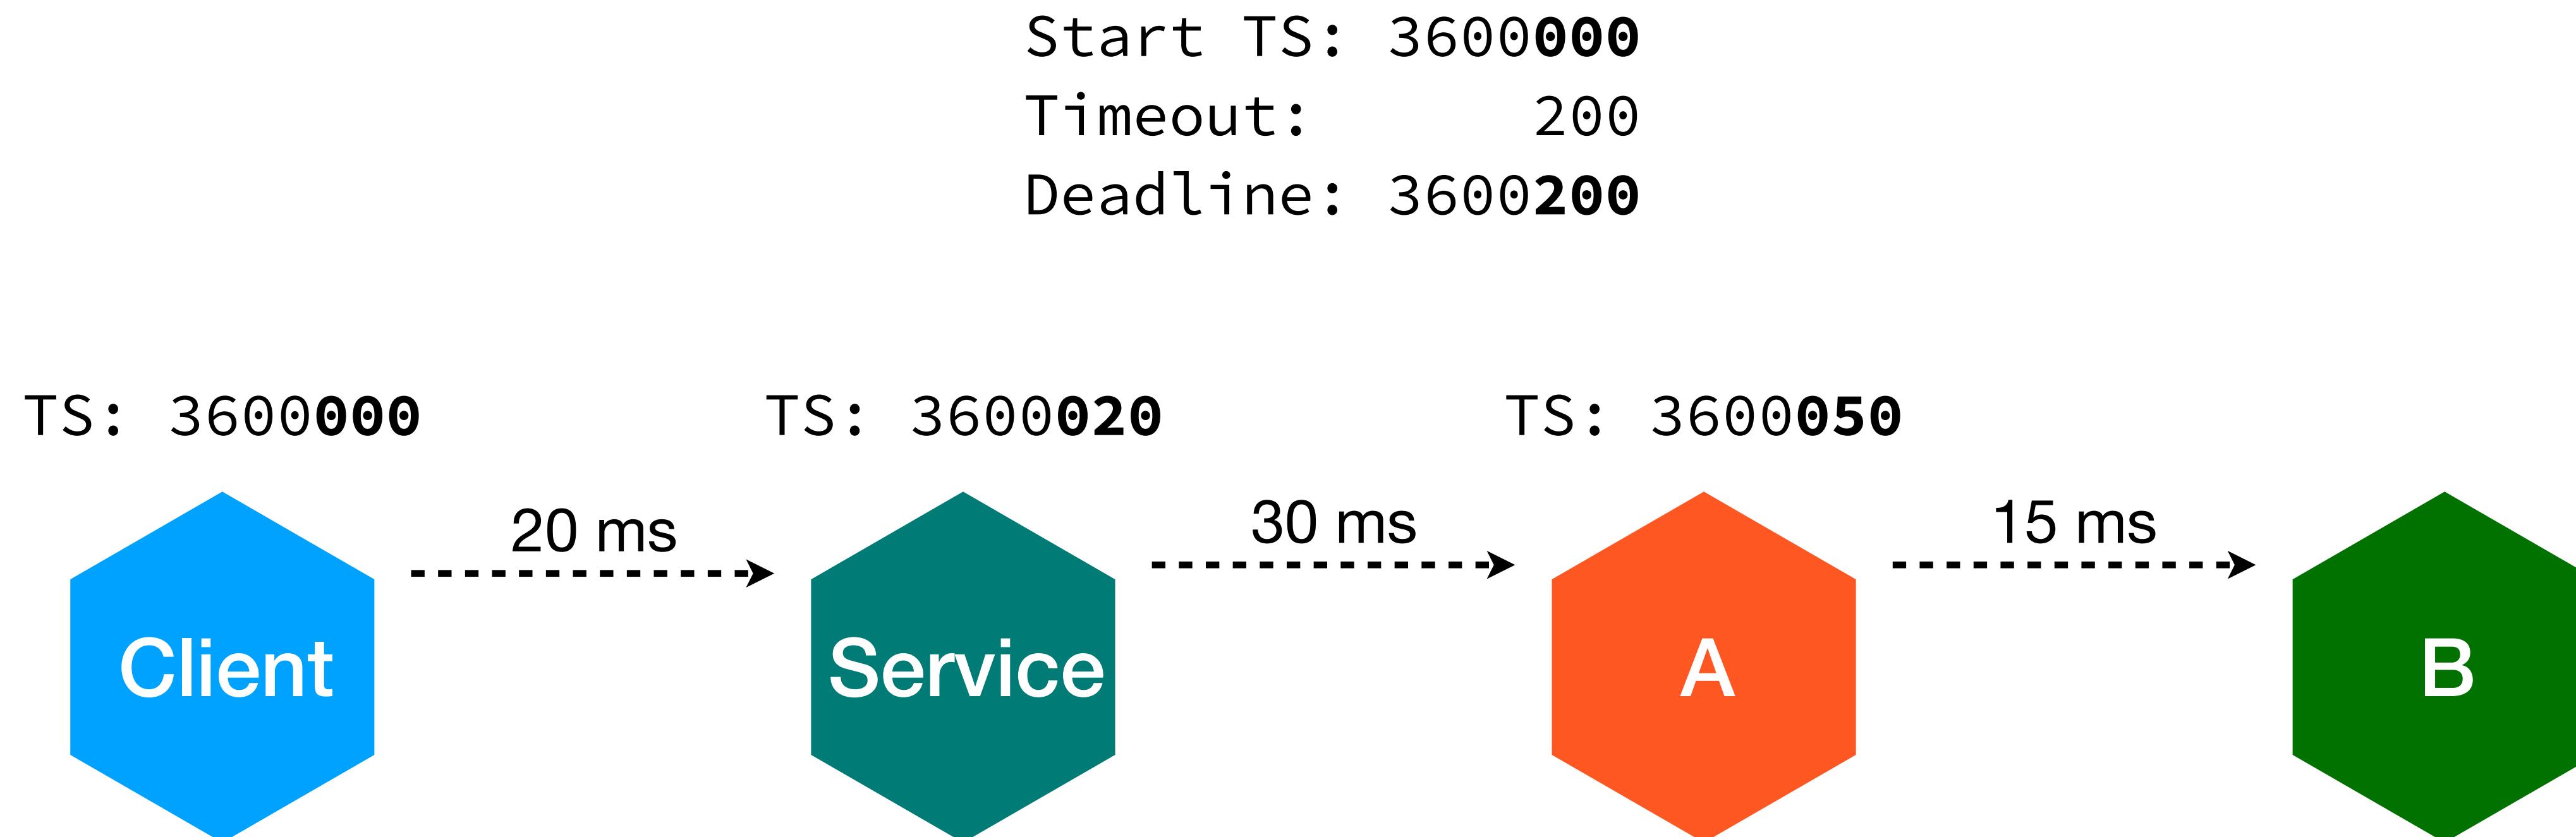


# Deadline propagation

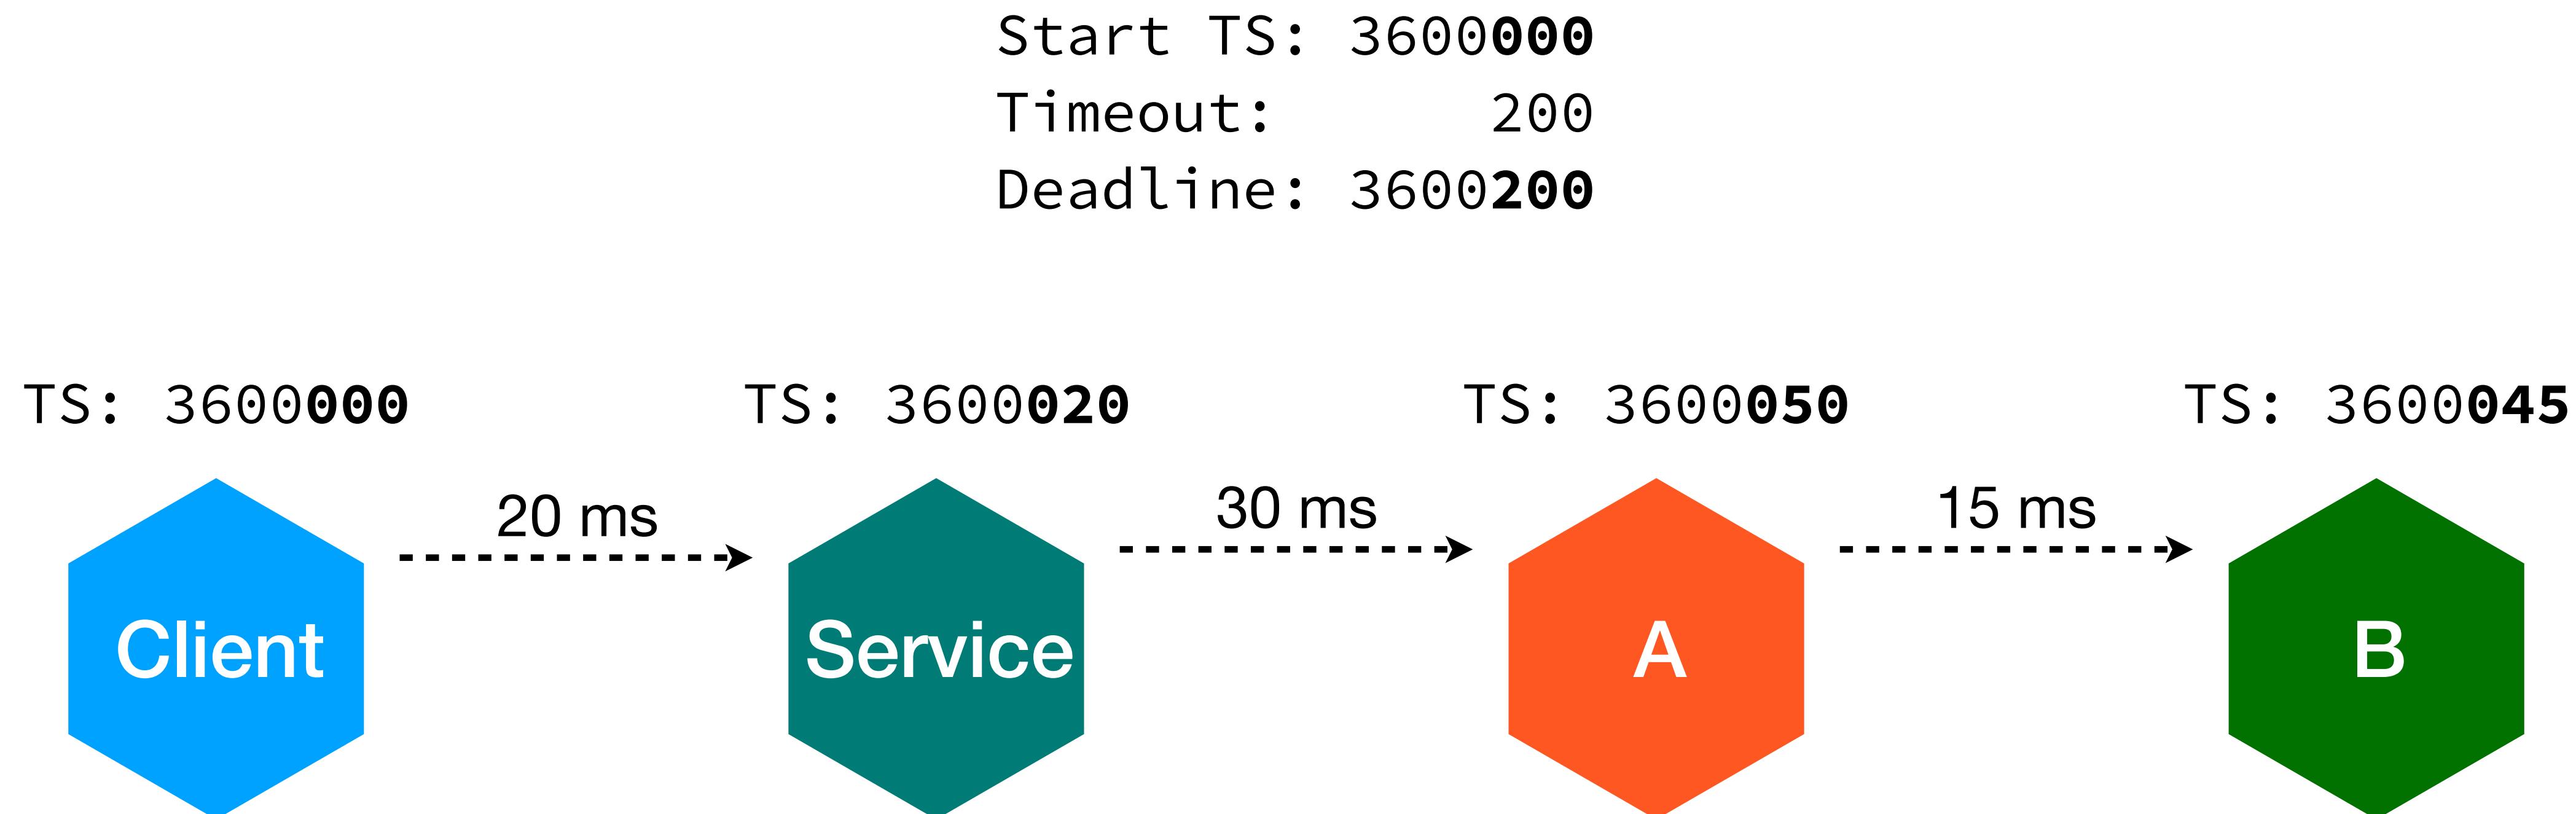
Start TS: 3600000  
Timeout: 200  
Deadline: 3600200



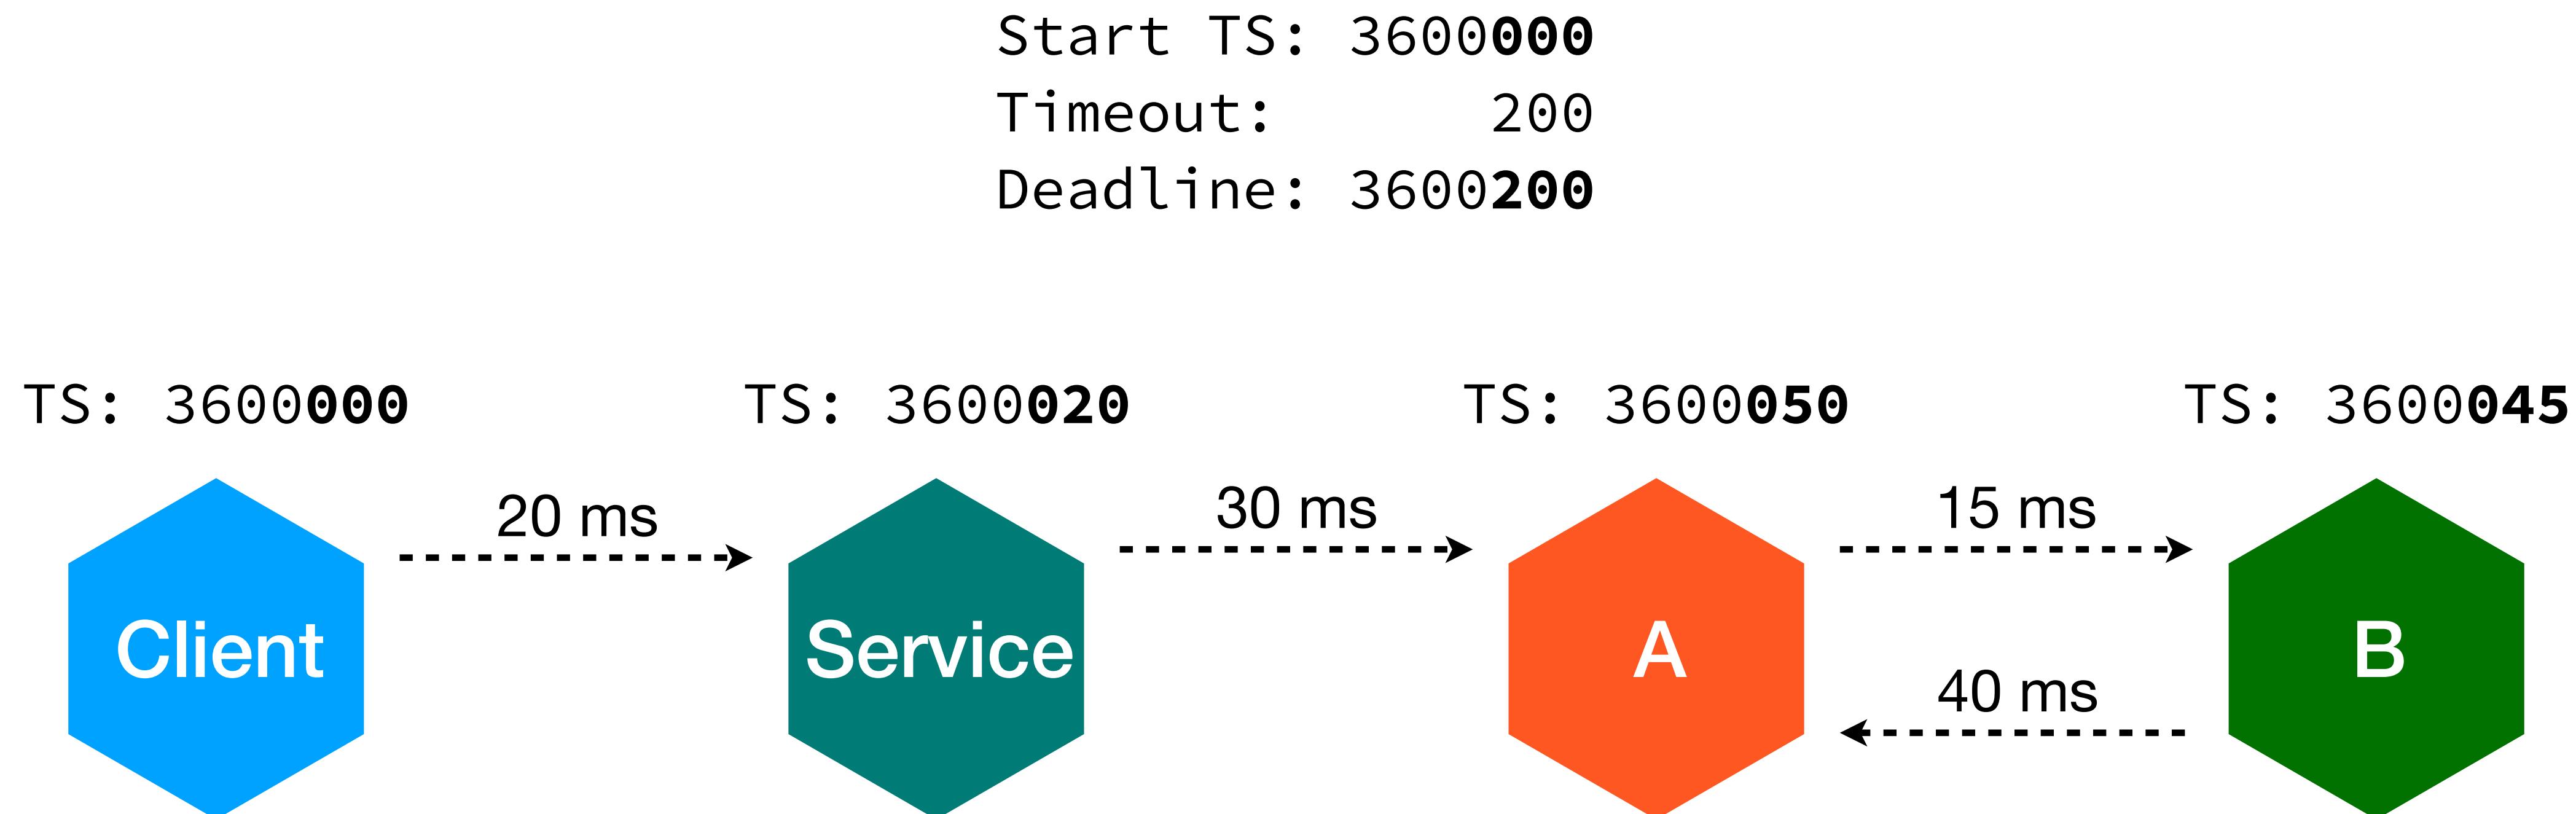
# Deadline propagation



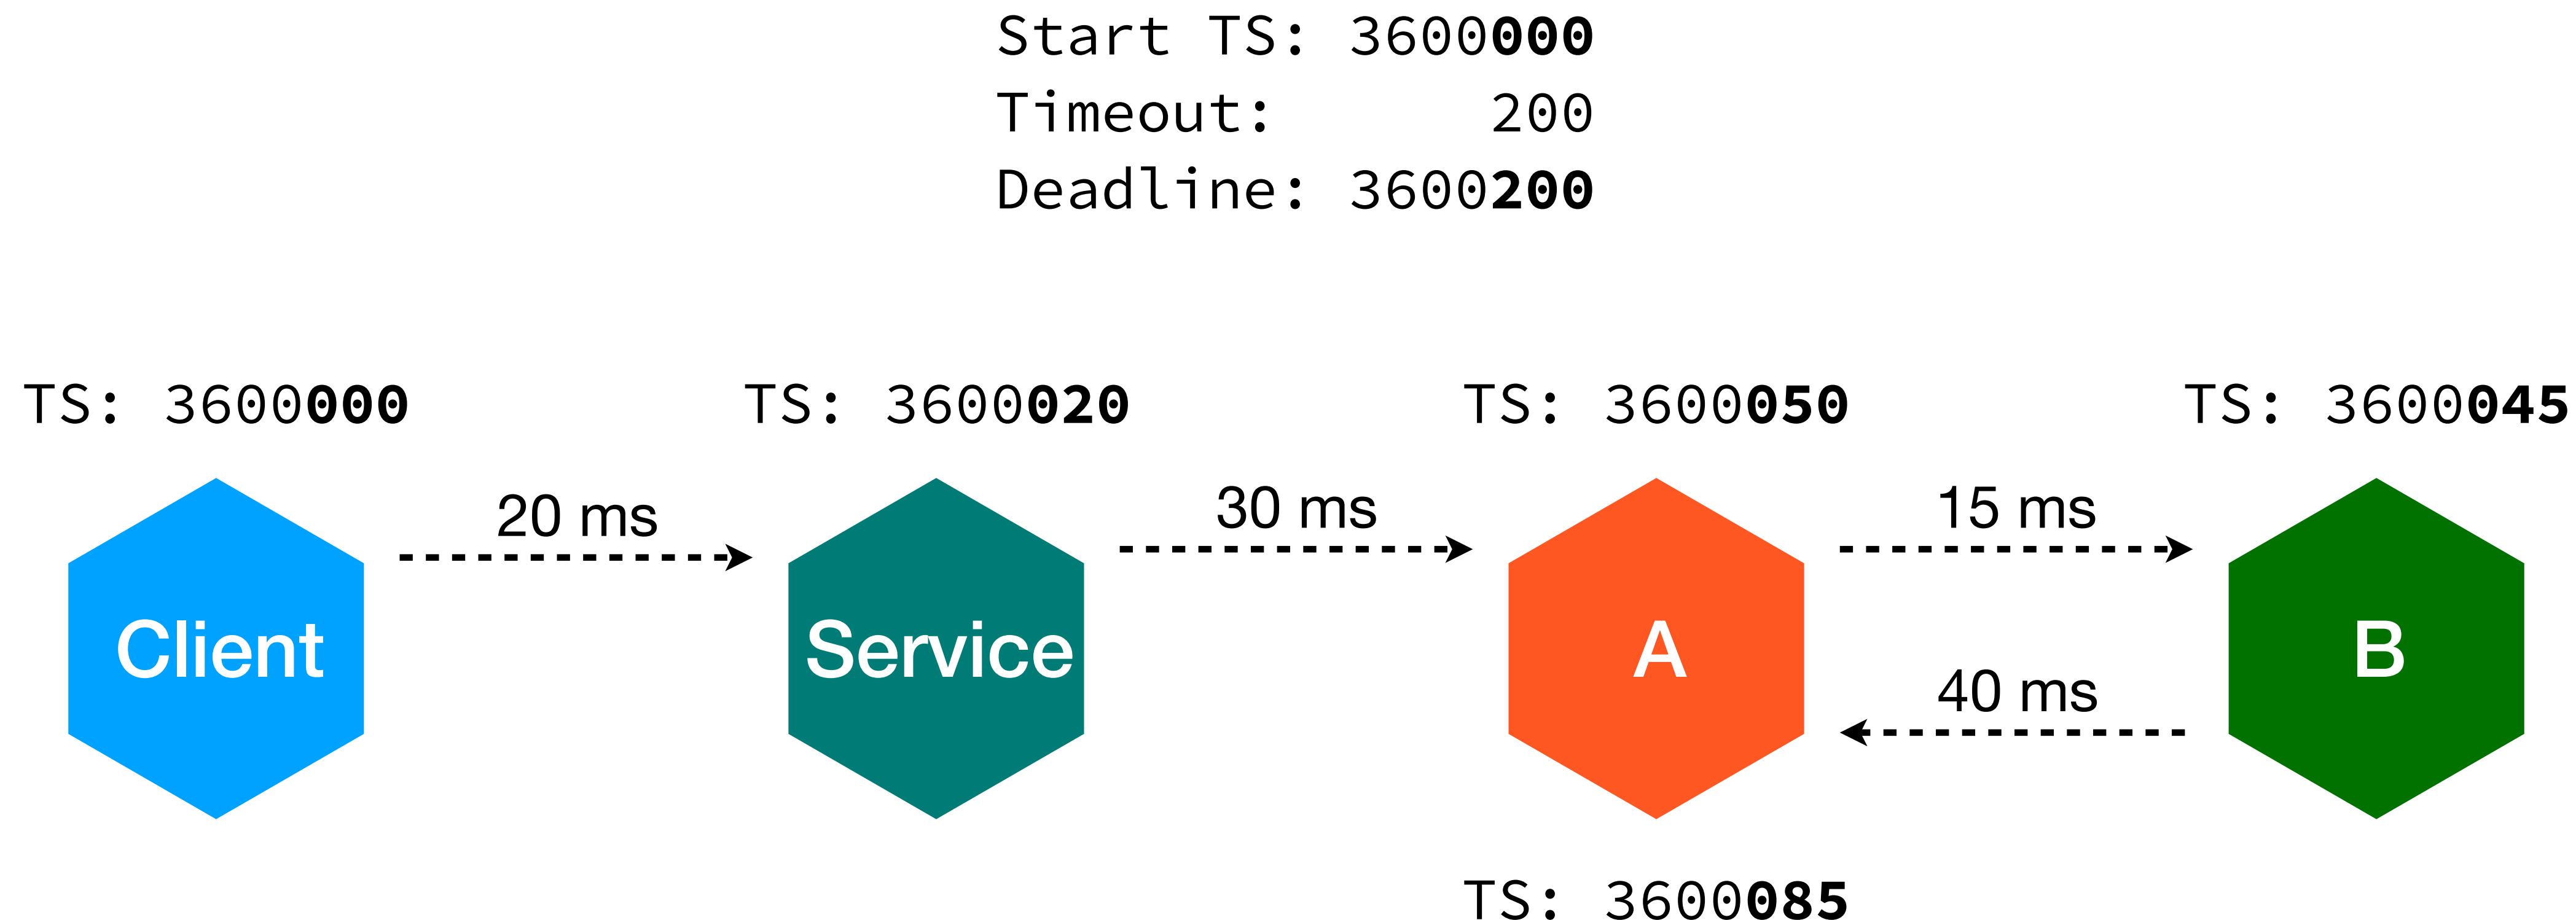
# Deadline propagation



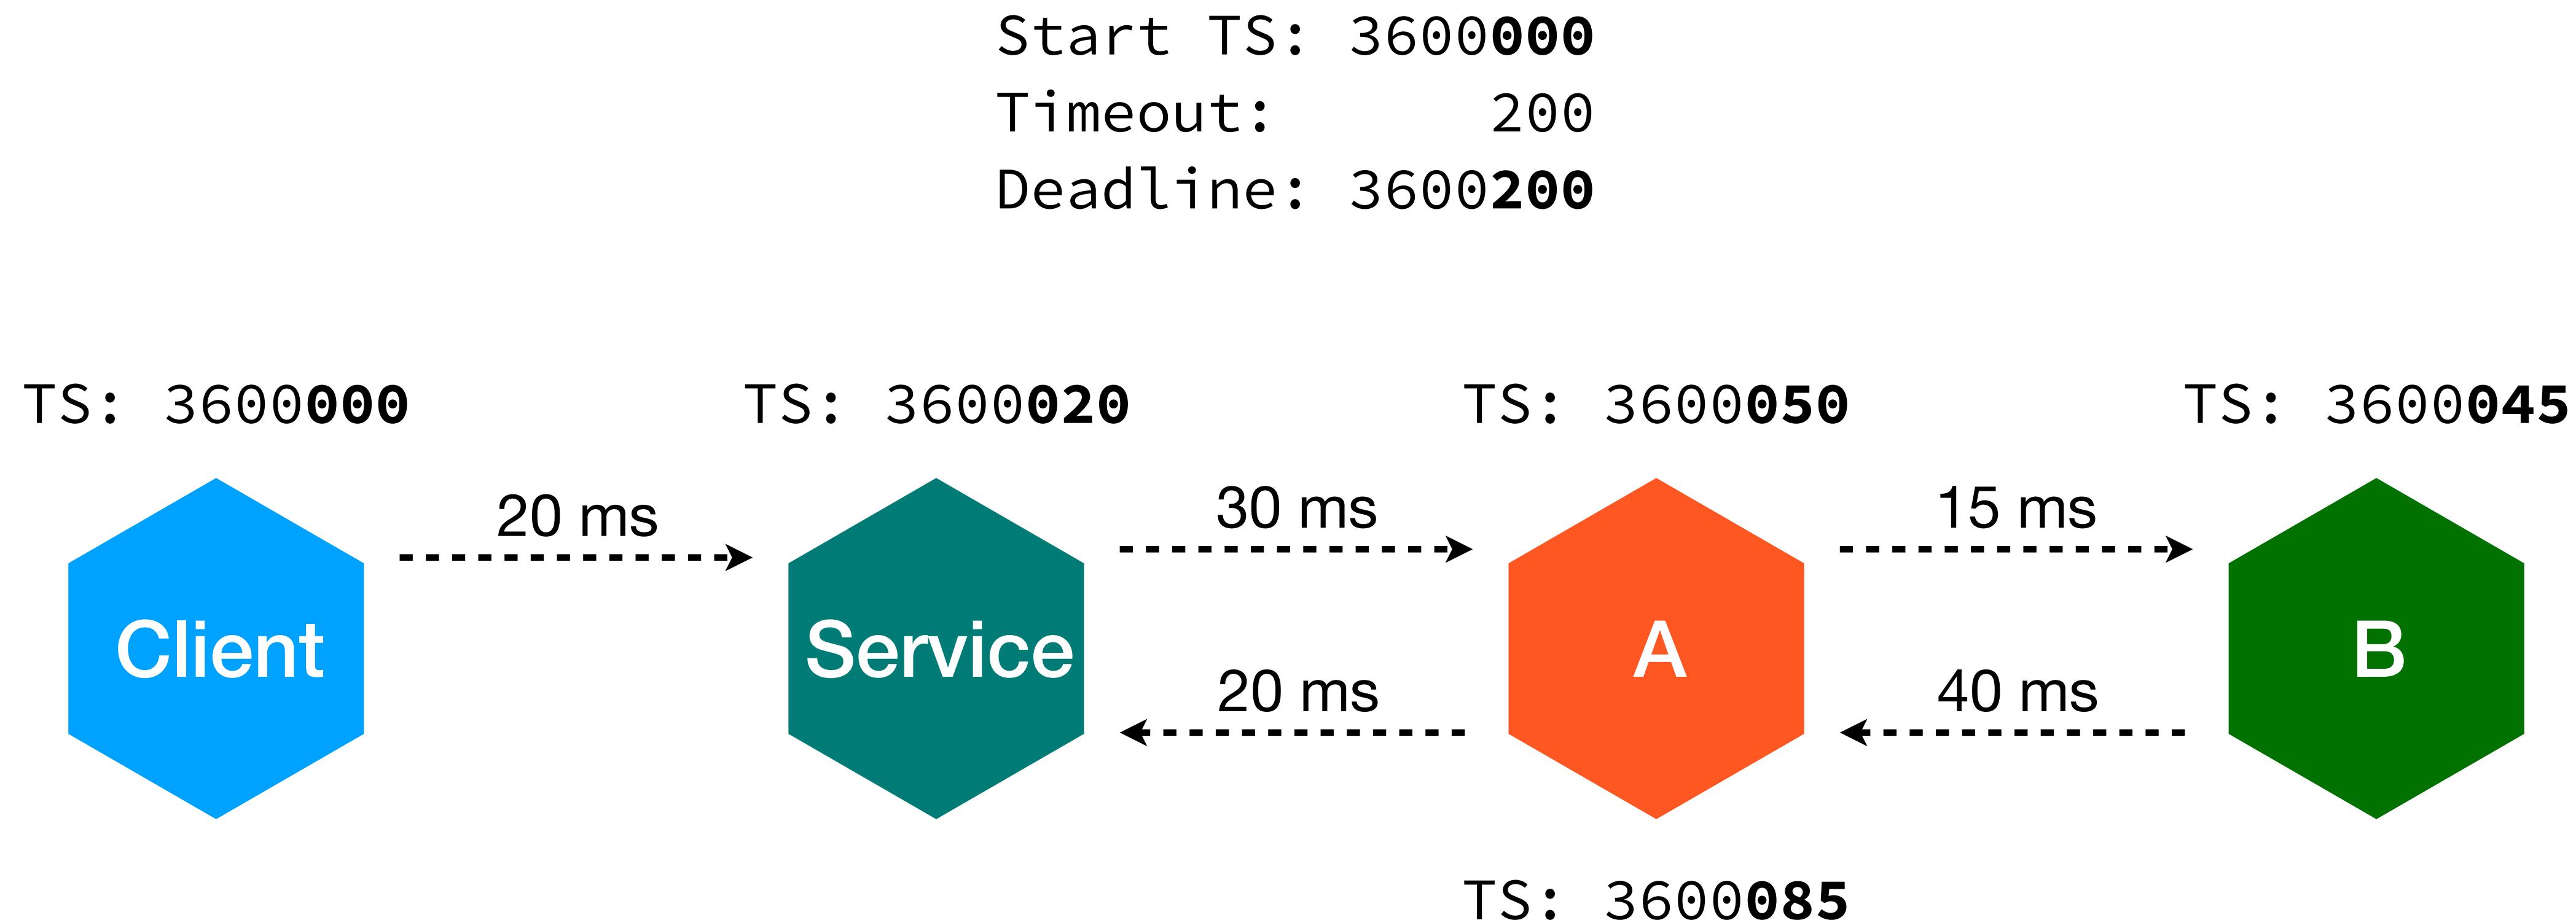
# Deadline propagation



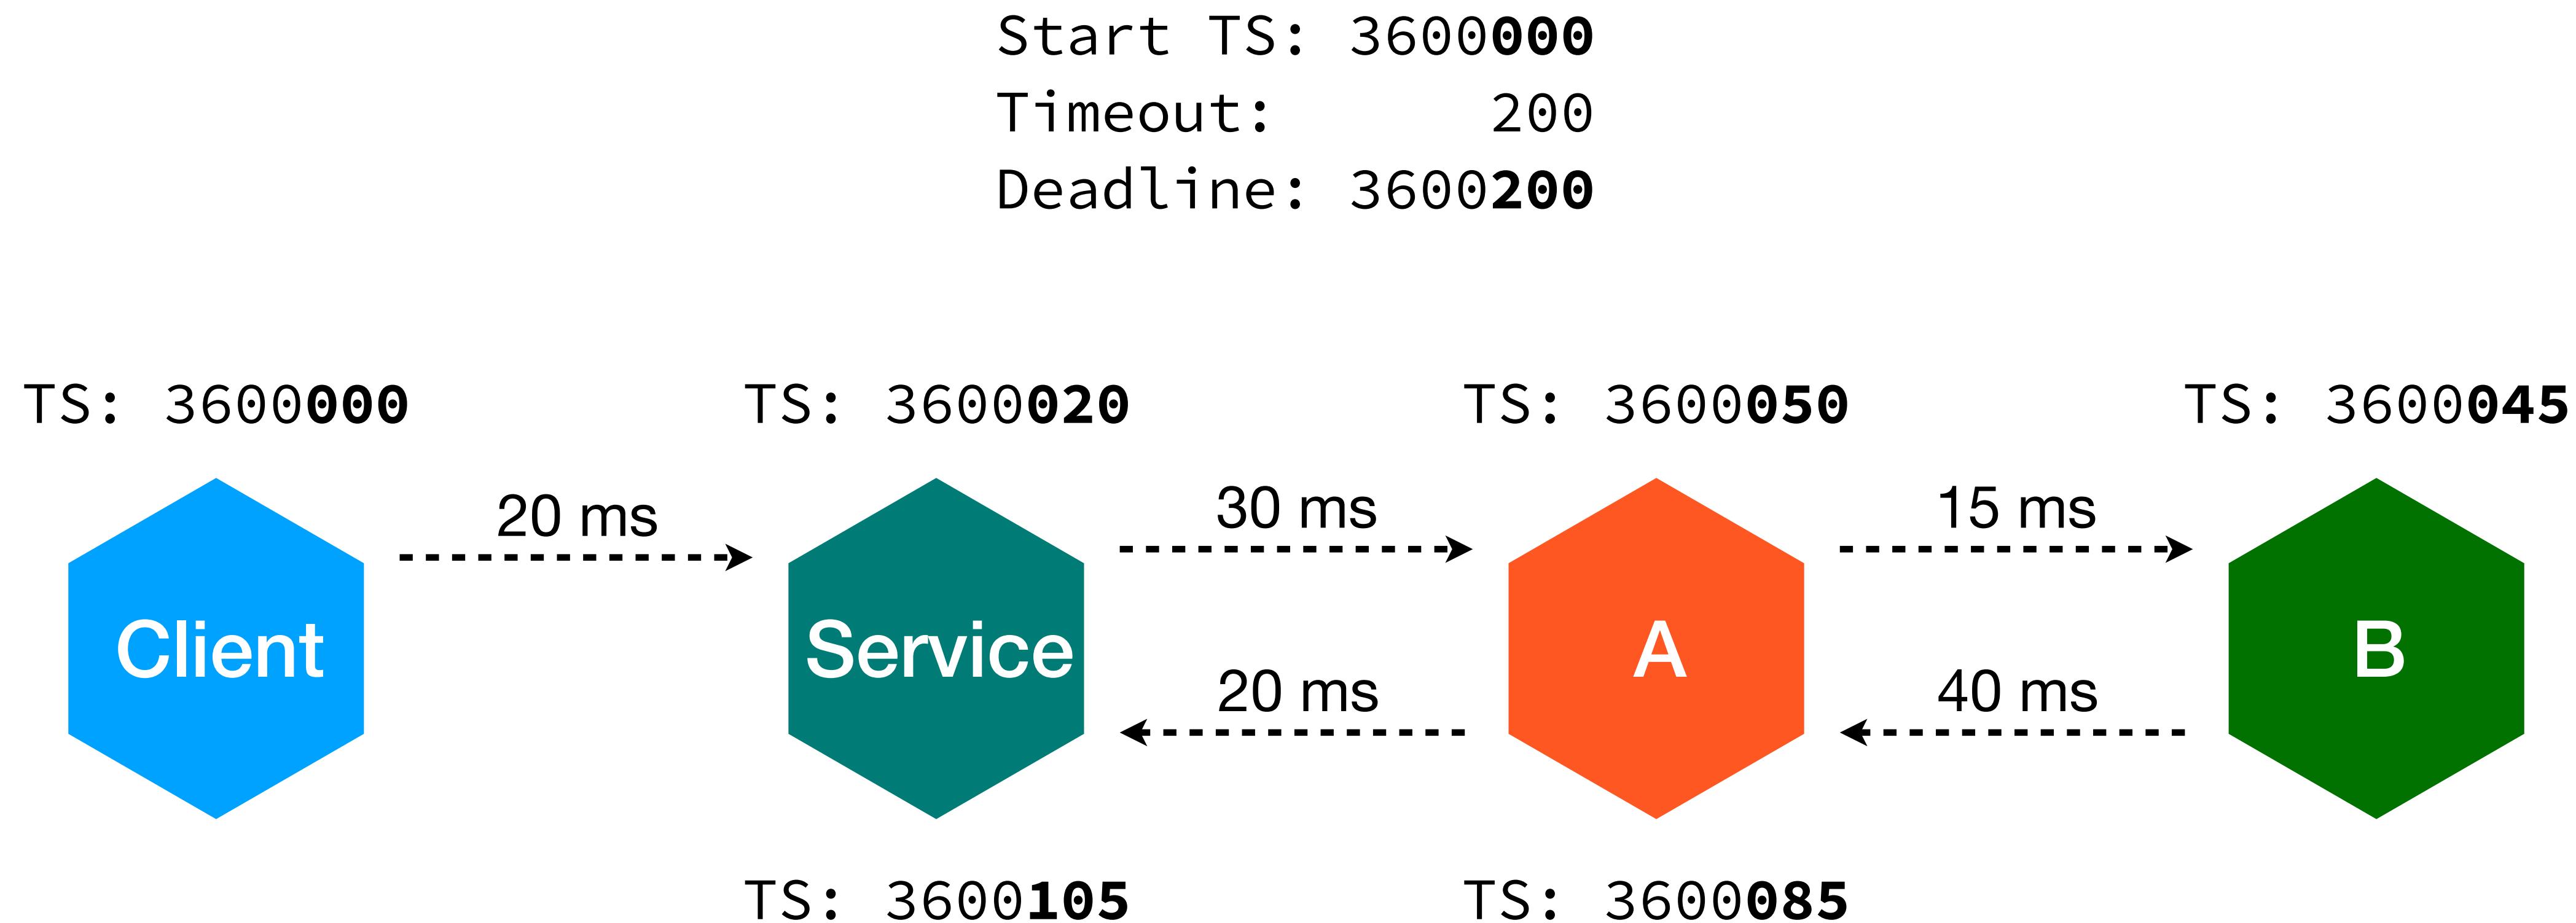
# Deadline propagation



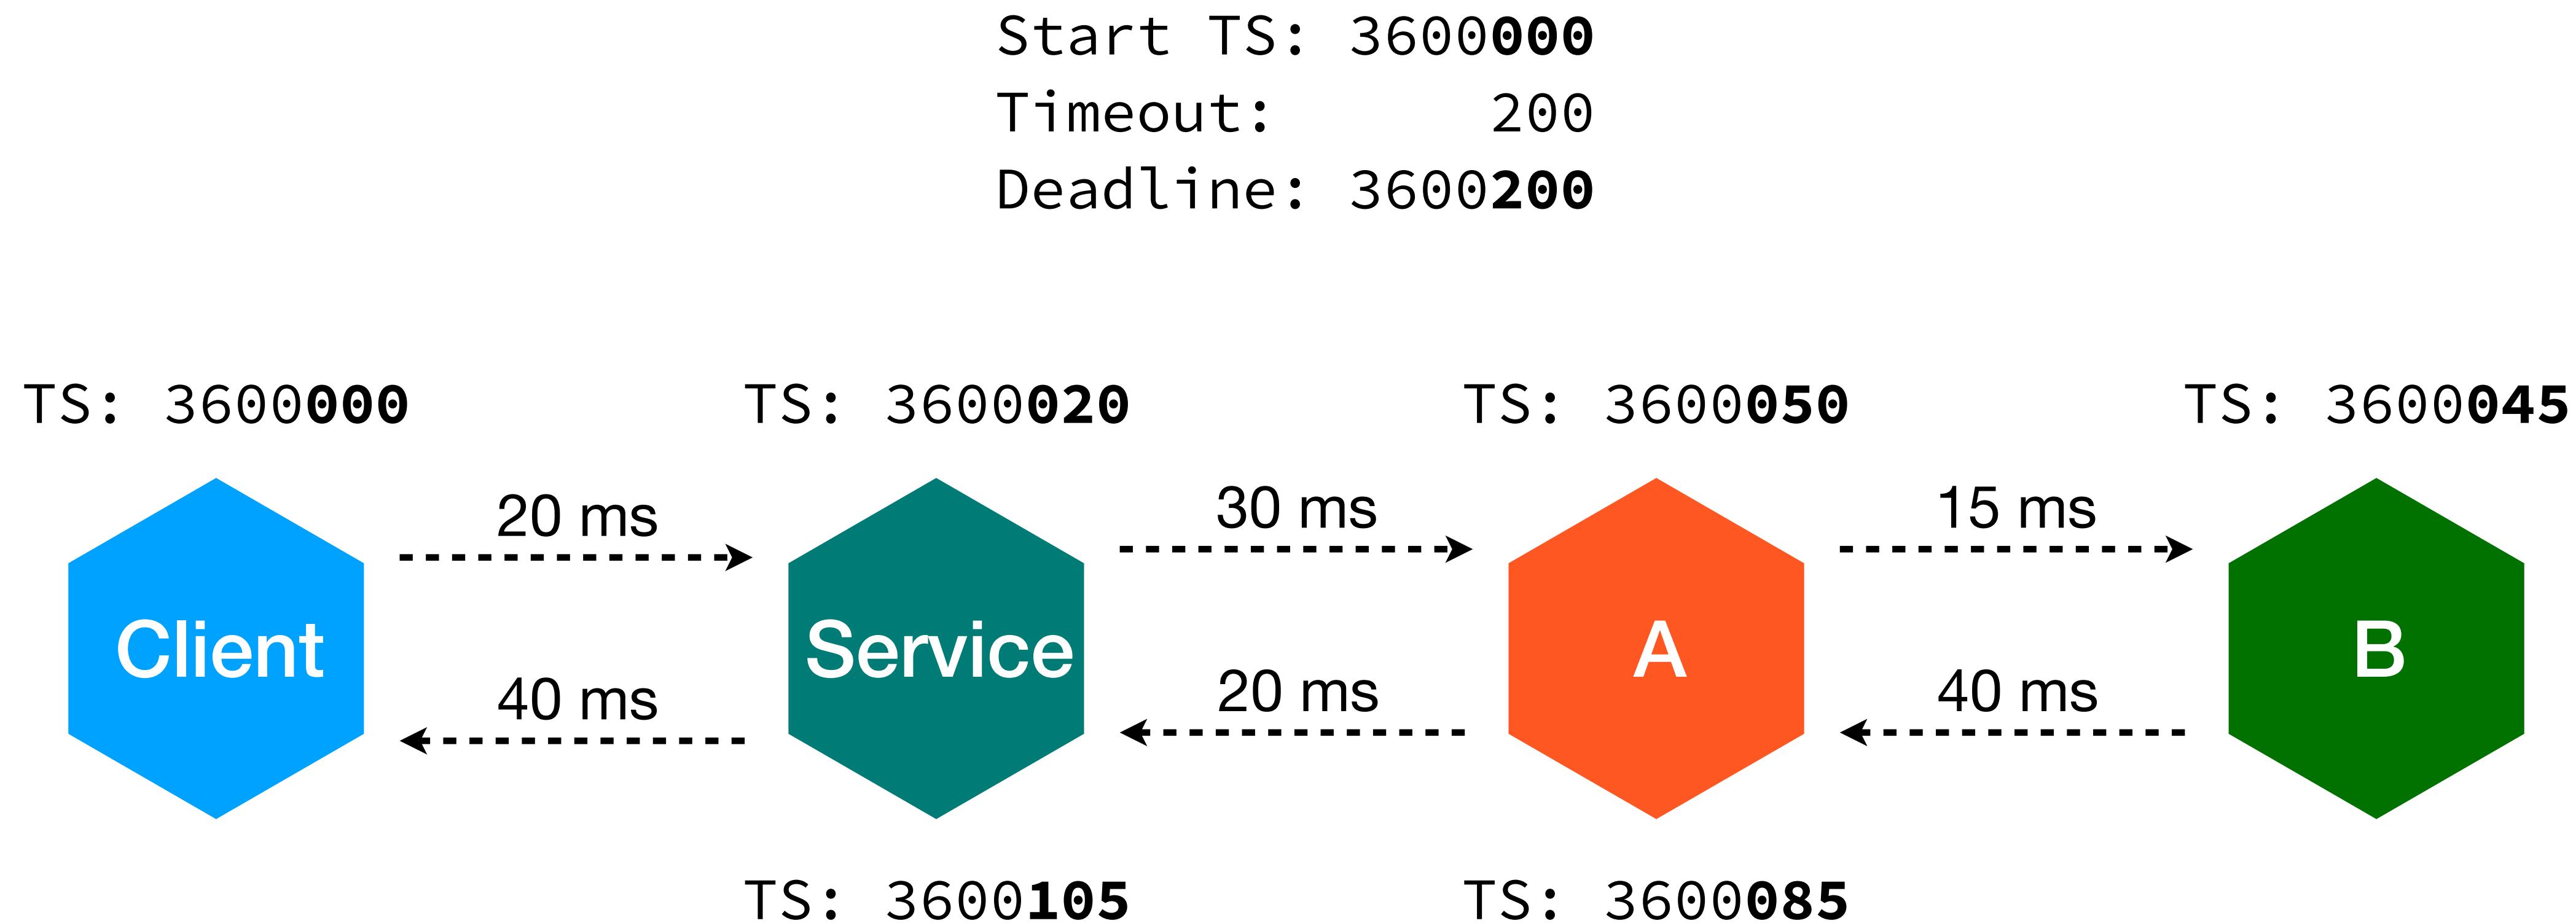
# Deadline propagation



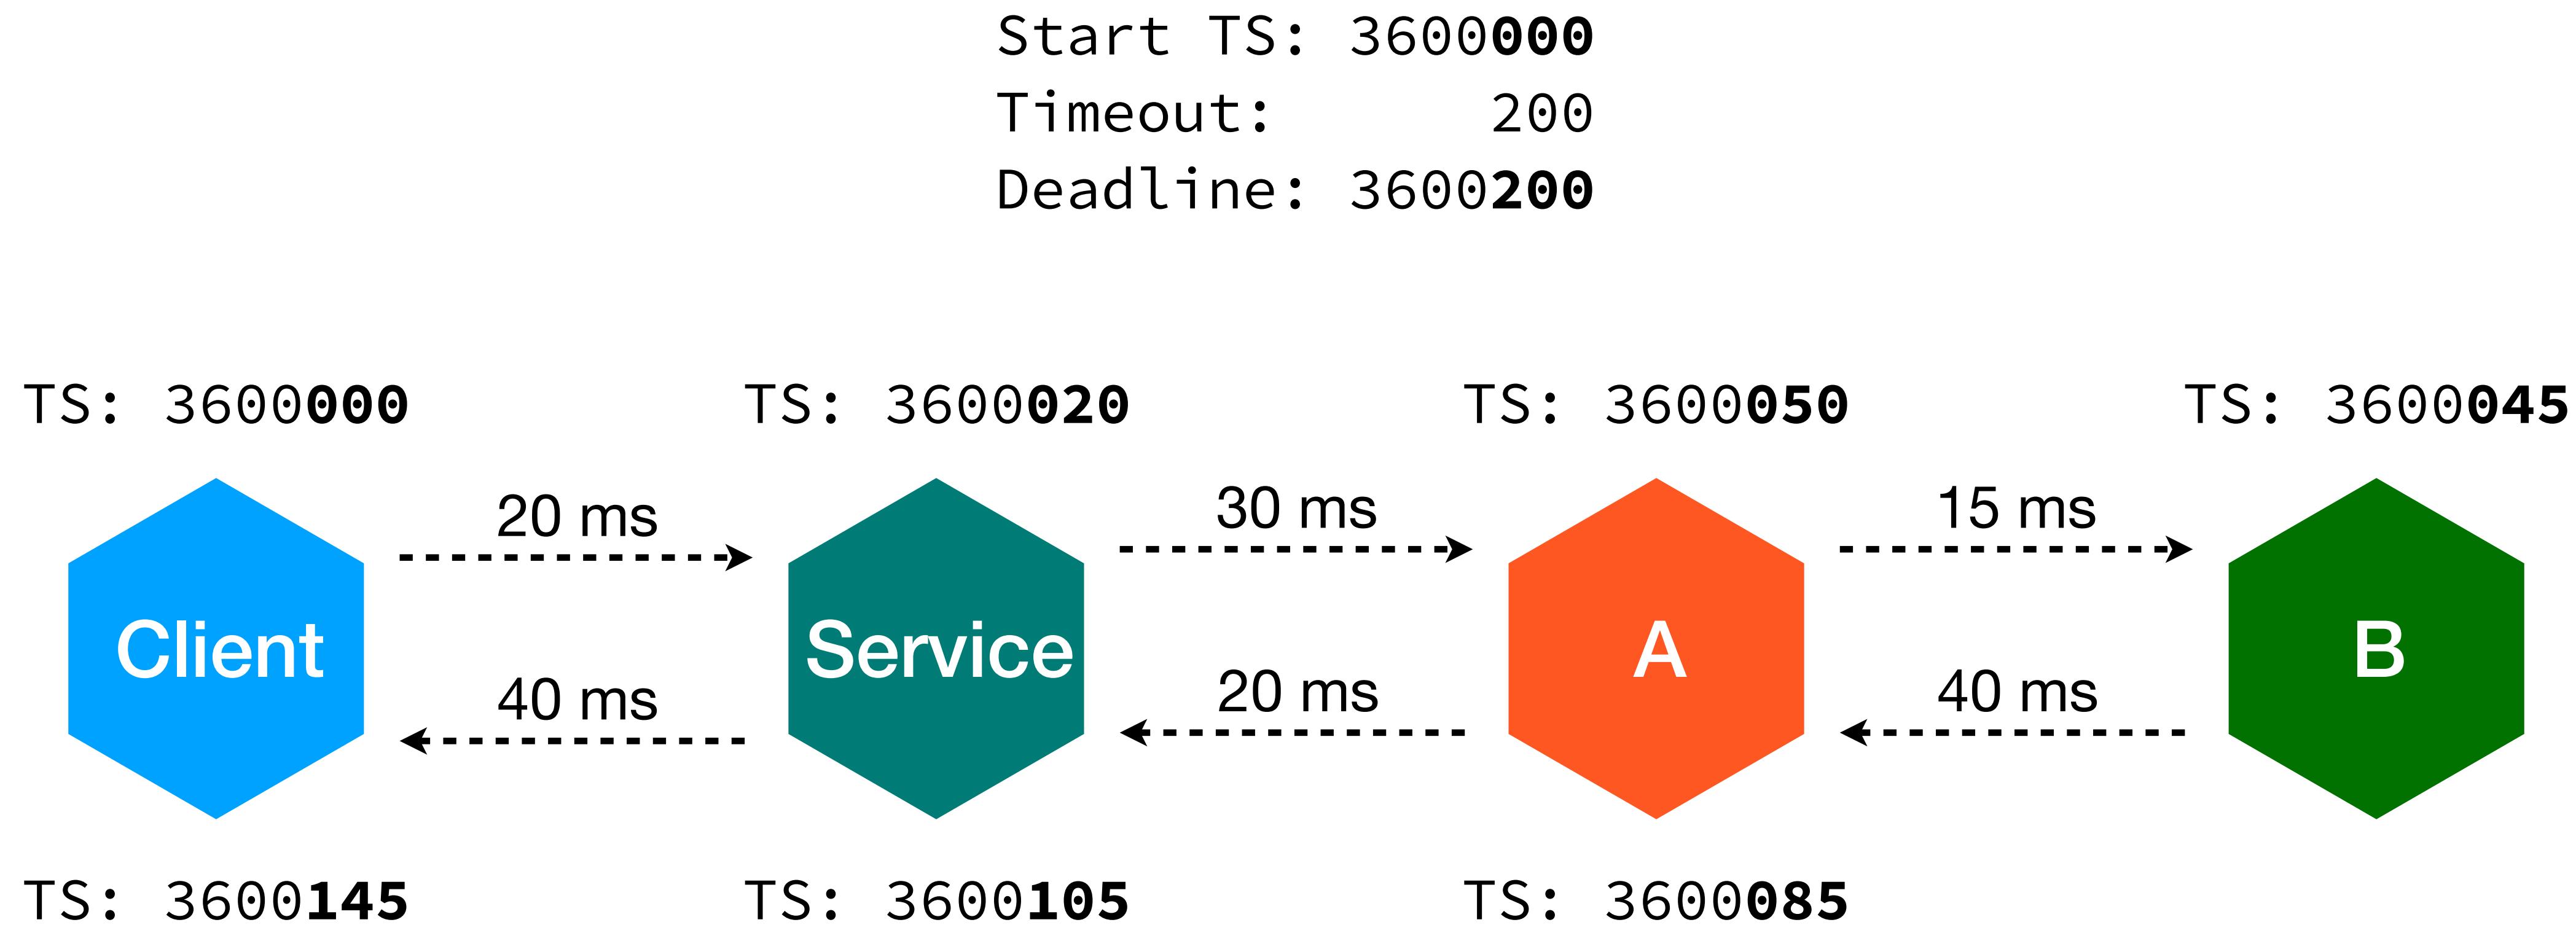
# Deadline propagation



# Deadline propagation

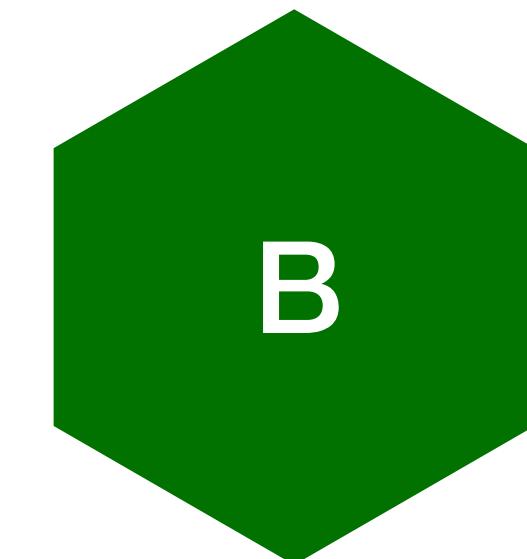
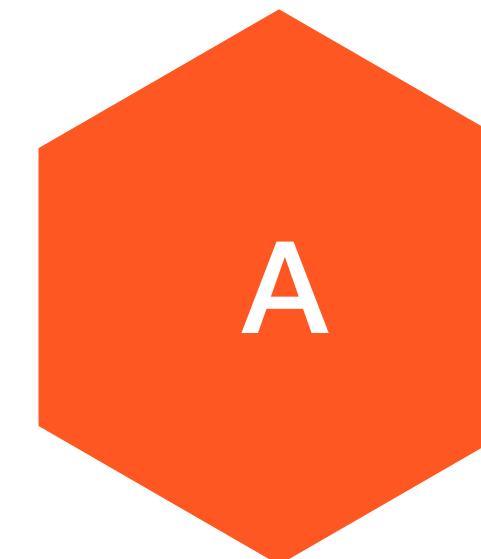


# Deadline propagation



# Deadline propagation

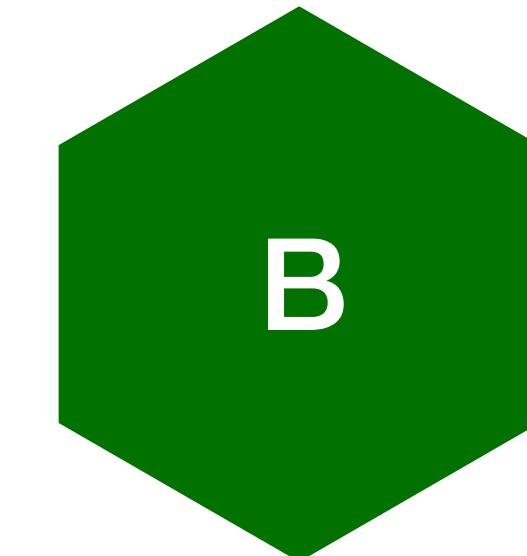
Start TS: **3600000**  
Timeout: 200  
Deadline: **3600200**



# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

TS: 3600000



# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

TS: 3600000



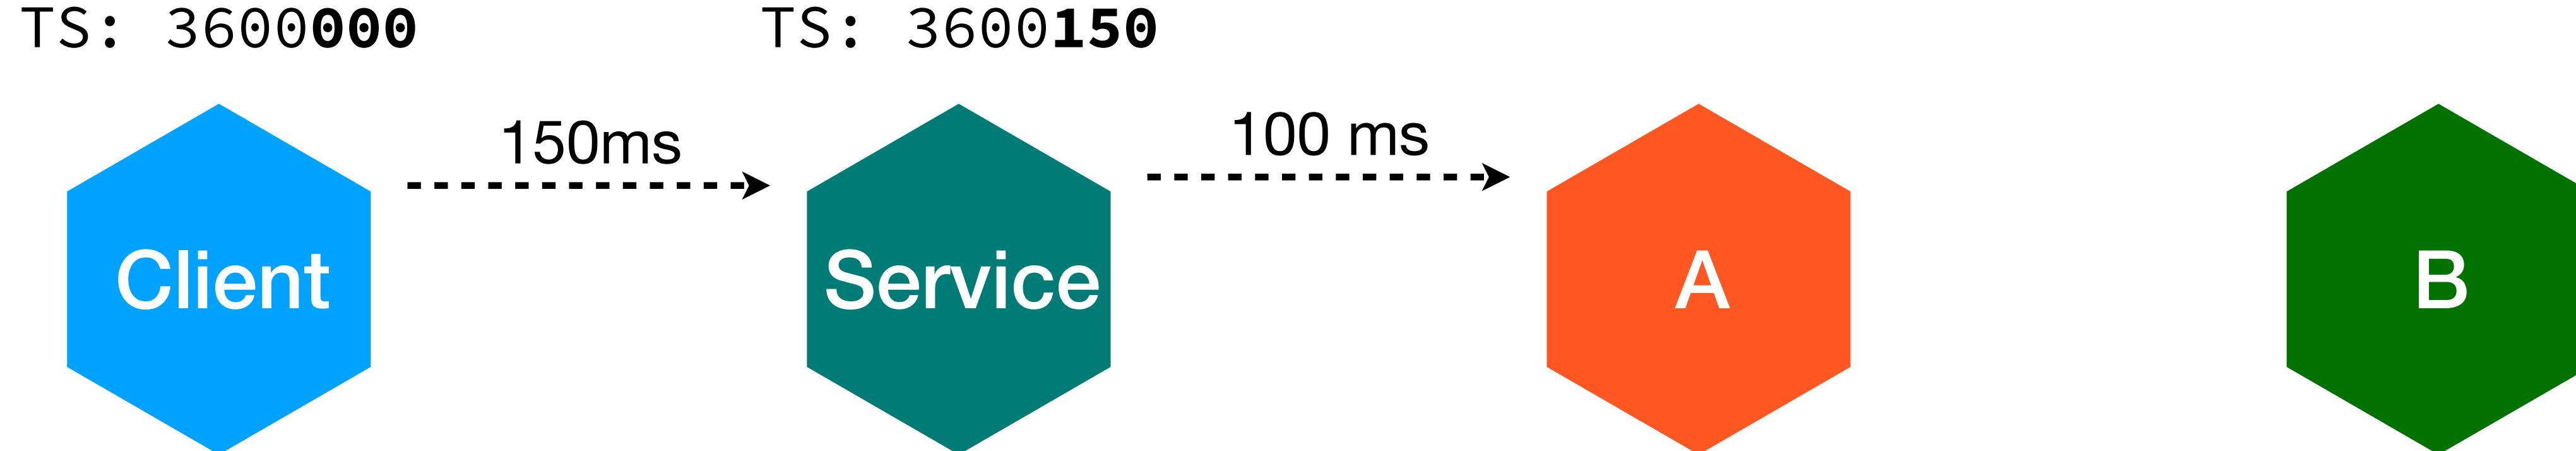
# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

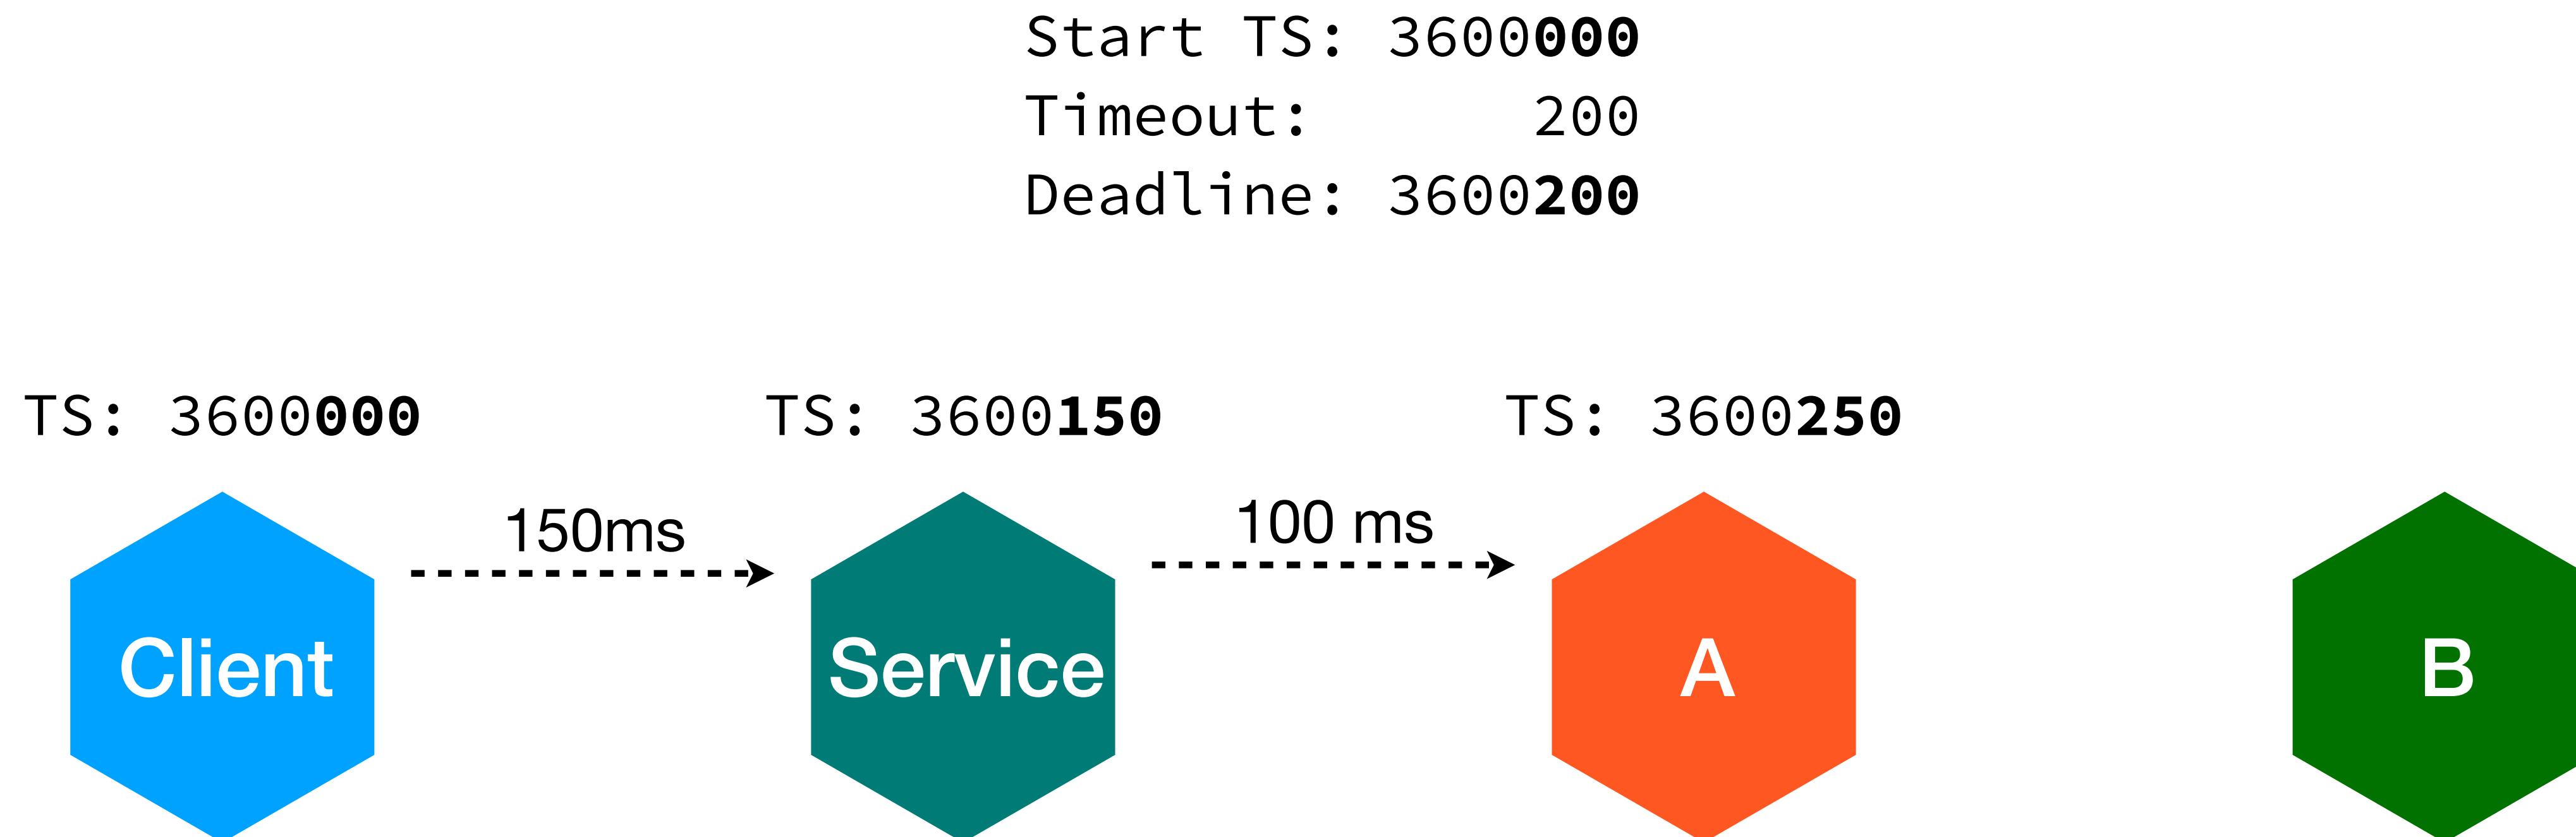


# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

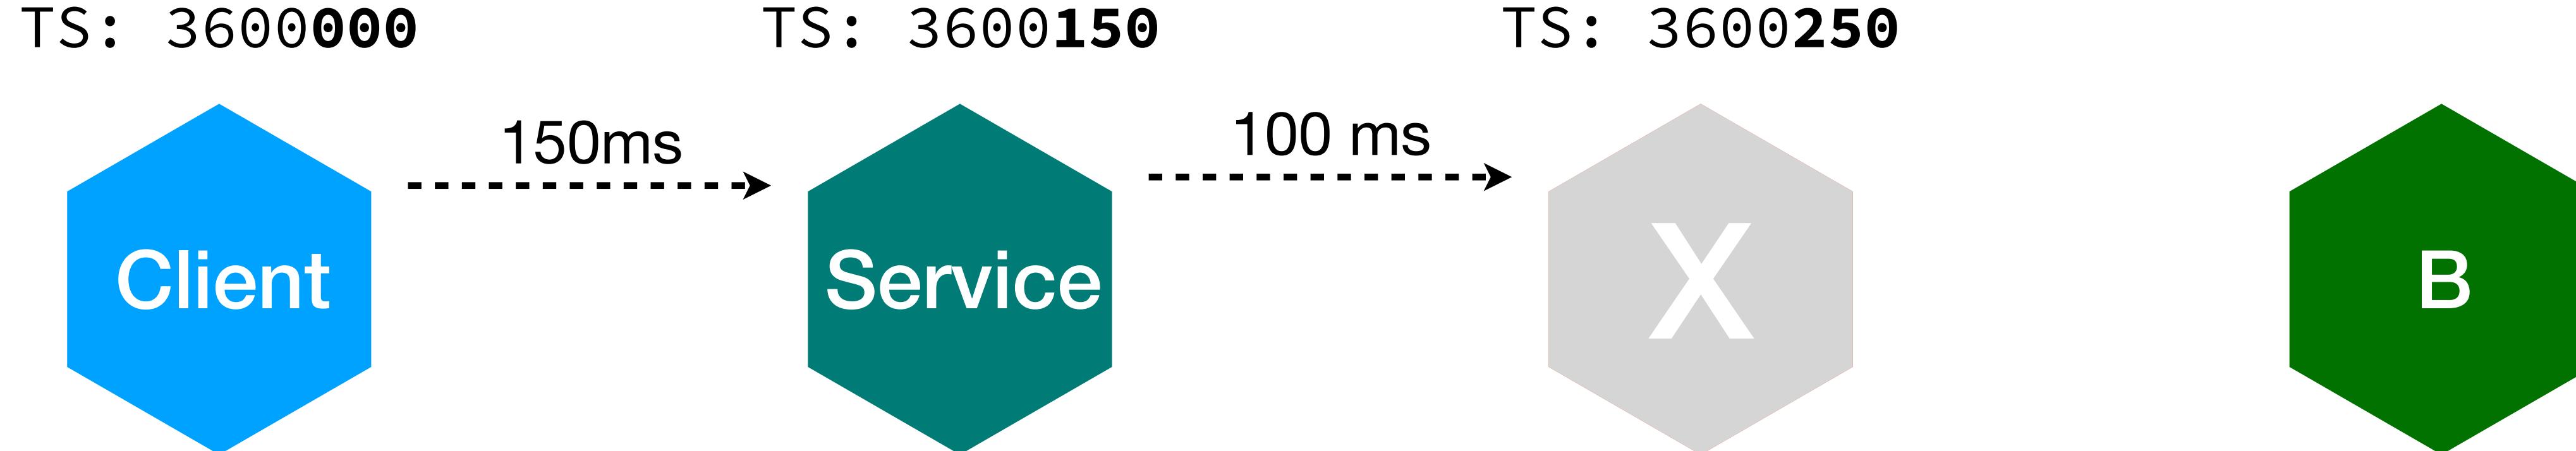


# Deadline propagation



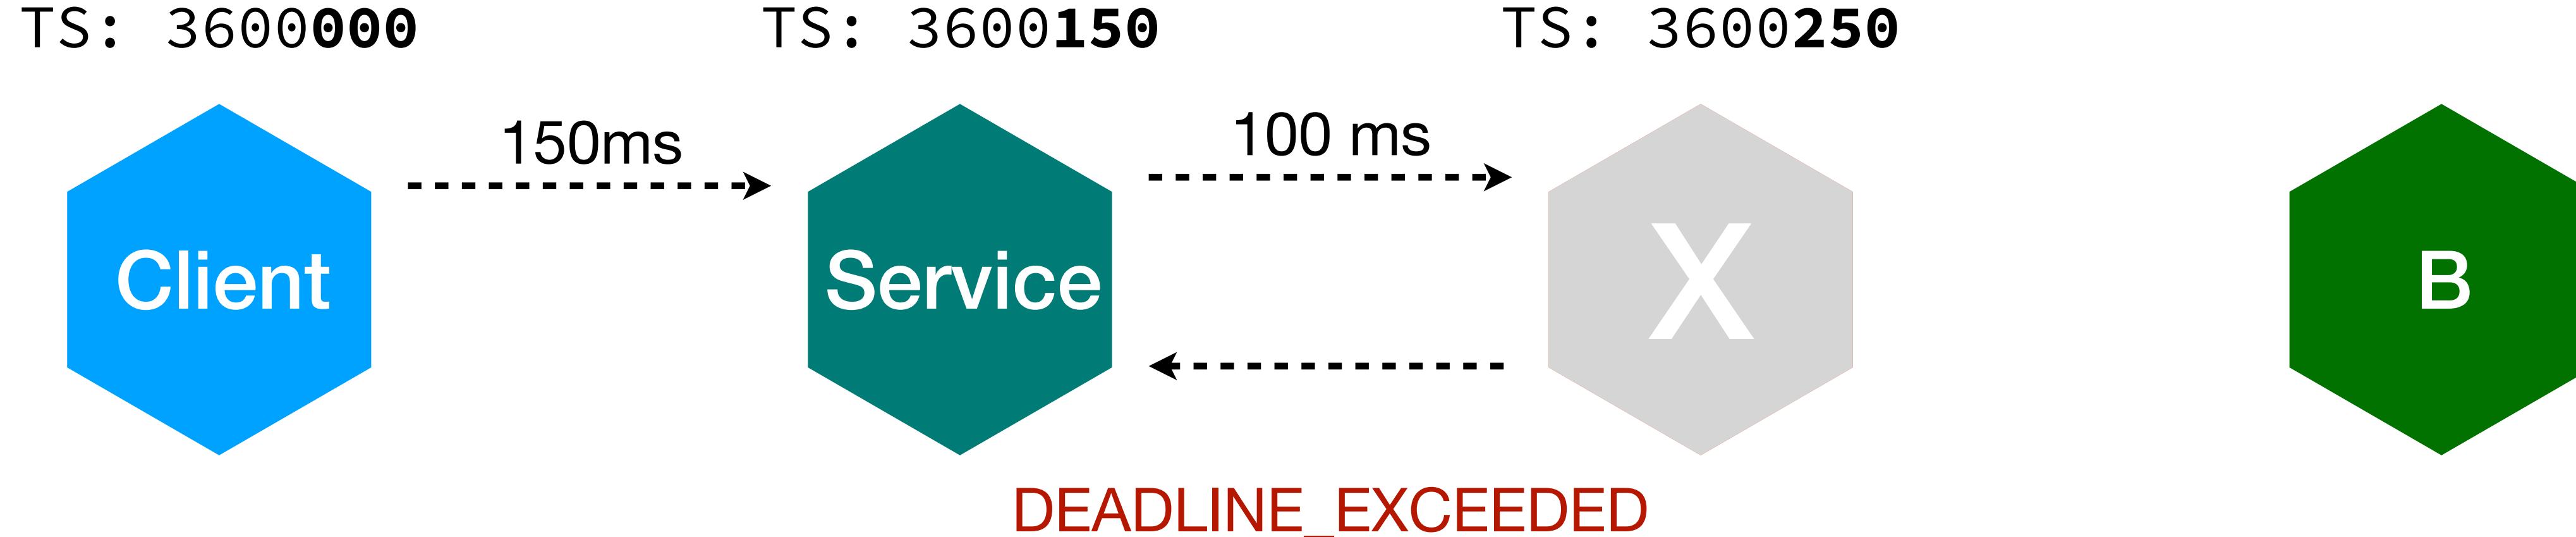
# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200

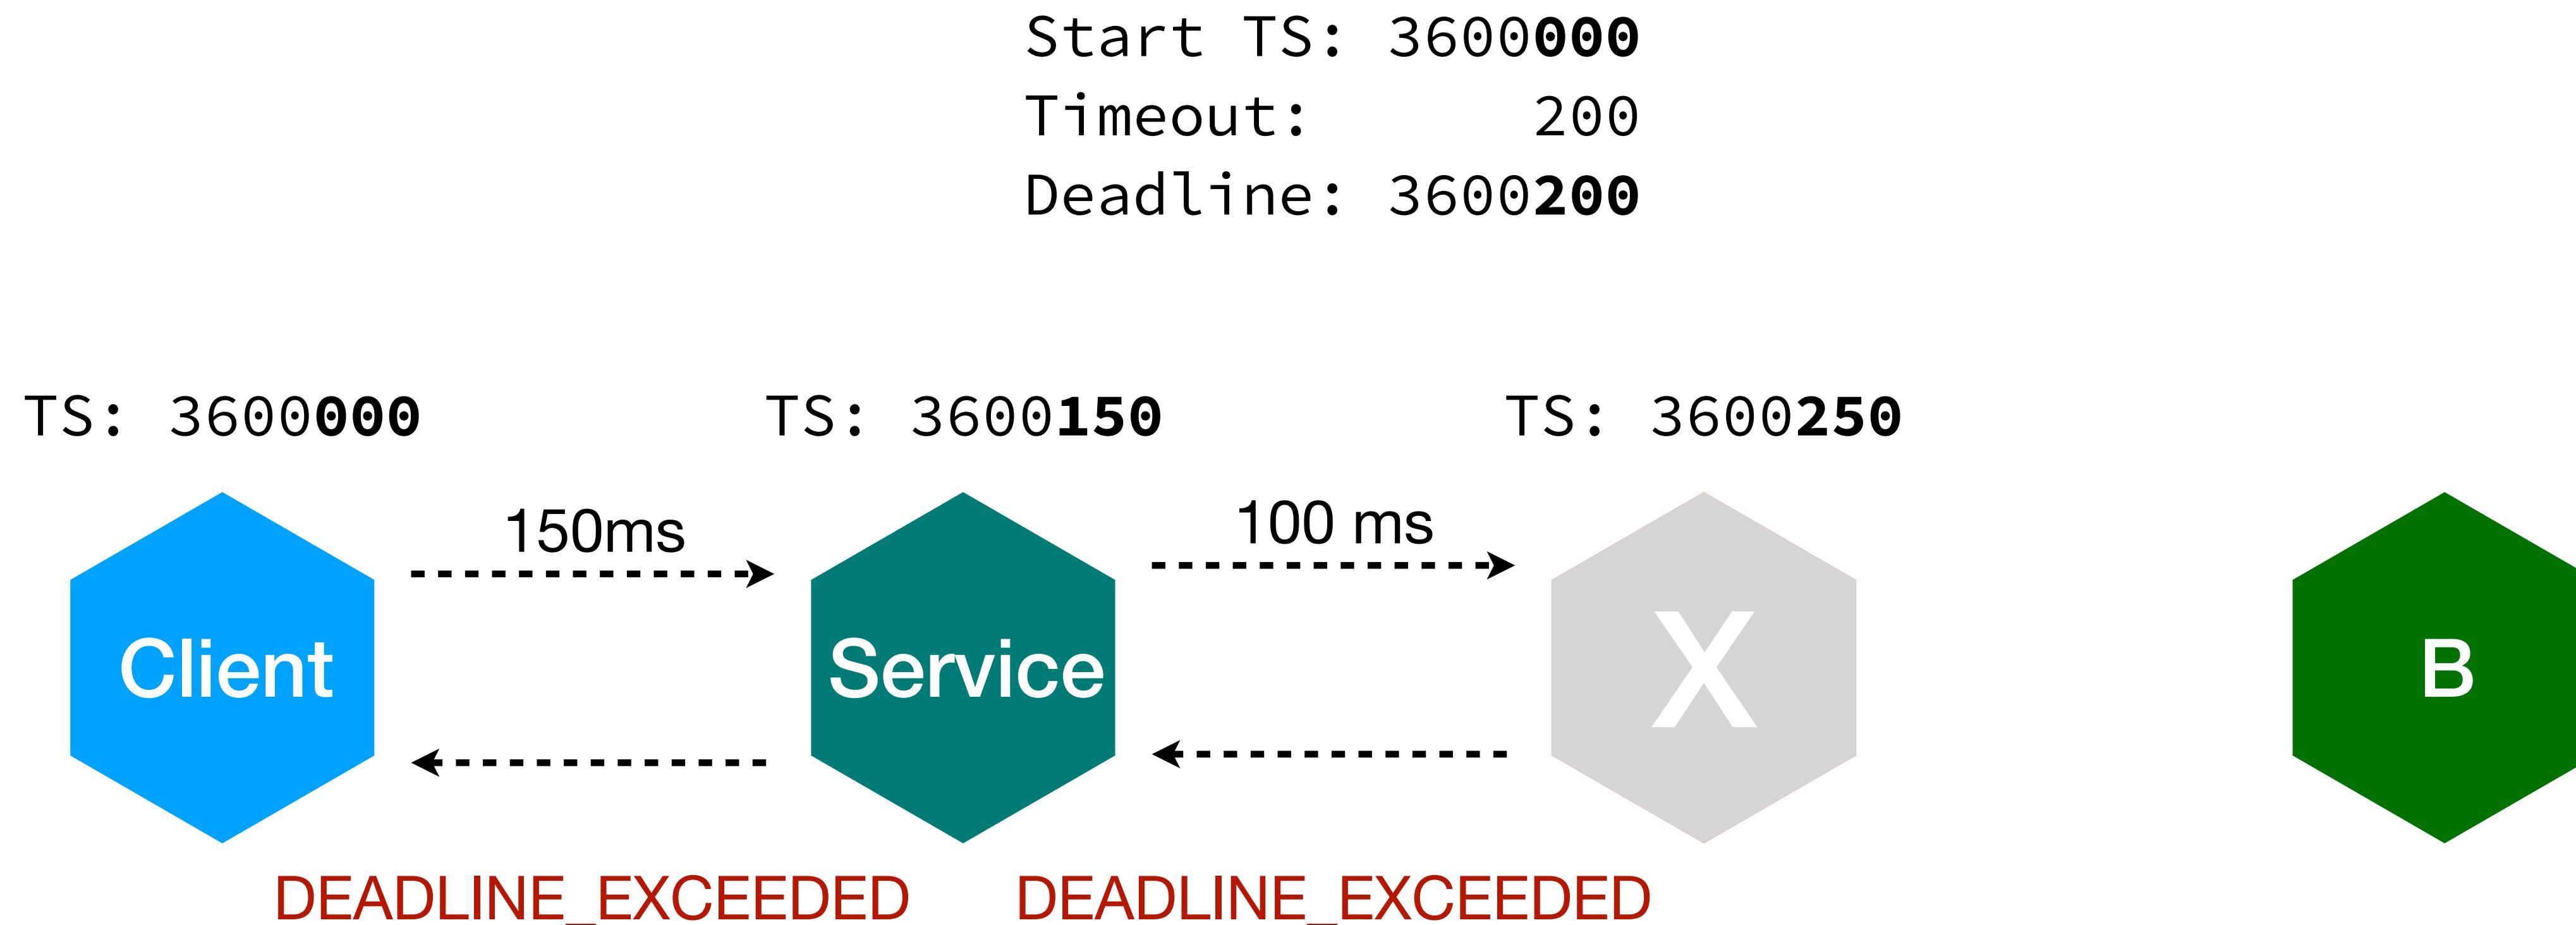


# Deadline propagation

Start TS: 3600000  
Timeout: 200  
Deadline: 3600200



# Deadline propagation



# Cancellation

# Cancellation

- Can be initiated by both client/server

# Cancellation

- Can be initiated by both client/server
- Immediately terminates the RPC

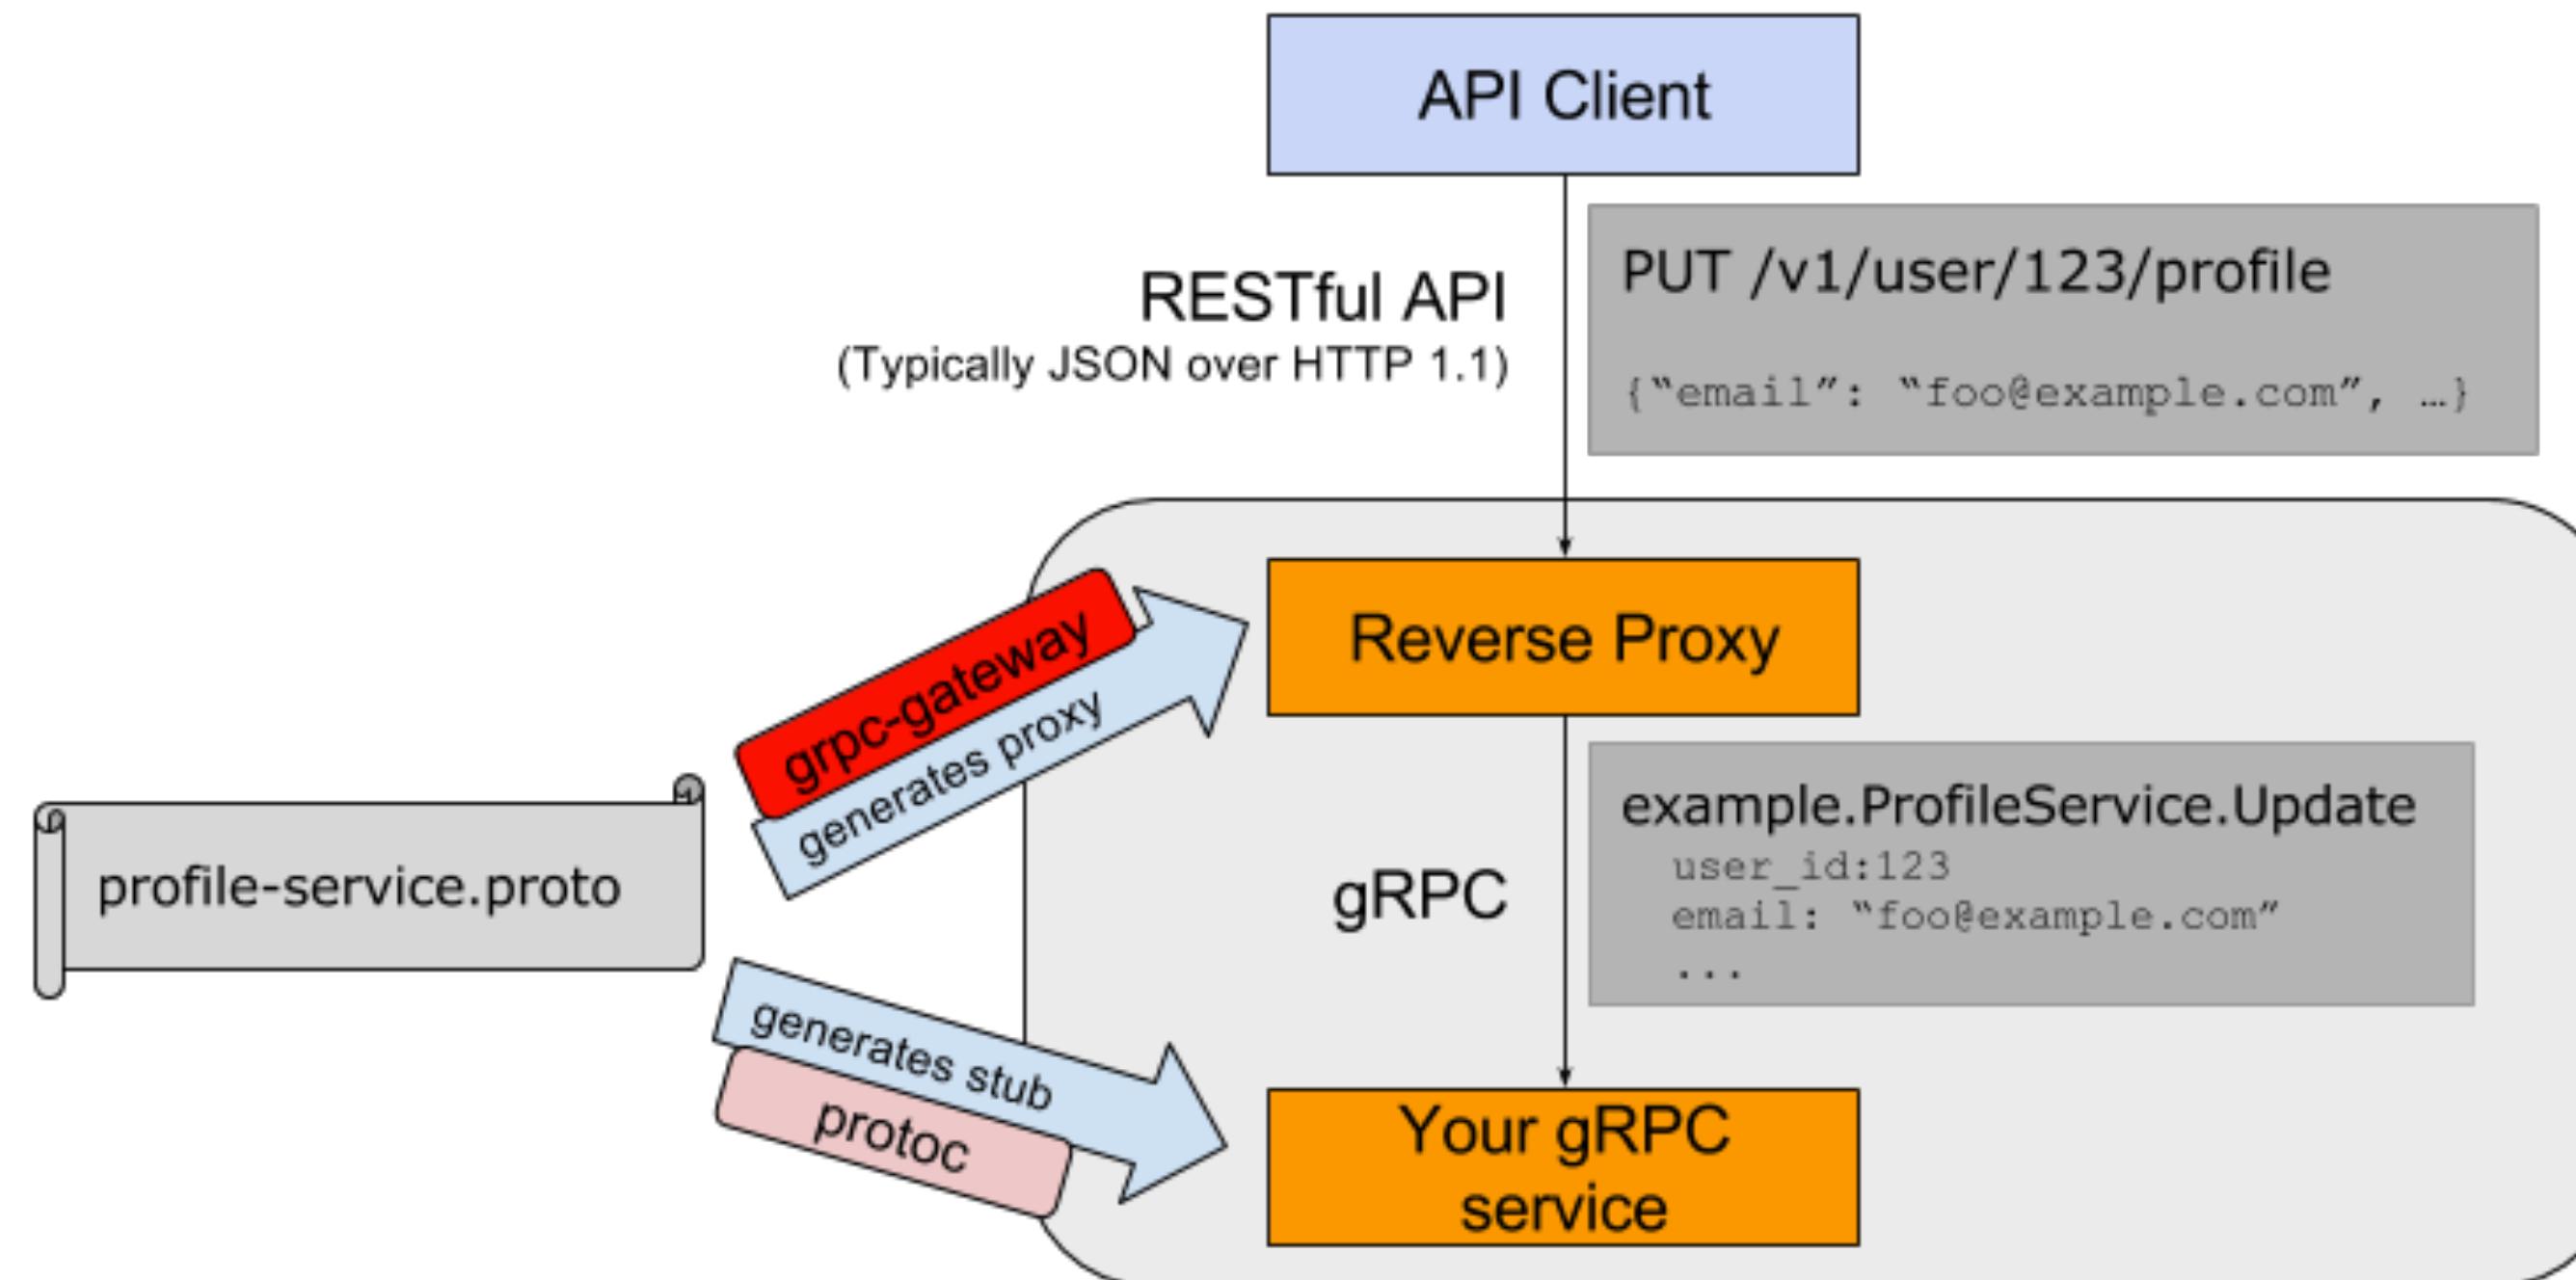
# Cancellation

- Can be initiated by both client/server
- Immediately terminates the RPC
- It is not a roll-back

# Cancellation

- Can be initiated by both client/server
- Immediately terminates the RPC
- It is not a roll-back
- Automatically cascaded

# Backward compatibility



<https://github.com/grpc-ecosystem/grpc-gateway>

# gRPC language support

- C++
- Python
- Java
- Go
- Ruby
- C#
- JS (Node)
- Android Java
- Objective-C
- PHP

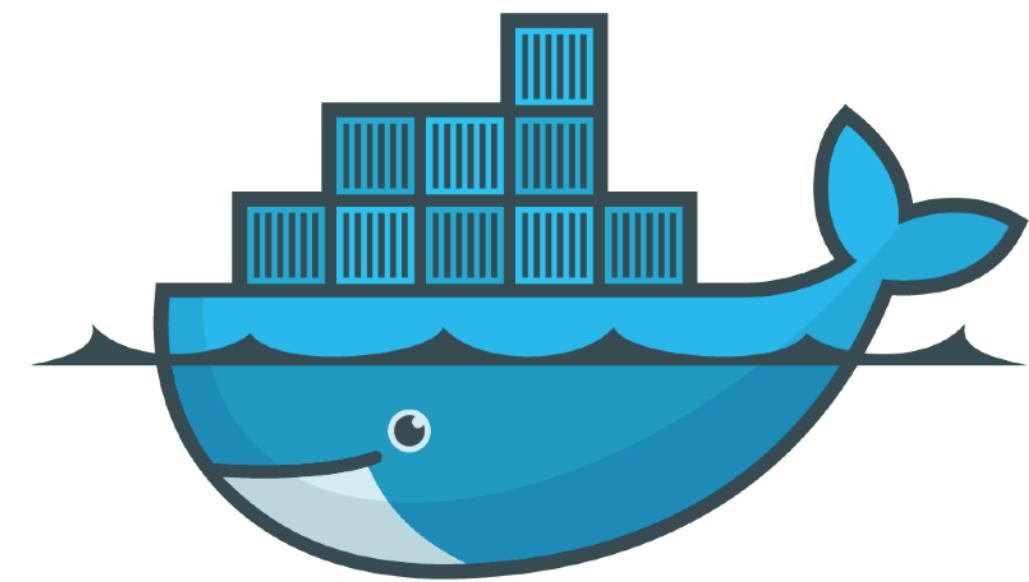
# gRPC Platform support

- Linux
- macOS
- Windows
- Android
- iOS

# Success stories

Google

# Success stories

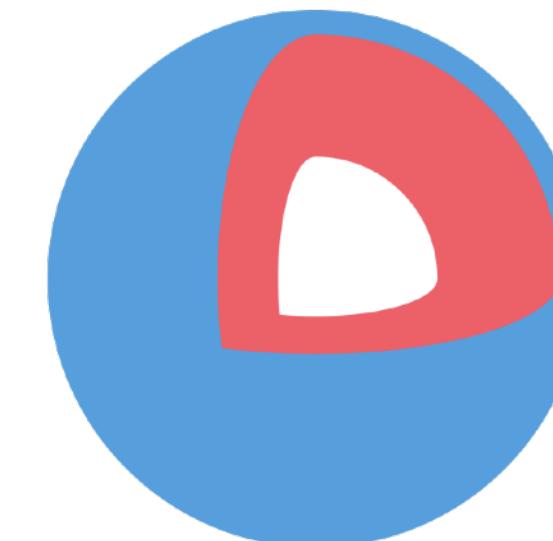


docker



Square

carbon3D™



Core OS

JUNIPER  
NETWORKS®

NETFLIX

# Benefits

# Benefits

- Focus on the API design & contract

# Benefits

- Focus on the API design & contract
- HTTP2 is awesome

# Benefits

- Focus on the API design & contract
- HTTP2 is awesome
- Bi-directional streaming

# Benefits

- Focus on the API design & contract
- HTTP2 is awesome
- Bi-directional streaming
- Freedom to pick any language

# Benefits

- Focus on the API design & contract
- HTTP2 is awesome
- Bi-directional streaming
- Freedom to pick any language
- Service-to-Service and Service-to-Mobile friendly

# Benefits

- Focus on the API design & contract
- HTTP2 is awesome
- Bi-directional streaming
- Freedom to pick any language
- Service-to-Service and Service-to-Mobile friendly
- Production ready

**Fin.**

**Thank you.**